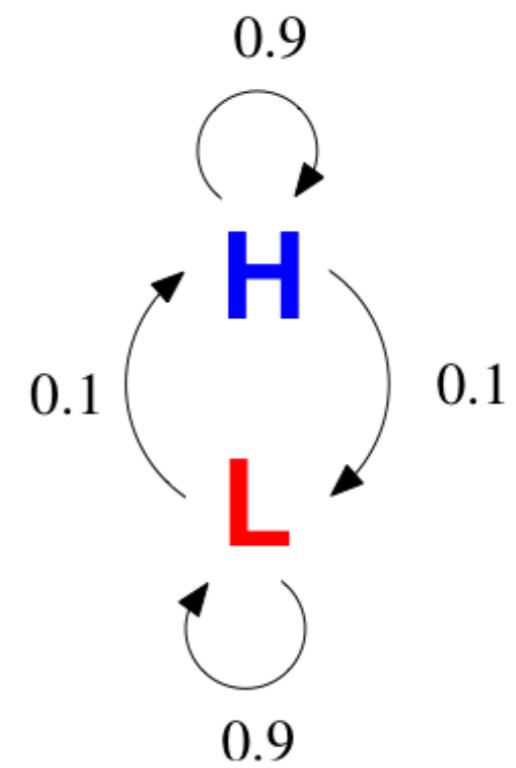
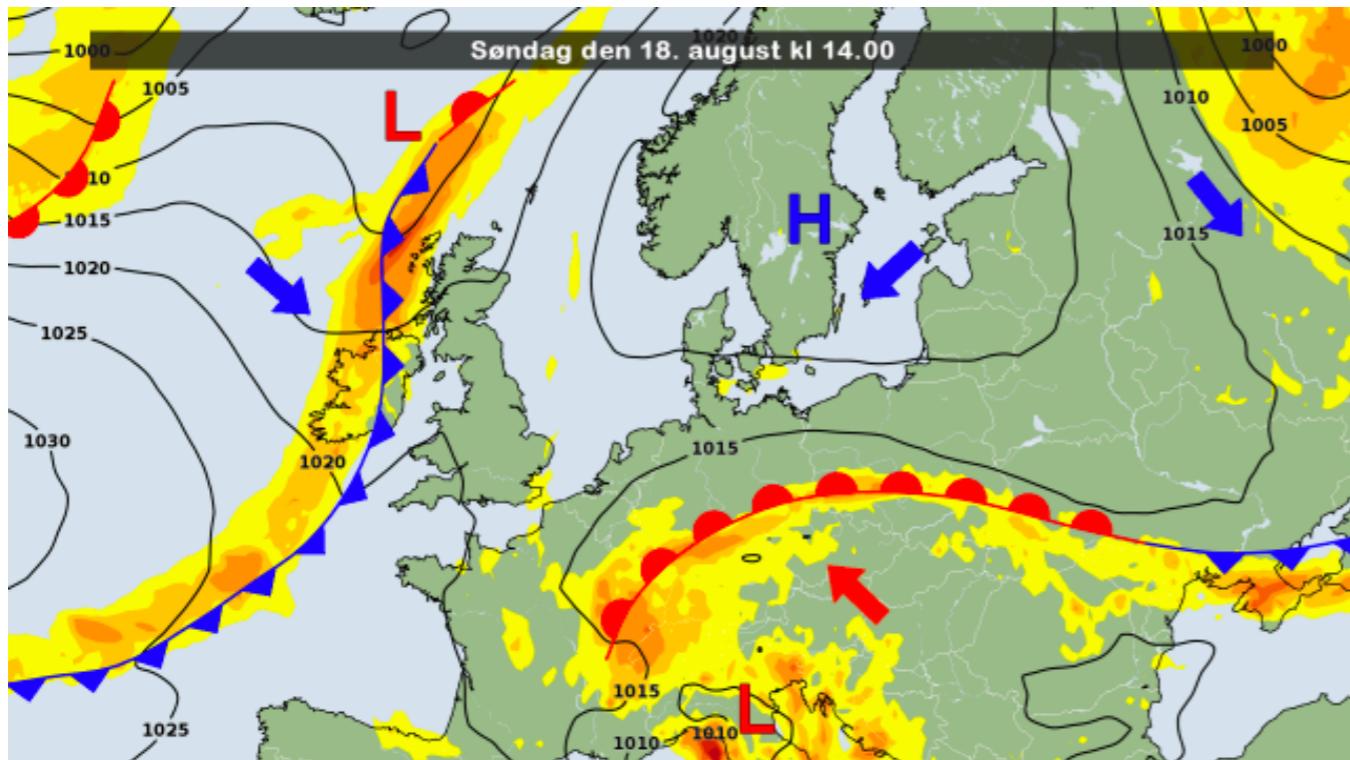


Hidden Markov models

Week 3

Kasper Munch

A weather model

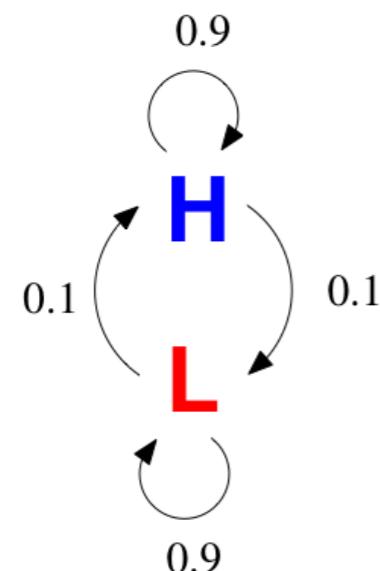


Mandag Tirsdag Onsdag Torsdag Fredag Lørdag Søndag Mandag Tirsdag Onsdag Torsdag Fredag

Hidden Markov model

Transitions and emissions

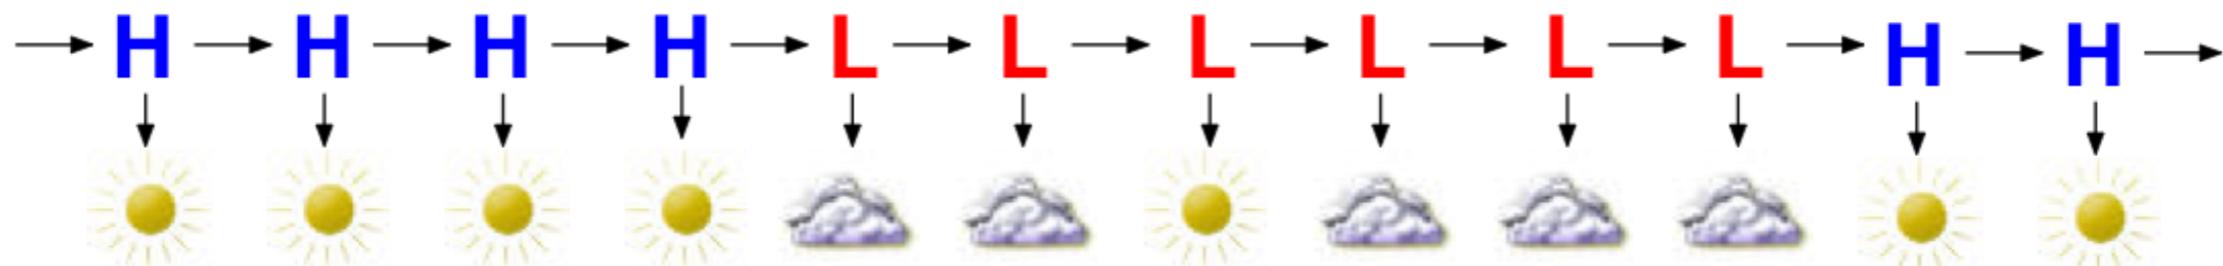
Transition probabilities:



Emission probabilities:

- H**: $P(\text{☀️}) = 0.7, P(\text{☁️}) = 0.3$
- L**: $P(\text{☀️}) = 0.3, P(\text{☁️}) = 0.7$

Hidden Markov chain:



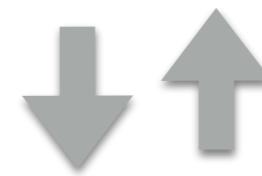
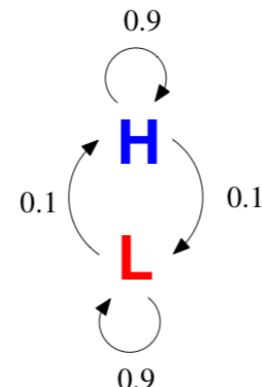
Watching the weather channel

Finding the most likely model parameters

Observations:



Model parameters:



$$\begin{aligned} \mathbf{H}: \quad P(\text{sun}) &= 0.7, & P(\text{cloud}) &= 0.3 \\ \mathbf{L}: \quad P(\text{sun}) &= 0.3, & P(\text{cloud}) &= 0.7 \end{aligned}$$

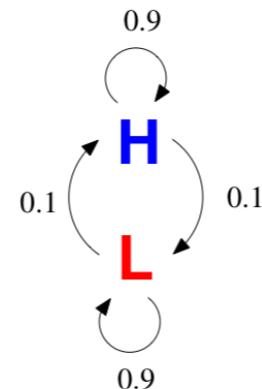
Guessing air pressure yourself

Finding the most likely hidden states

Observations:



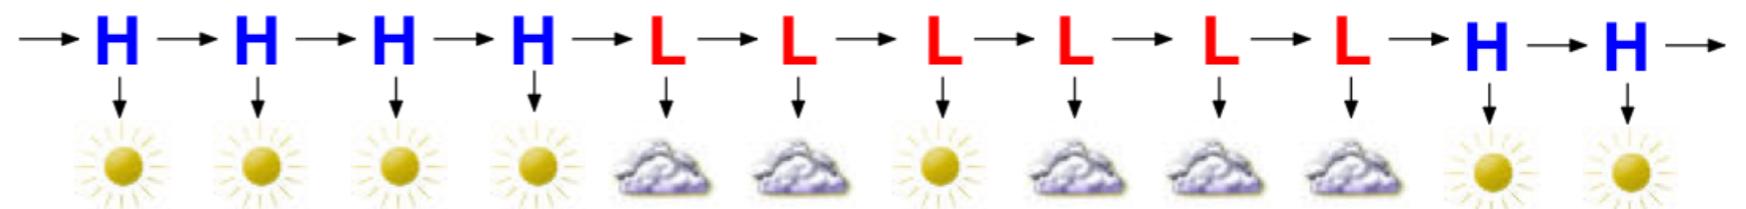
Model parameters:



$$\begin{aligned} \mathbf{H}: \quad P(\text{sun}) &= 0.7, & P(\text{cloud}) &= 0.3 \\ \mathbf{L}: \quad P(\text{sun}) &= 0.3, & P(\text{cloud}) &= 0.7 \end{aligned}$$



Most likely sequence
of hidden states:



Hidden Markov model

Why are they useful?

1. We can make the HMM **describe data** by optimising its **parameters**.
2. Using the HMM, we can estimate the **most likely sequence of hidden states**.
3. Using the HMM, we can estimate the **probability of each hidden state at each position** in the sequence

Hidden Markov model

Why are they useful?

2. We can make the HMM **describe data** by optimising its **parameters**.
1. Using the HMM, we can estimate the **most likely sequence of hidden states**.
3. Using the HMM, we can estimate the **probability of each hidden state at each position** in the sequence

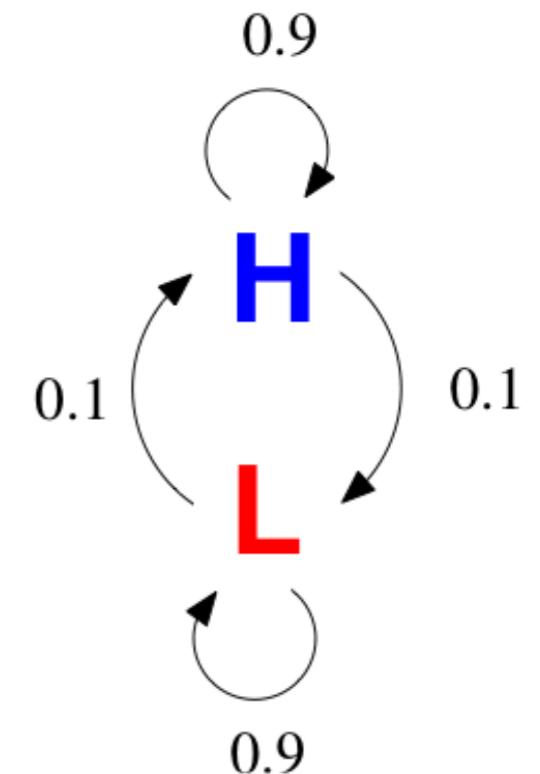
Markov chain

→ H → H → H → H → L → L → L → L → L → L → H → H →

The Markov property:

The probability of the state at the next position in the chain only depends on the state at the current position.

i.e. the process has no memory.



Markov chain

$$\pi_1 \quad \pi_2 \quad \pi_3 \quad \pi_4 \quad \pi_5 \quad \pi_6 \quad \cdots \quad \pi_L$$

· **H** → **H** → **L** → **L** → **L** → **L** → **L** → **L**

$$\begin{aligned} P(\pi) &= P(\pi_L | \pi_{L-1}) \ P(\pi_{L-1} | \pi_{L-2}) \ \cdots \ P(\pi_2 | \pi_1) \ P(\pi_1) \\ &= P(\pi_1) \prod_{i=2}^L P(\pi_i | \pi_{i-1}) \end{aligned}$$

5 min discussion

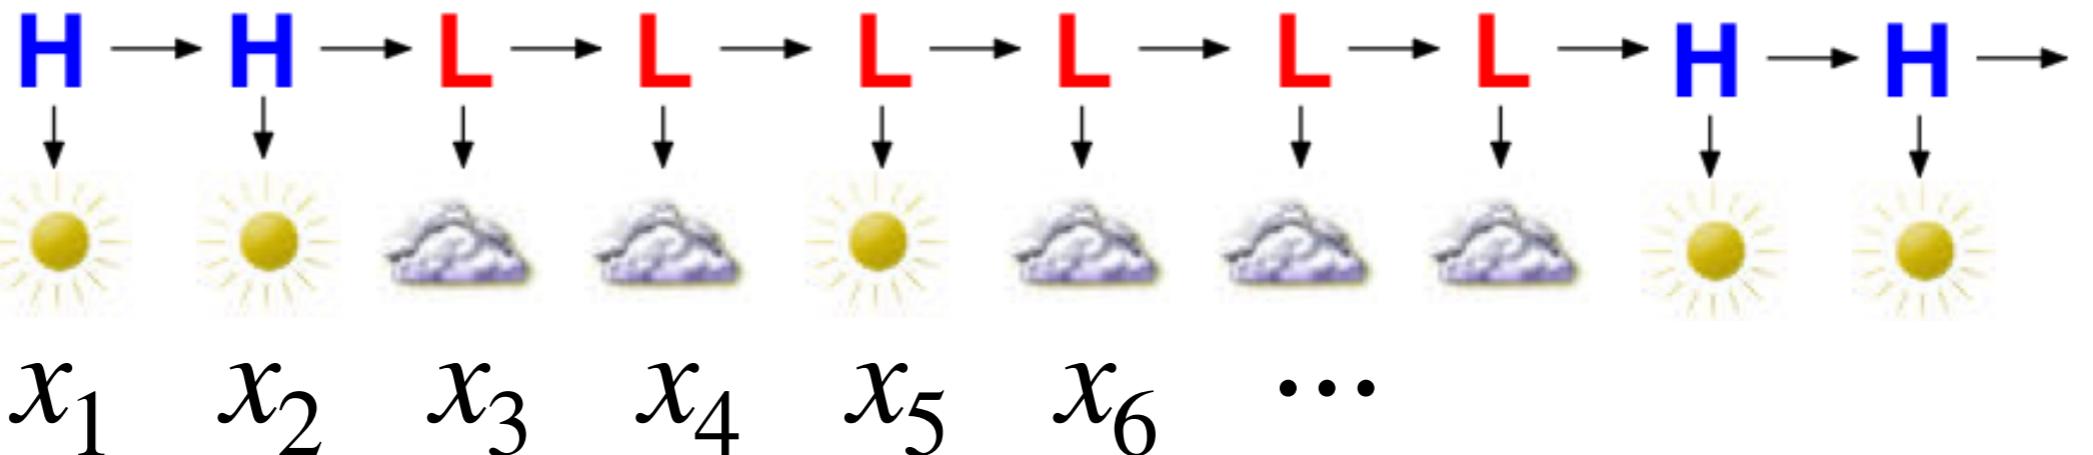
- What is a state and what is an observable?
- What is hidden and what is observed?
- What is a Markov chain?

Hidden Markov chain

Hidden Markov chain: π



$\pi_1 \quad \pi_2 \quad \pi_3 \quad \pi_4 \quad \pi_5 \quad \pi_6 \quad \dots$



Observed sequence: x

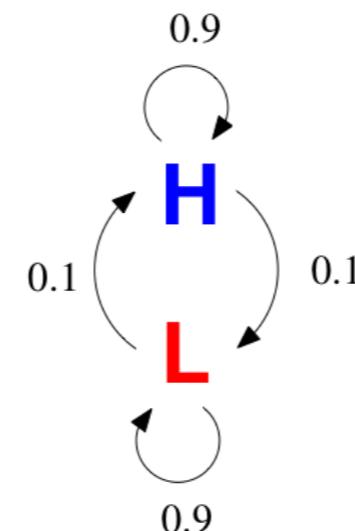
Hidden Markov model

State path: $\pi_1 \quad \pi_2 \quad \pi_3 \quad \pi_4 \quad \pi_5 \quad \pi_6 \quad \dots$

Sequence symbols: $x_1 \quad x_2 \quad x_3 \quad x_4 \quad x_5 \quad x_6 \quad \dots$

Transition probability:

$$a_{kl} = P(\pi_i = l \mid \pi_{i-1} = k)$$



Emission probability:

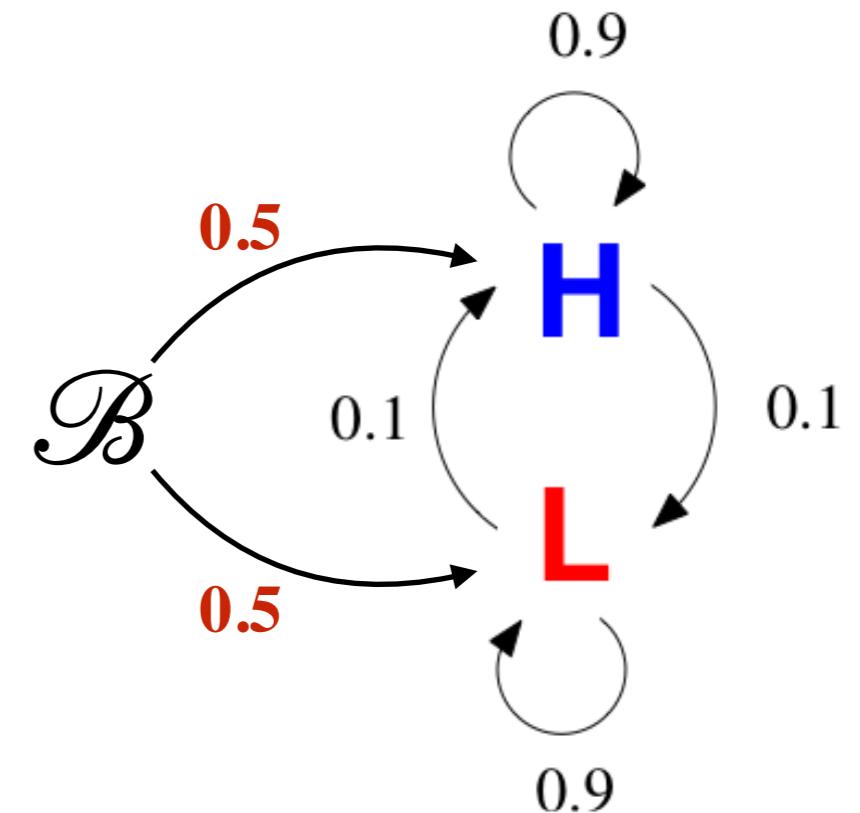
$$e_k(b) = P(x_i = b \mid \pi_i = k)$$

H: $P(\text{☀}) = 0.7, \quad P(\text{☁}) = 0.3$

L: $P(\text{☀}) = 0.3, \quad P(\text{☁}) = 0.7$

A begin state

Transition probabilities from begin state reflect prior belief:

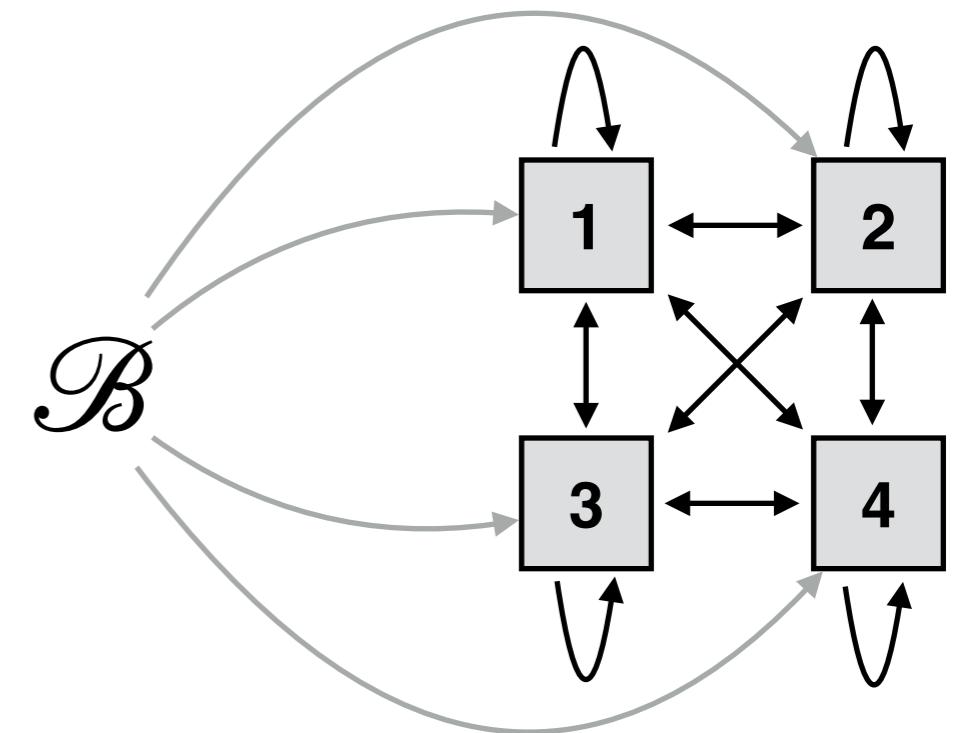


Transition probabilities from the null state represents prior belief

Joint probability

Of the observables x and hidden states π

π_1	π_2	π_3	π_4	π_5	π_6	\cdots			
1	2	2	1	1	4	4	1	3	3
\downarrow									
A	G	G	T	C	A	T	C	A	T
x_1	x_2	x_3	x_4	x_5	x_6	\cdots			

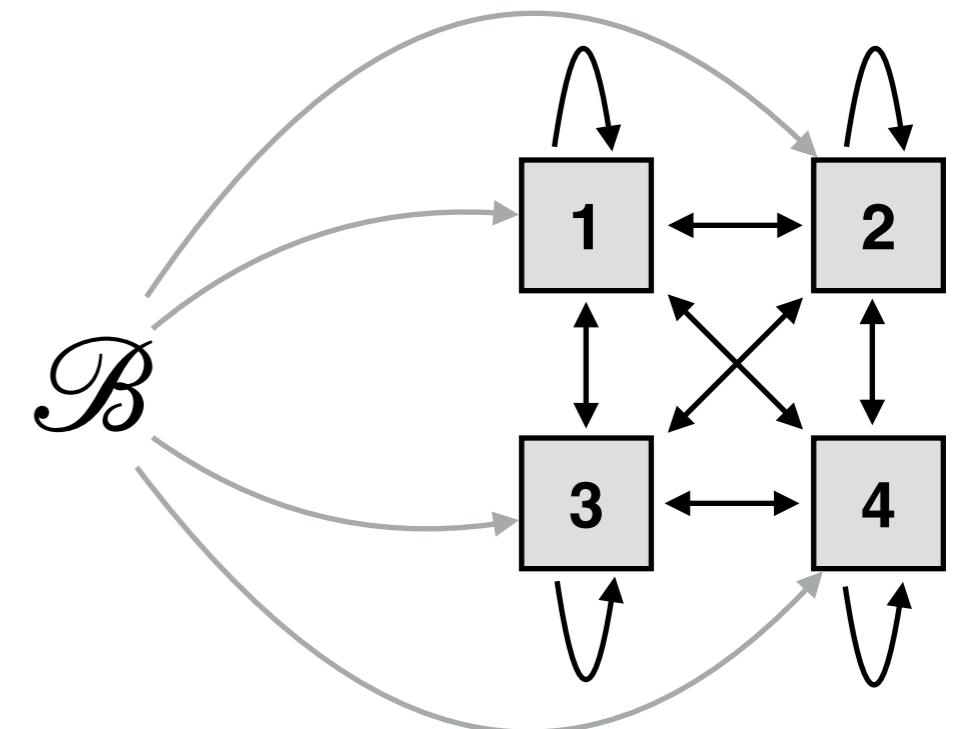
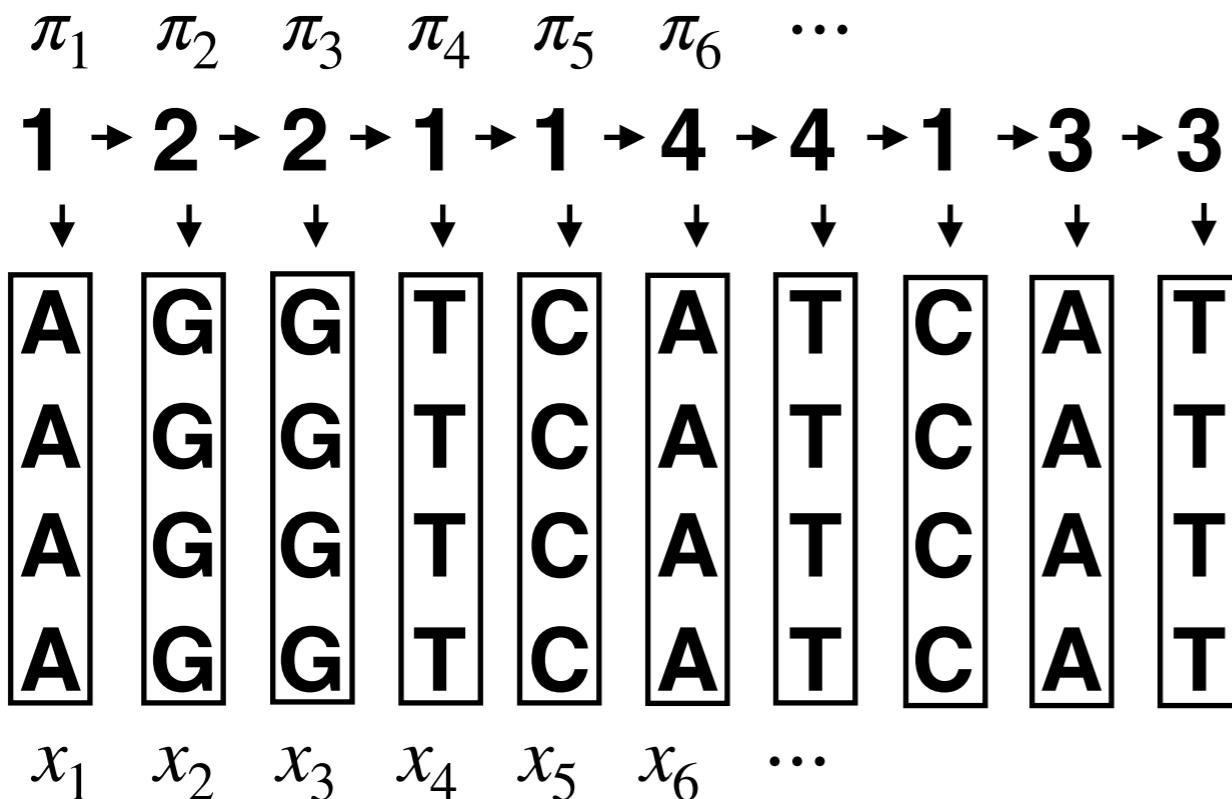


Joint probability of x and π :

$$P(x, \pi) = a_{0\pi_1} \prod_{i=1}^L e_{\pi_i}(x_i) a_{\pi_i \pi_{i+1}}$$

Joint probability

Of the observables x and hidden states π



$$P(x, \pi) = a_{0\pi_1} \prod_{i=1}^L e_{\pi_i}(x_i) a_{\pi_i \pi_{i+1}}$$

The probability of our sequence

Across *all* possible hidden state paths

How likely is it that our observables are generated by our hidden Markov model?

The probability of our sequence

Across **all possible** hidden state paths

$$P(x) = \sum_{\pi} P(x, \pi) \quad P(x, \pi) = a_{0\pi_1} \prod_{i=1}^L e_{\pi_i}(x_i) a_{\pi_i \pi_{i+1}}$$

The probability of our sequence

Across all possible hidden state paths

$$P(x) = \sum_{\pi} P(x, \pi) \quad P(x, \pi) = a_{0\pi_1} \prod_{i=1}^L e_{\pi_i}(x_i) a_{\pi_i \pi_{i+1}}$$

Problem: the number of paths grows exponentially with sequence length!

Luckily, we can use dynamic programming like in the Viterbi algorithm.

The probability of our sequence

Across all possible hidden state paths

$f_k(i)$: The joint **probability** of **all state paths** for the first i states and i observables, assuming that the last state is k .

The probability of our sequence

Across all possible hidden state paths

$f_k(i)$

- The joint **probability** of **all state paths** for the first i states and i observables, assuming that the last state is k .

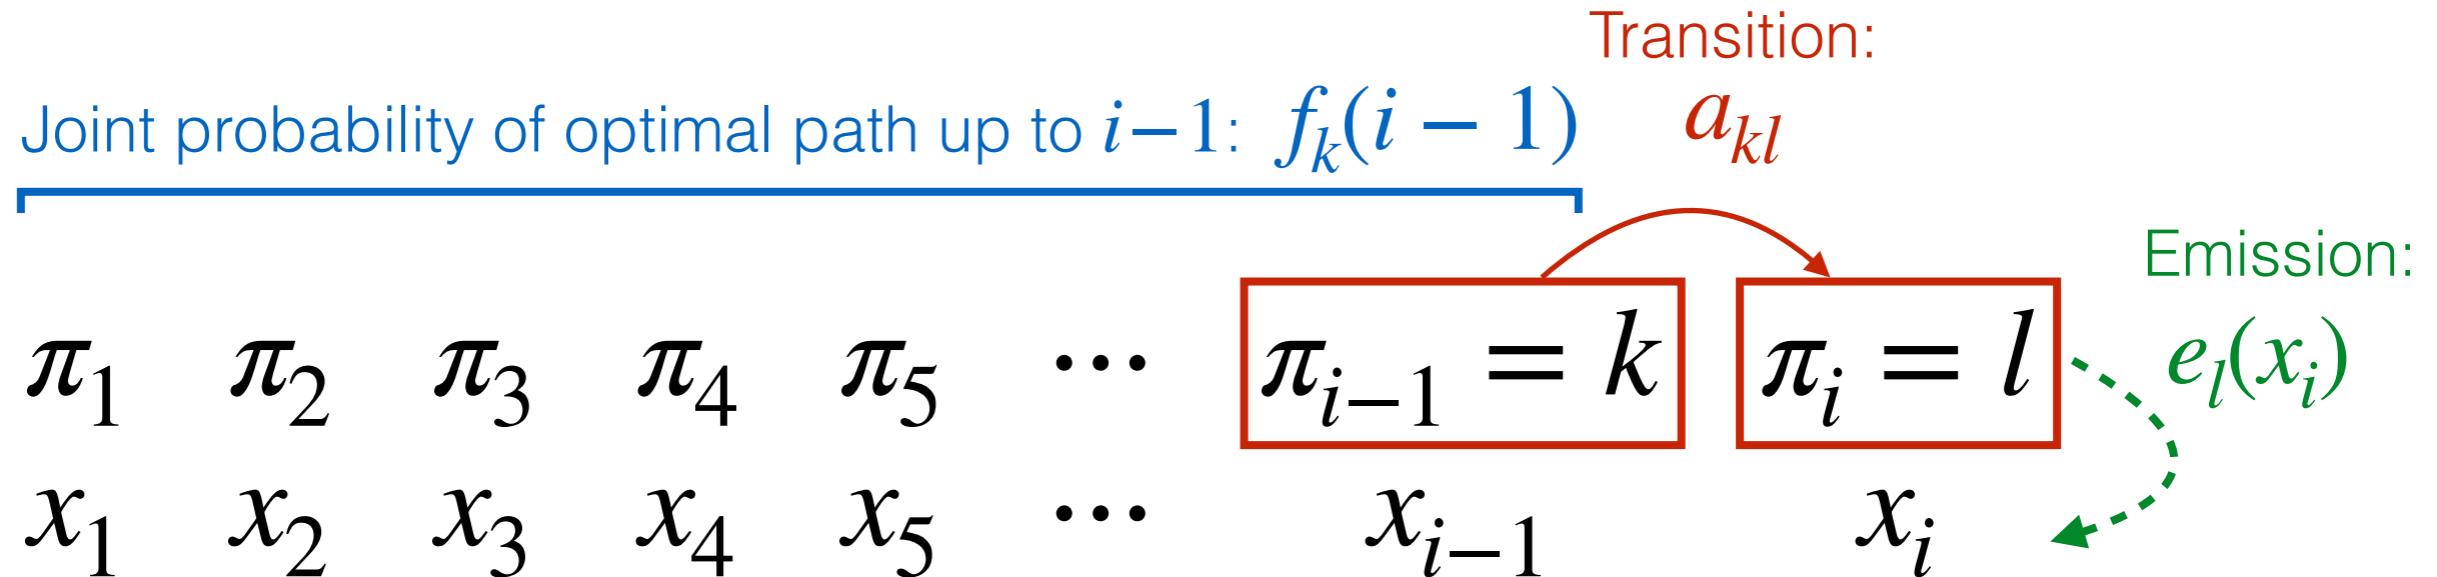
Joint probability of all paths up to i : $f_l(i)$

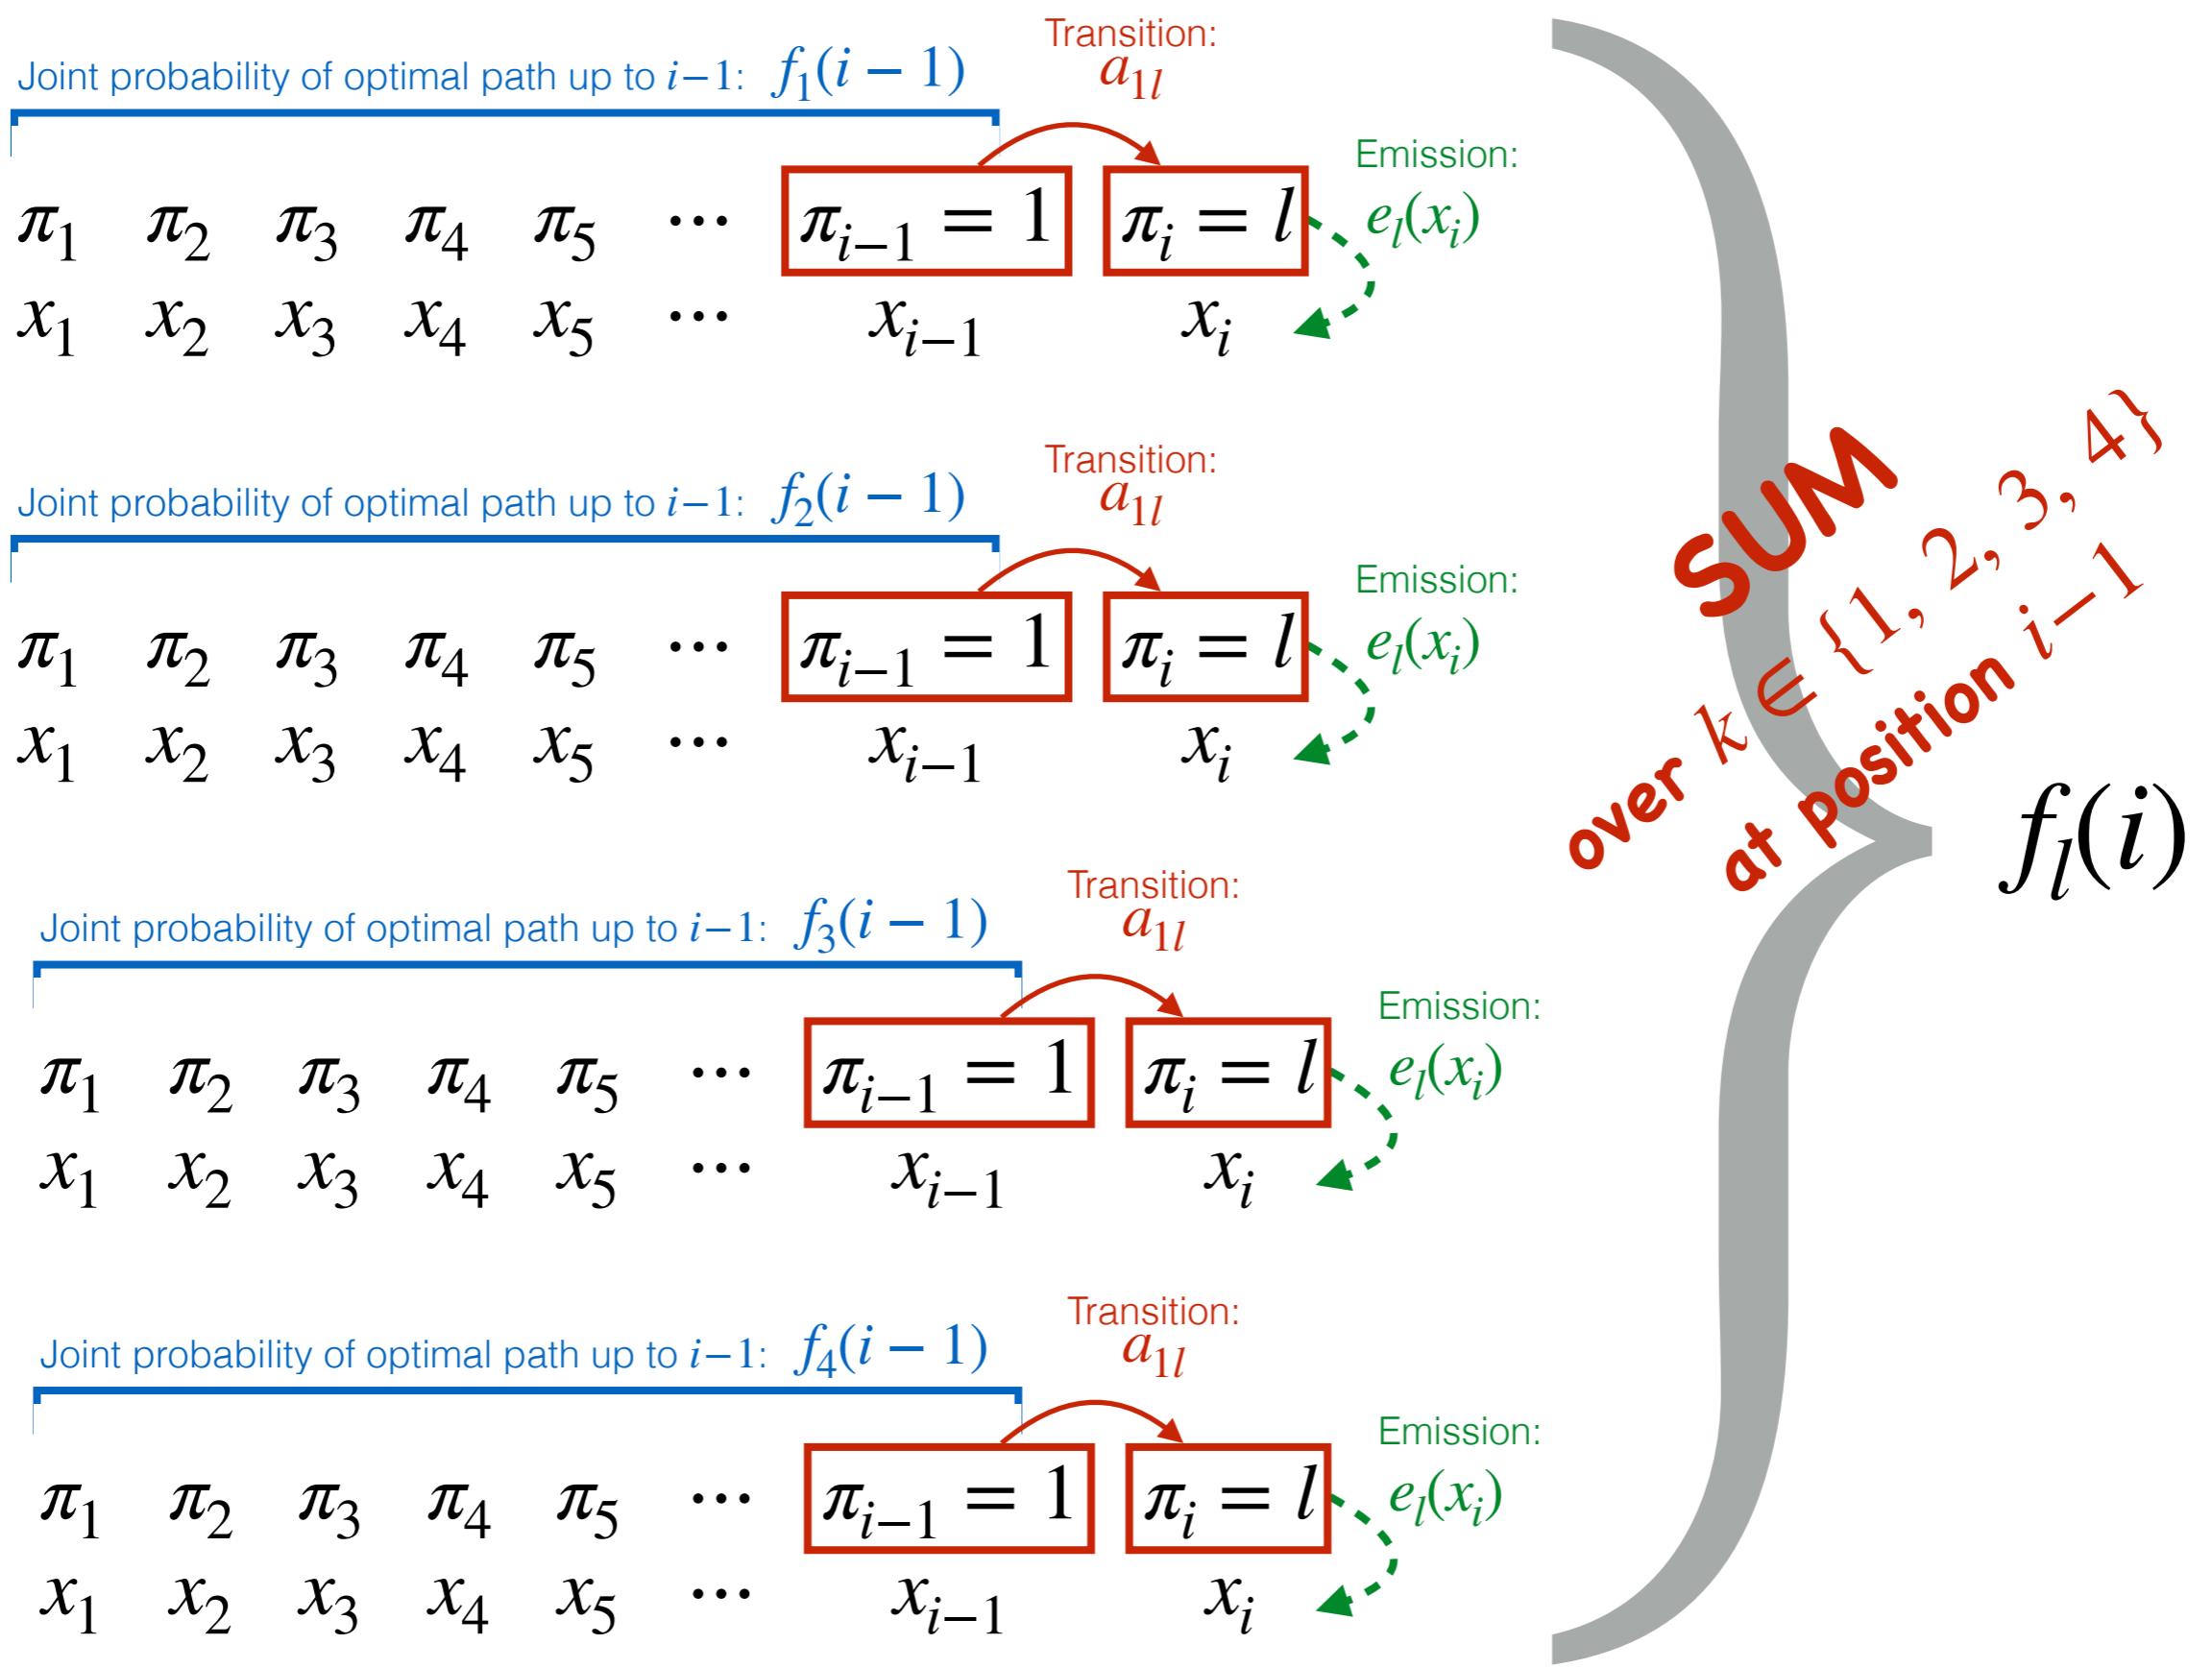
π_1	π_2	π_3	π_4	π_5	\dots	$\pi_i = l$
x_1	x_2	x_3	x_4	x_5	\dots	x_i

Most probable state path

Because of the Markov property we can compute its probability recursively:

$$f_l(i) = \sum_k f_k(i-1) a_{kl} e_l(x_i)$$





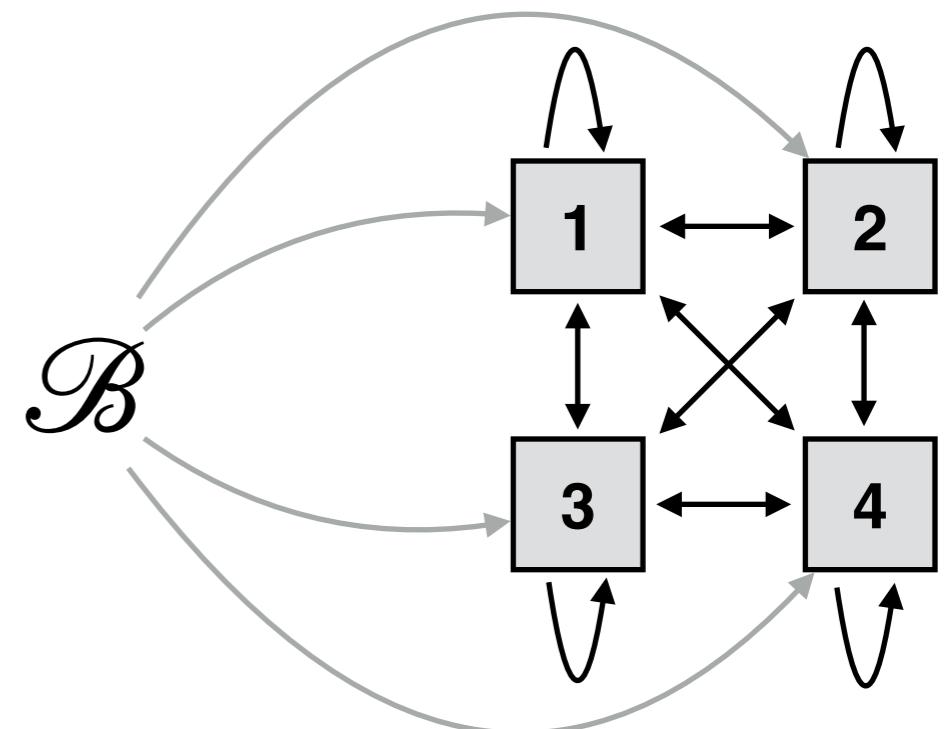
Recursion

Recursive case:

$$f_l(i) = e_l(x_i) \sum_k (f_k(i - 1) a_{kl})$$

Base case: when no sequence is used the state is *always* the null state:

$$\nu_0(\mathcal{B}) = 1$$



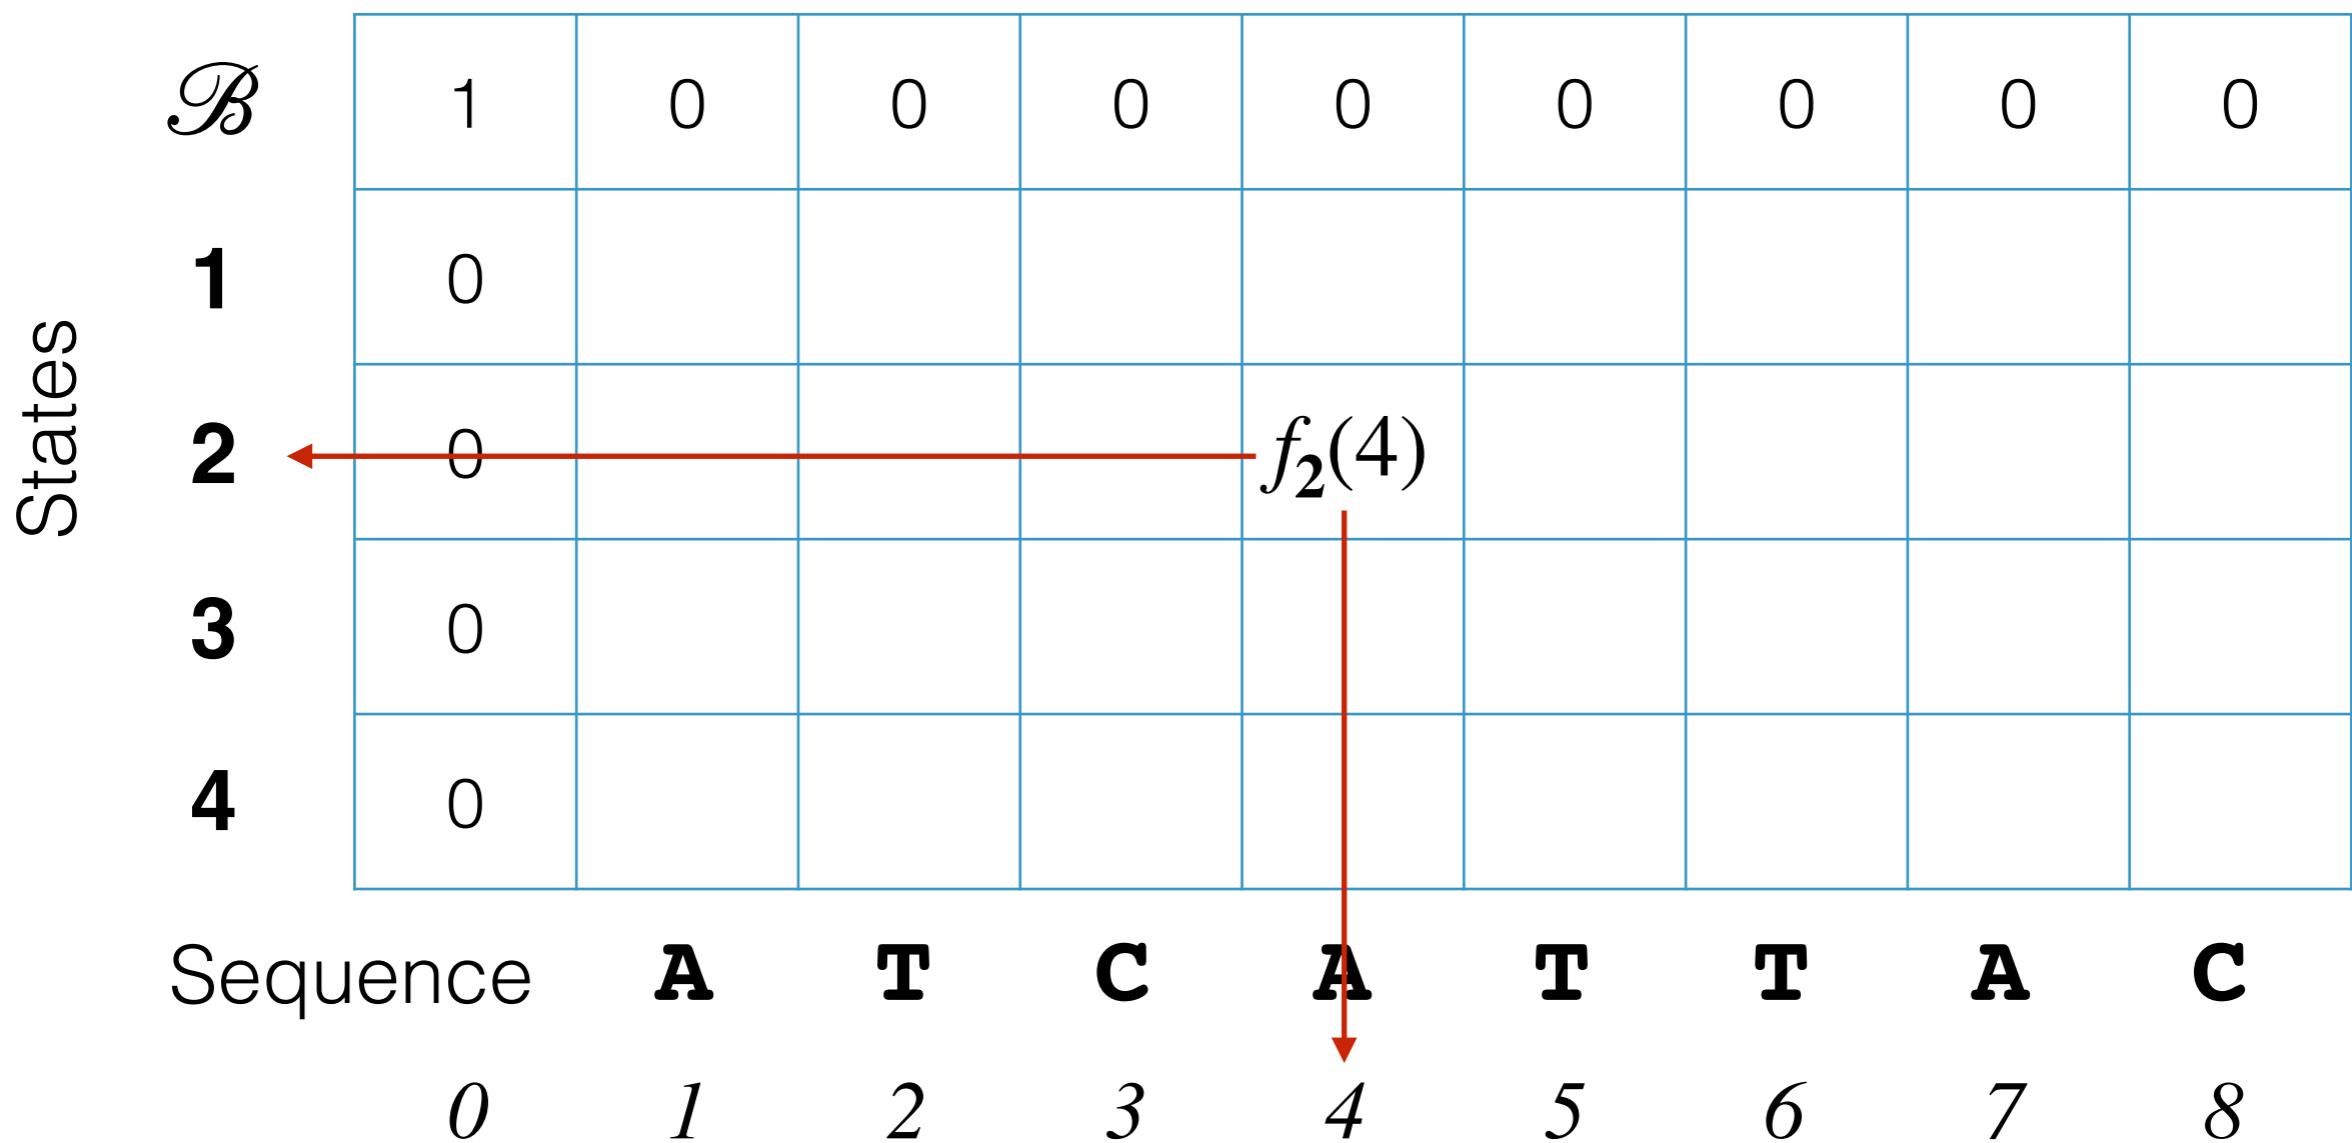
Recursion

$$f_l(i) = \sum_k f_k(i-1) a_{kl} e_l(x_i)$$

How do we avoid computing the same probability many many times?

Remember pairwise alignment?

Dynamic programming table for states and observables



Dynamic programming table for states and observables

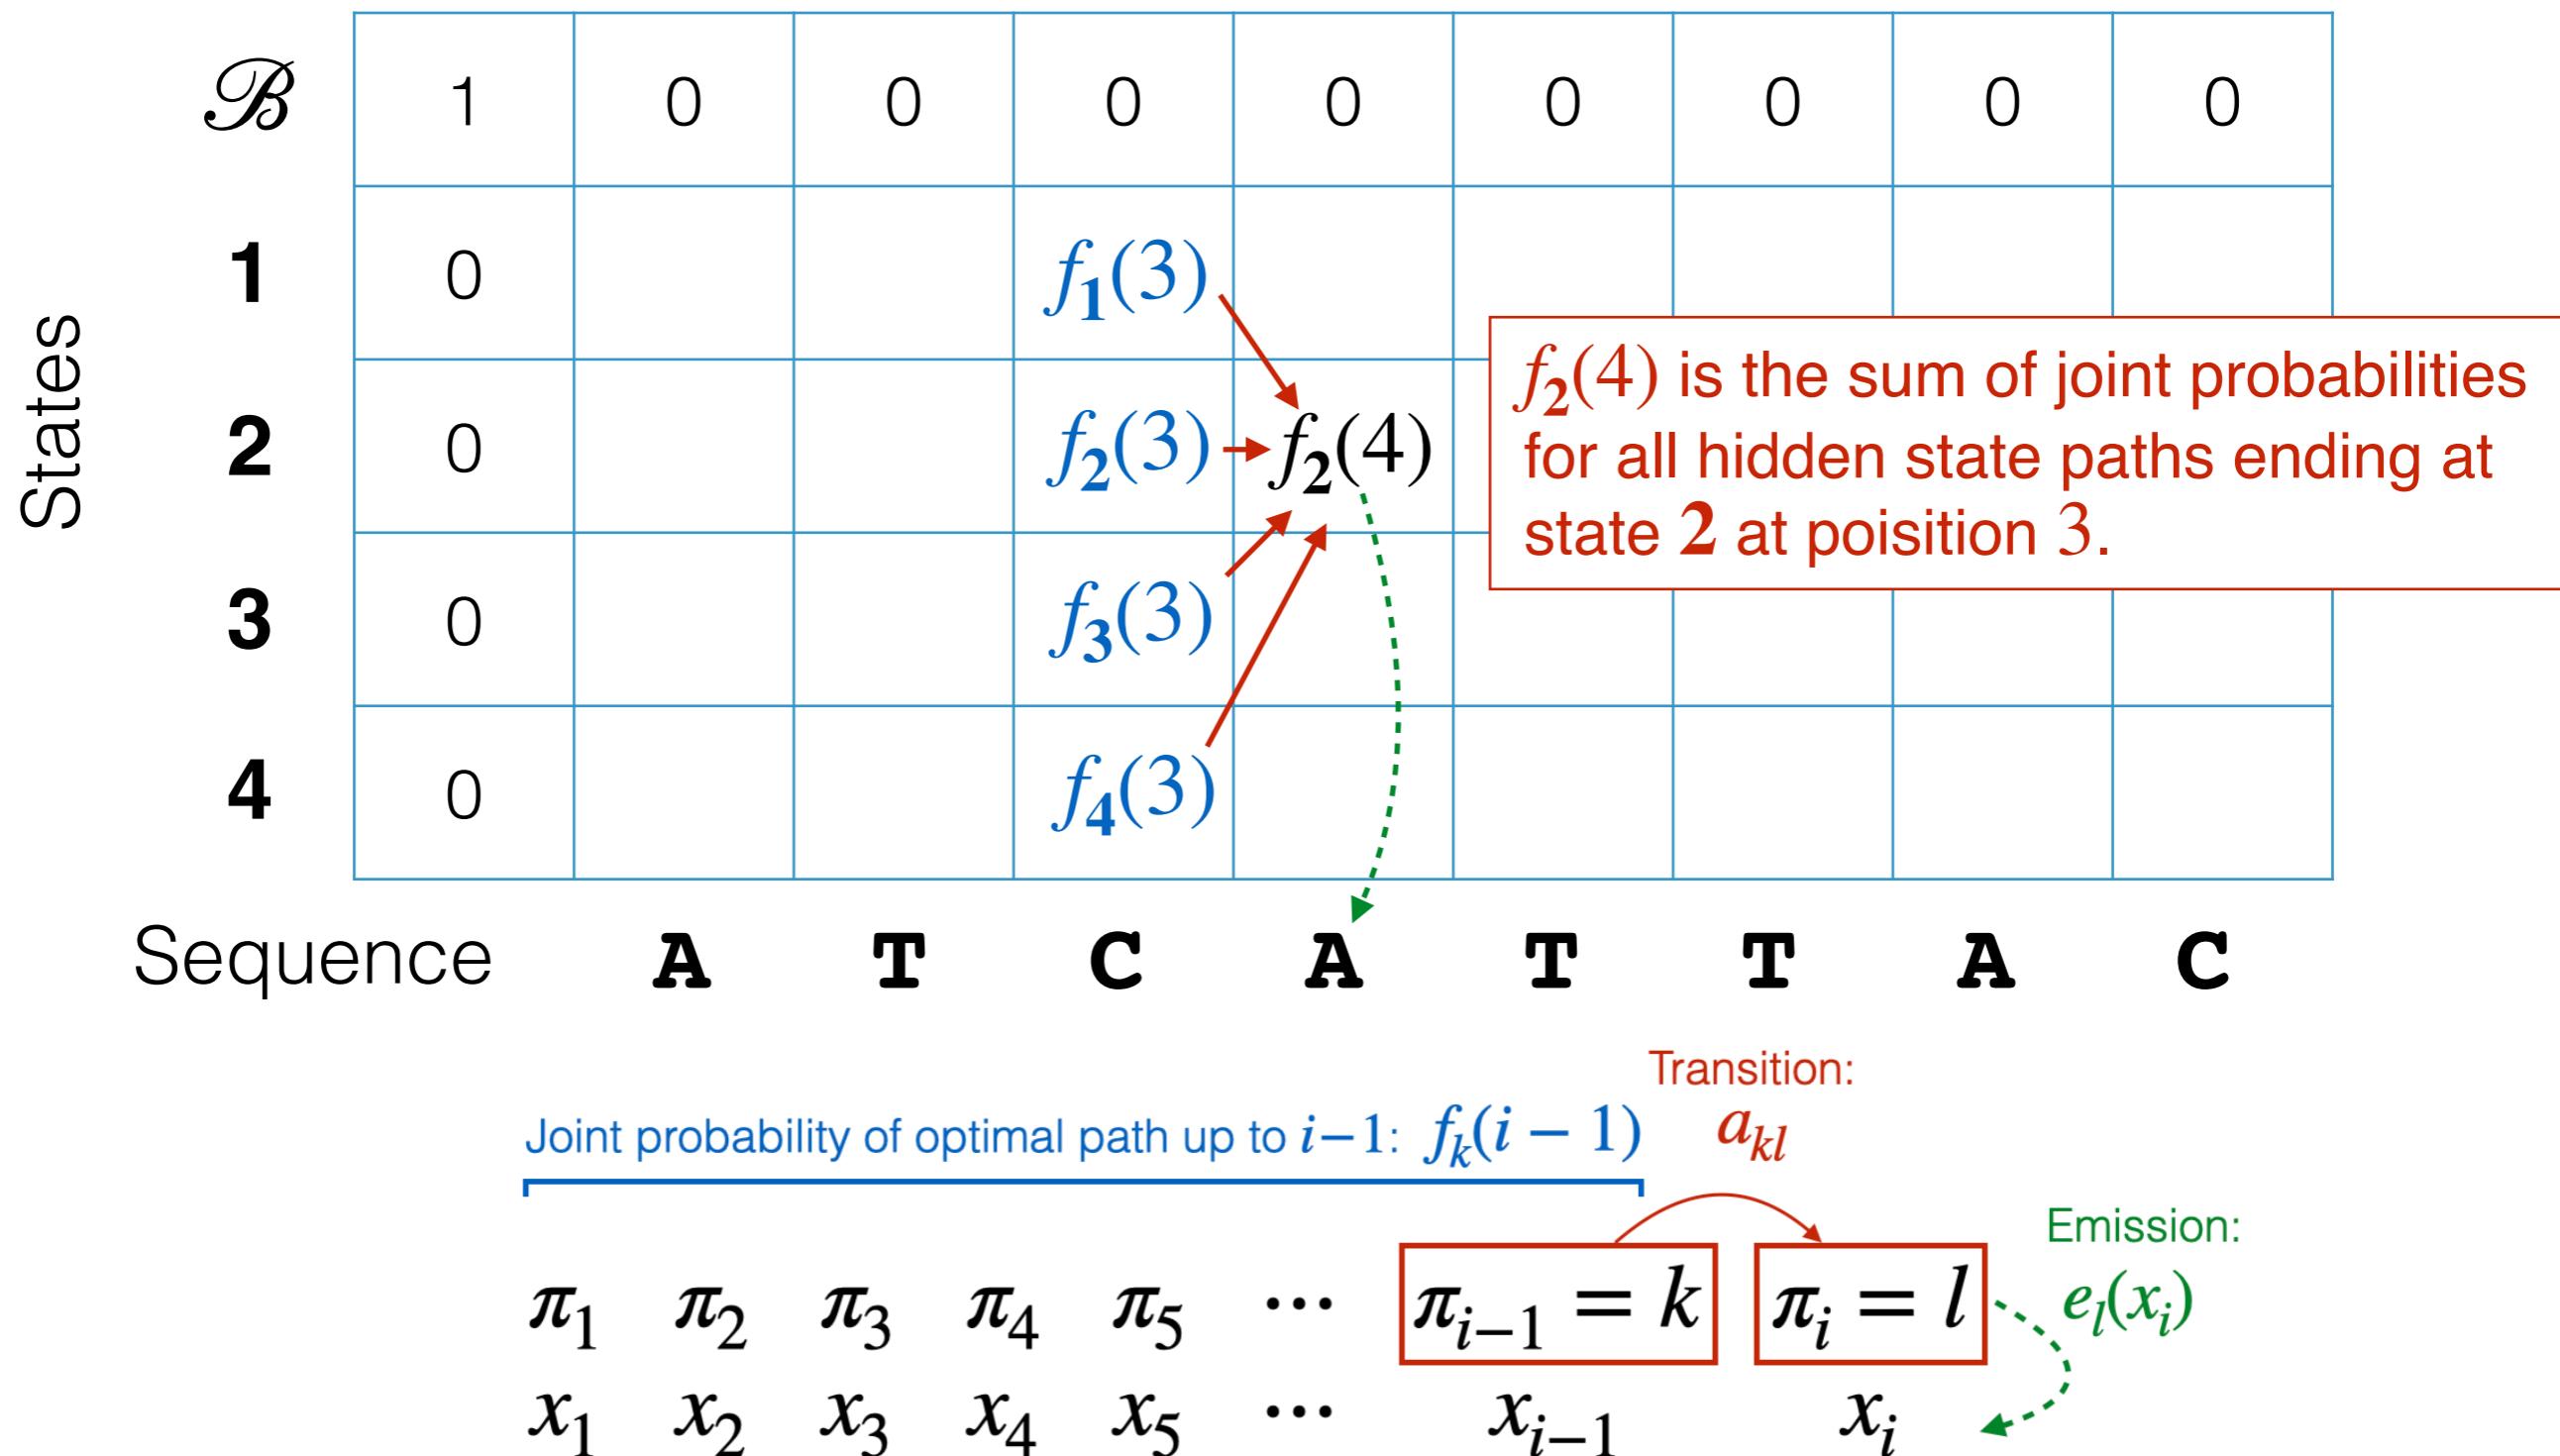
\mathcal{B}	1	0	0	0	0	0	0	0	0
1	0								$f_1(8)$
2	0								$f_2(8)$
3	0								$f_3(8)$
4	0								$f_4(8)$

States

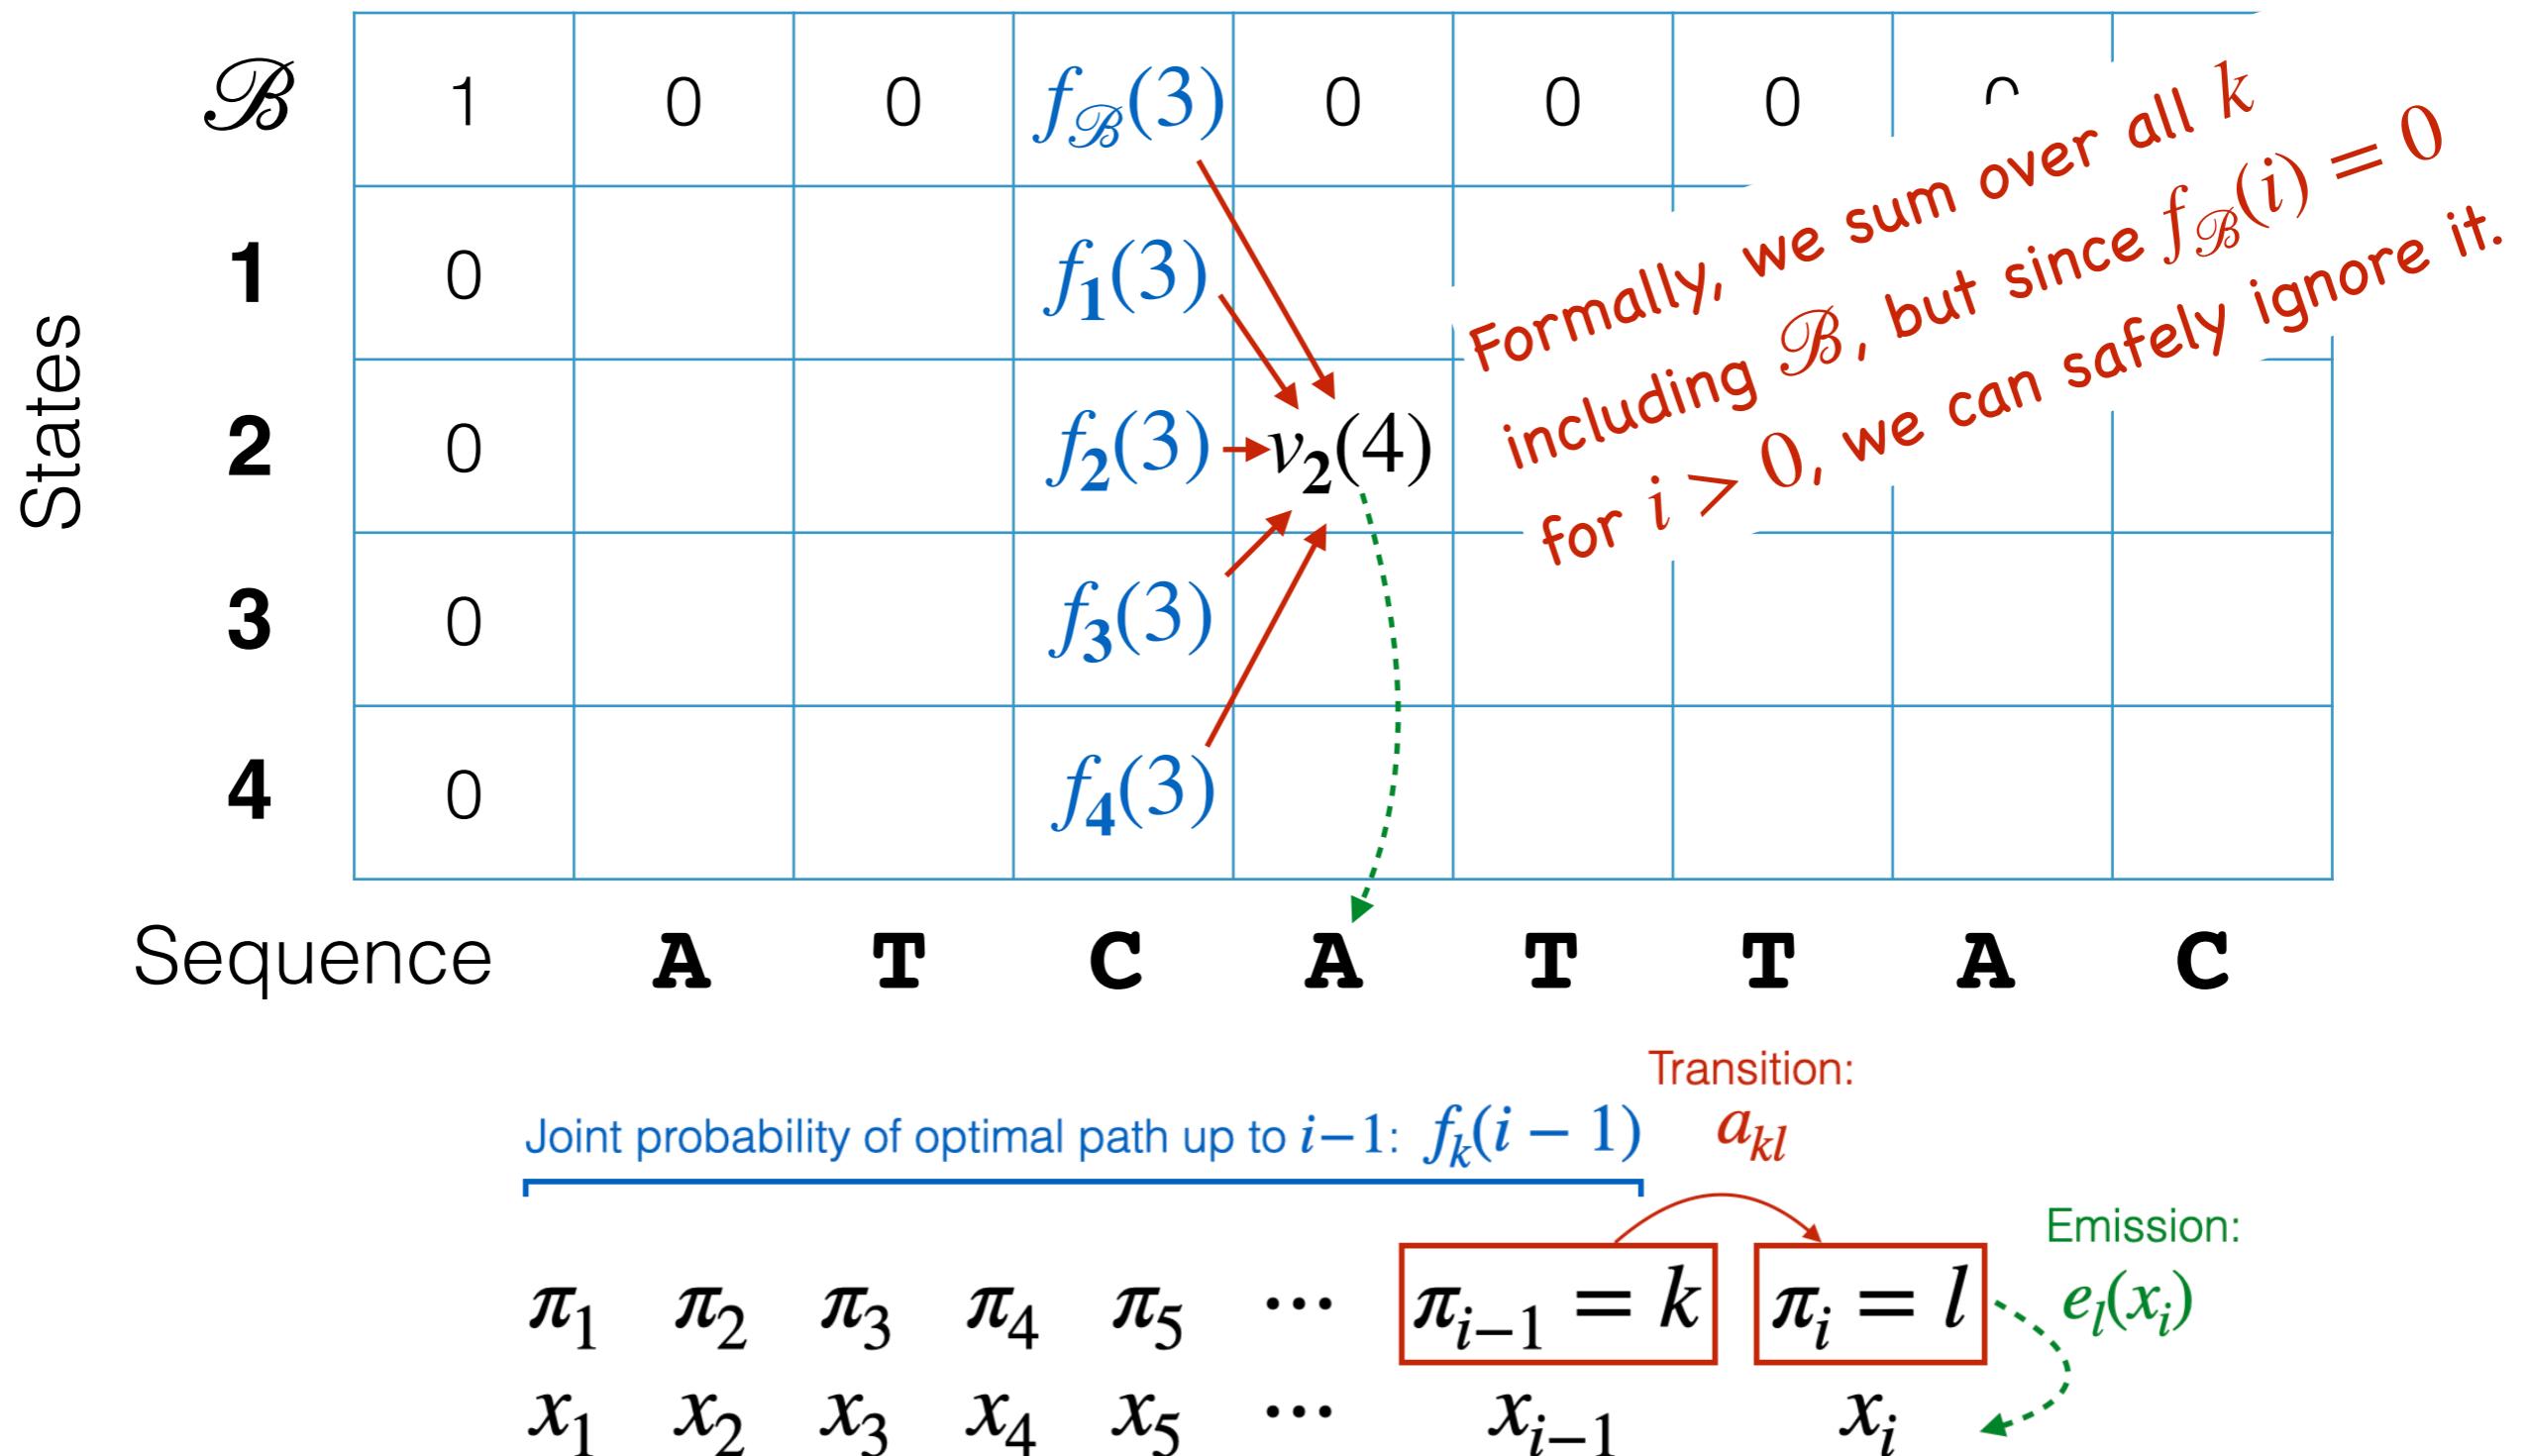
Sequence **A** **T** **C** **A** **T** **T** **A** **C**

Sum

Dynamic programming table for states and observables



Dynamic programming table for states and observables



Dynamic programming table for states and observables

\mathcal{B}	1	0	0	0	0	0	0	0	0
1	0								$f_1(8)$
2	0								$f_2(8)$
3	0								$f_3(8)$
4	0								$f_4(8)$

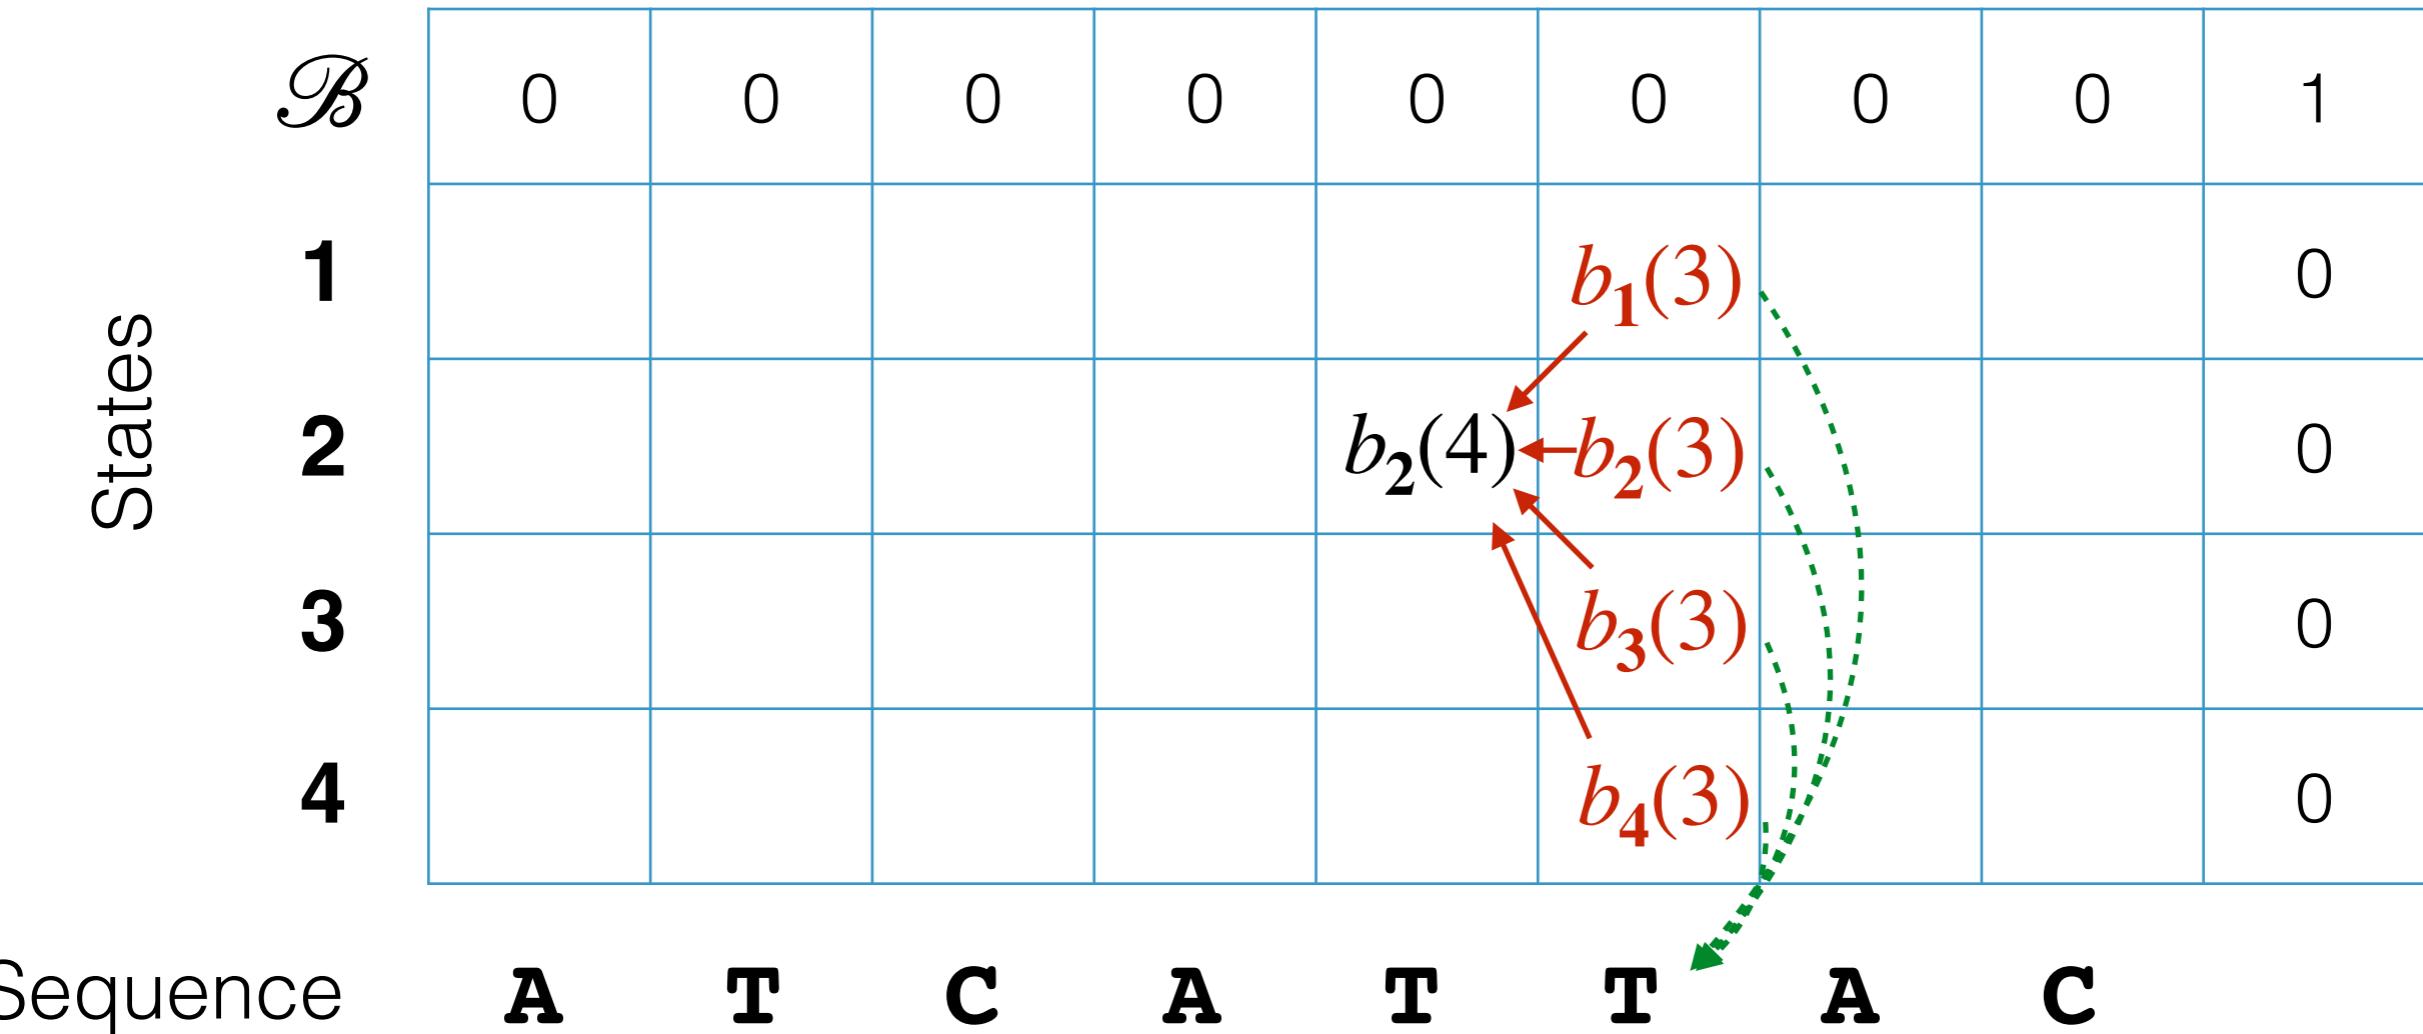
States

Sequence **A** **T** **C** **A** **T** **T** **A** **C**

Fill the table column by column, from left to right

Sum

Backward algorithm

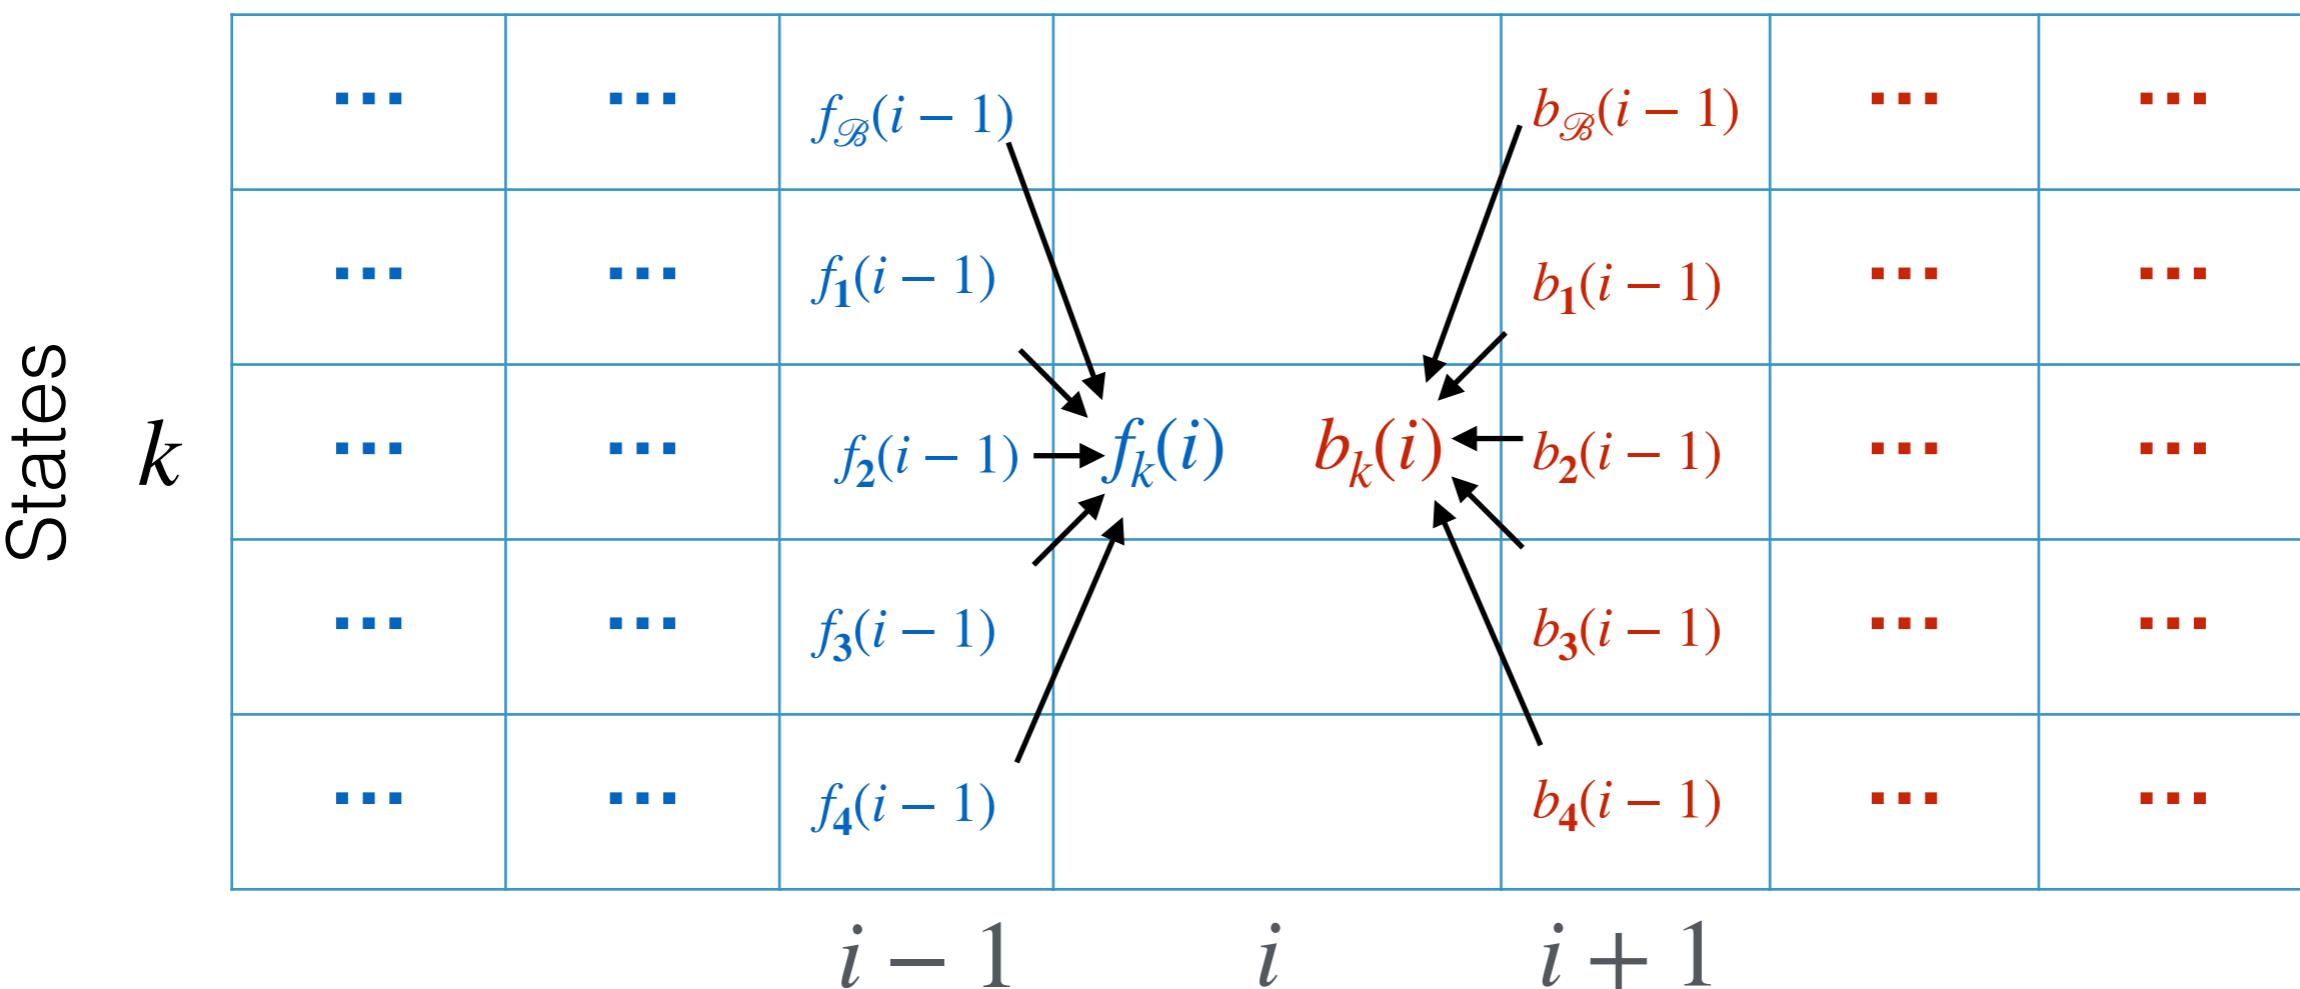


$$b_l(i) = \sum_k a_{kl} e_l(x_{i+1}) b_k(i+1)$$

Posterior probability

- of a state at a position in the sequence

Forward:



Backward:

The proportion of the total probability contributed by paths with state k at position i :

$$P(\pi_i = k | x) = \frac{P(\pi_i = k)}{P(x)} = \frac{f_k(i) b_k(i)}{P(x)}$$

Posterior decoding

Hidden states with max posterior probability

Posterior probability that the state at position i is k :

$$P(\pi_i = k | x) = \frac{f_k(i) b_k(i)}{P(x)}$$

Important: Contributions from many different paths. So the sequence of states with the highest posterior probability will not correspond to a valid path through the state model.

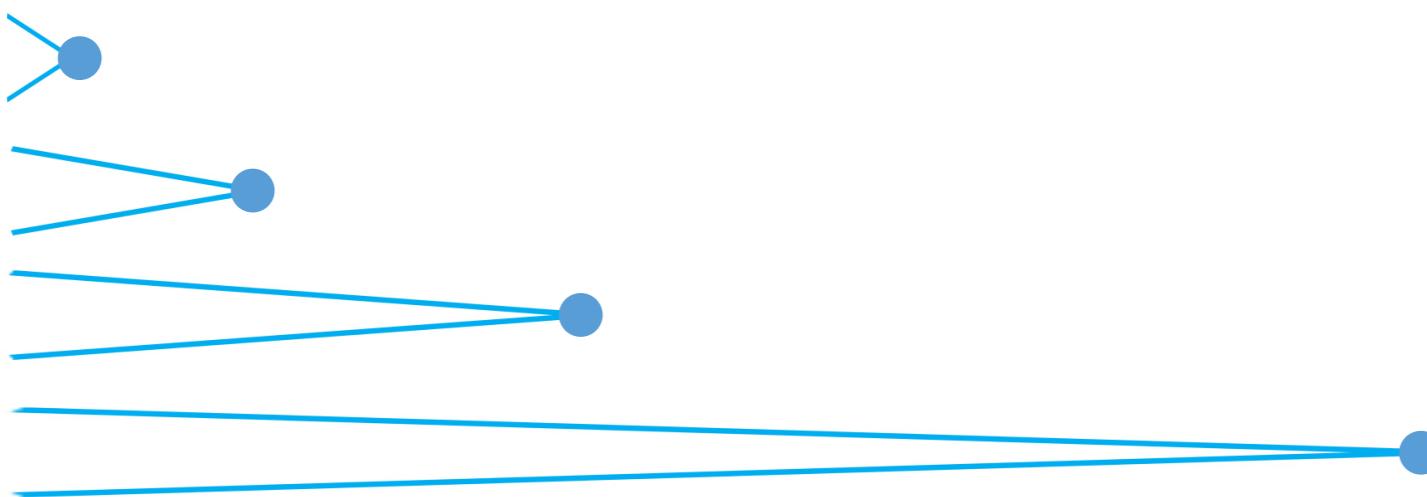
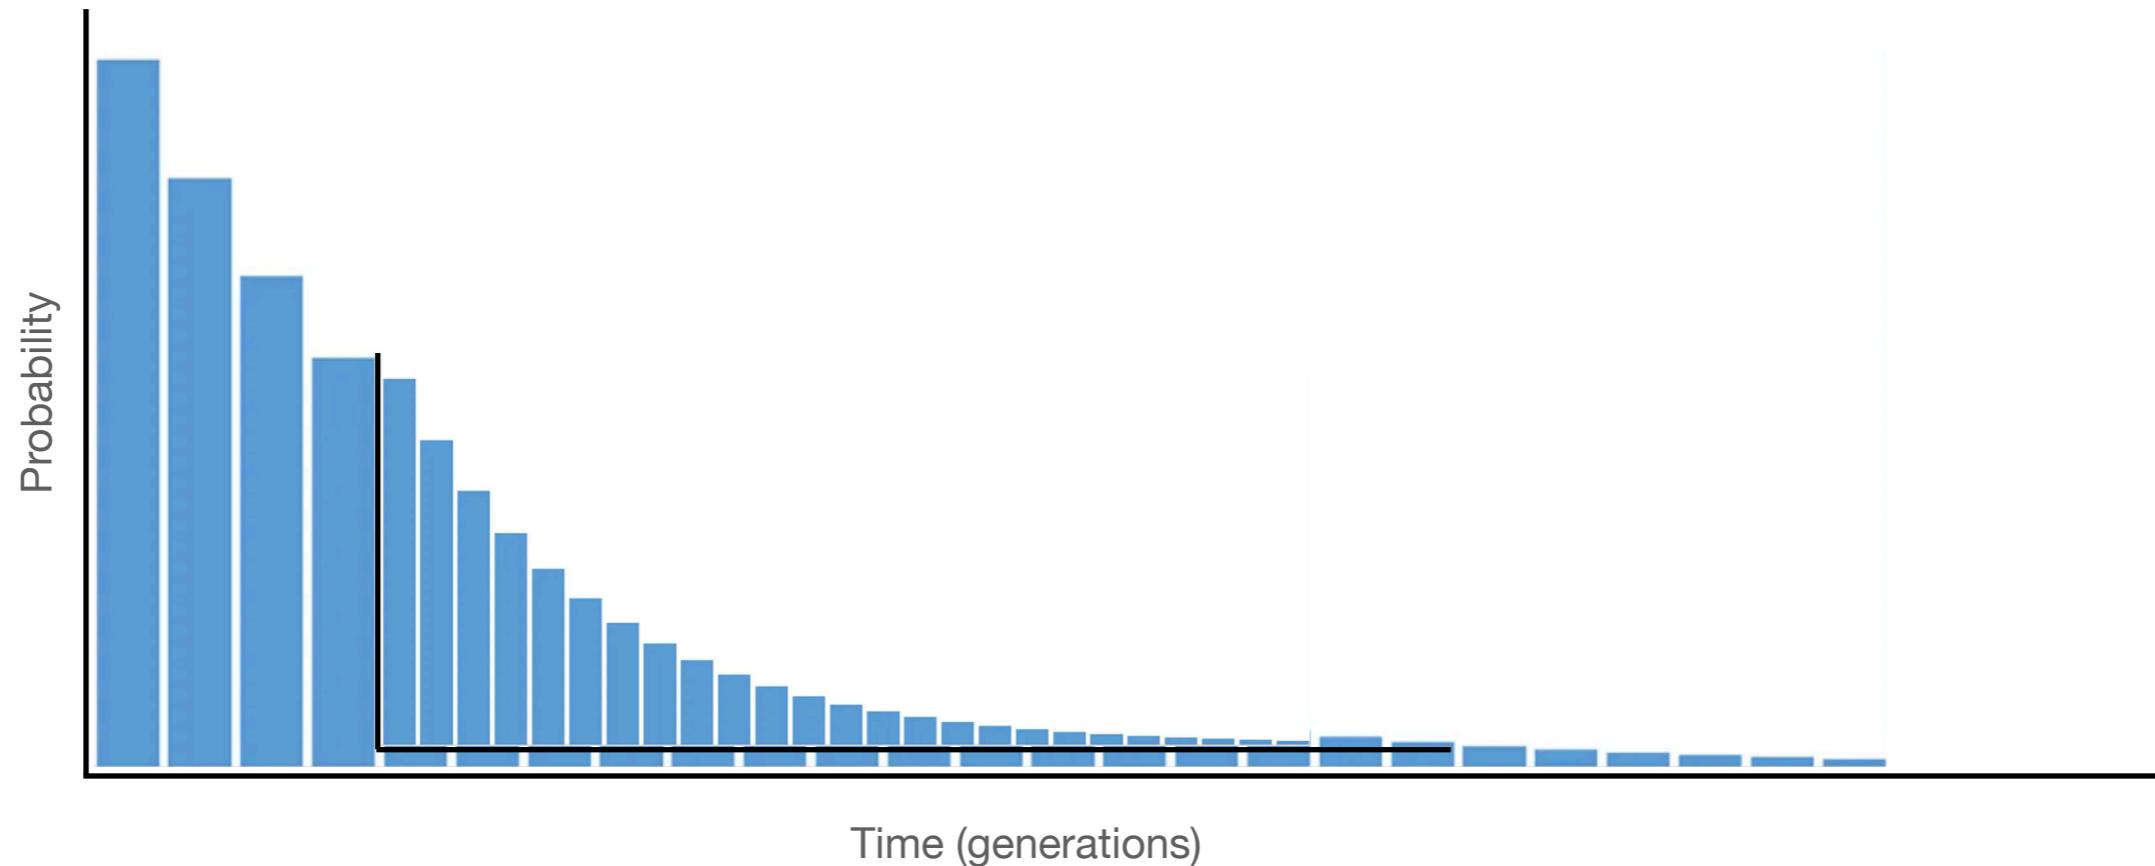
Demography with PSMC

Week 3

Kasper Munch

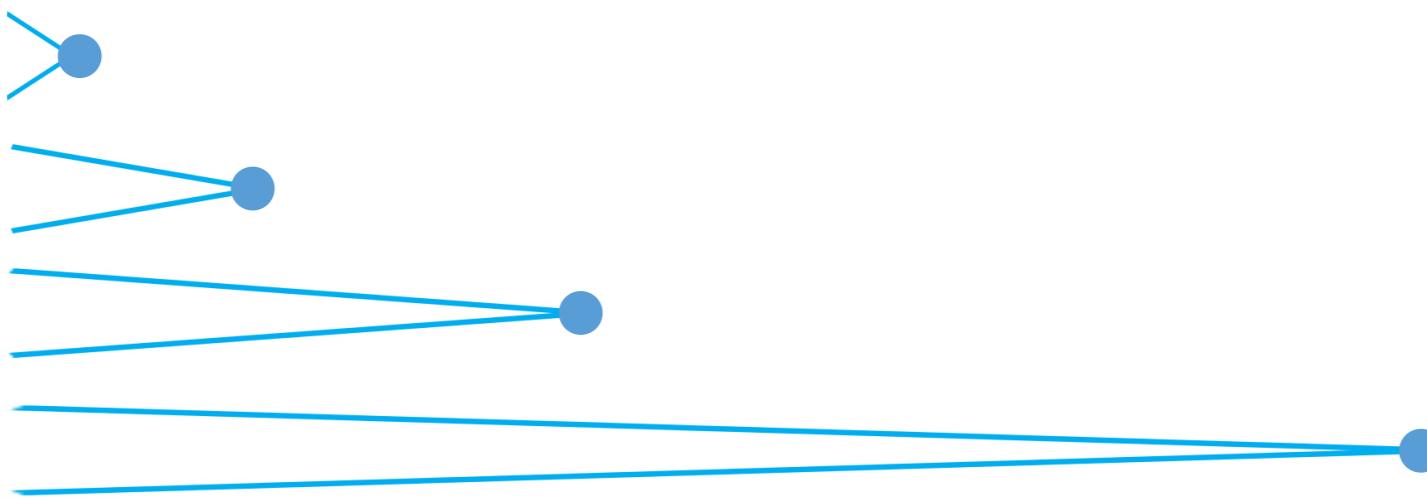
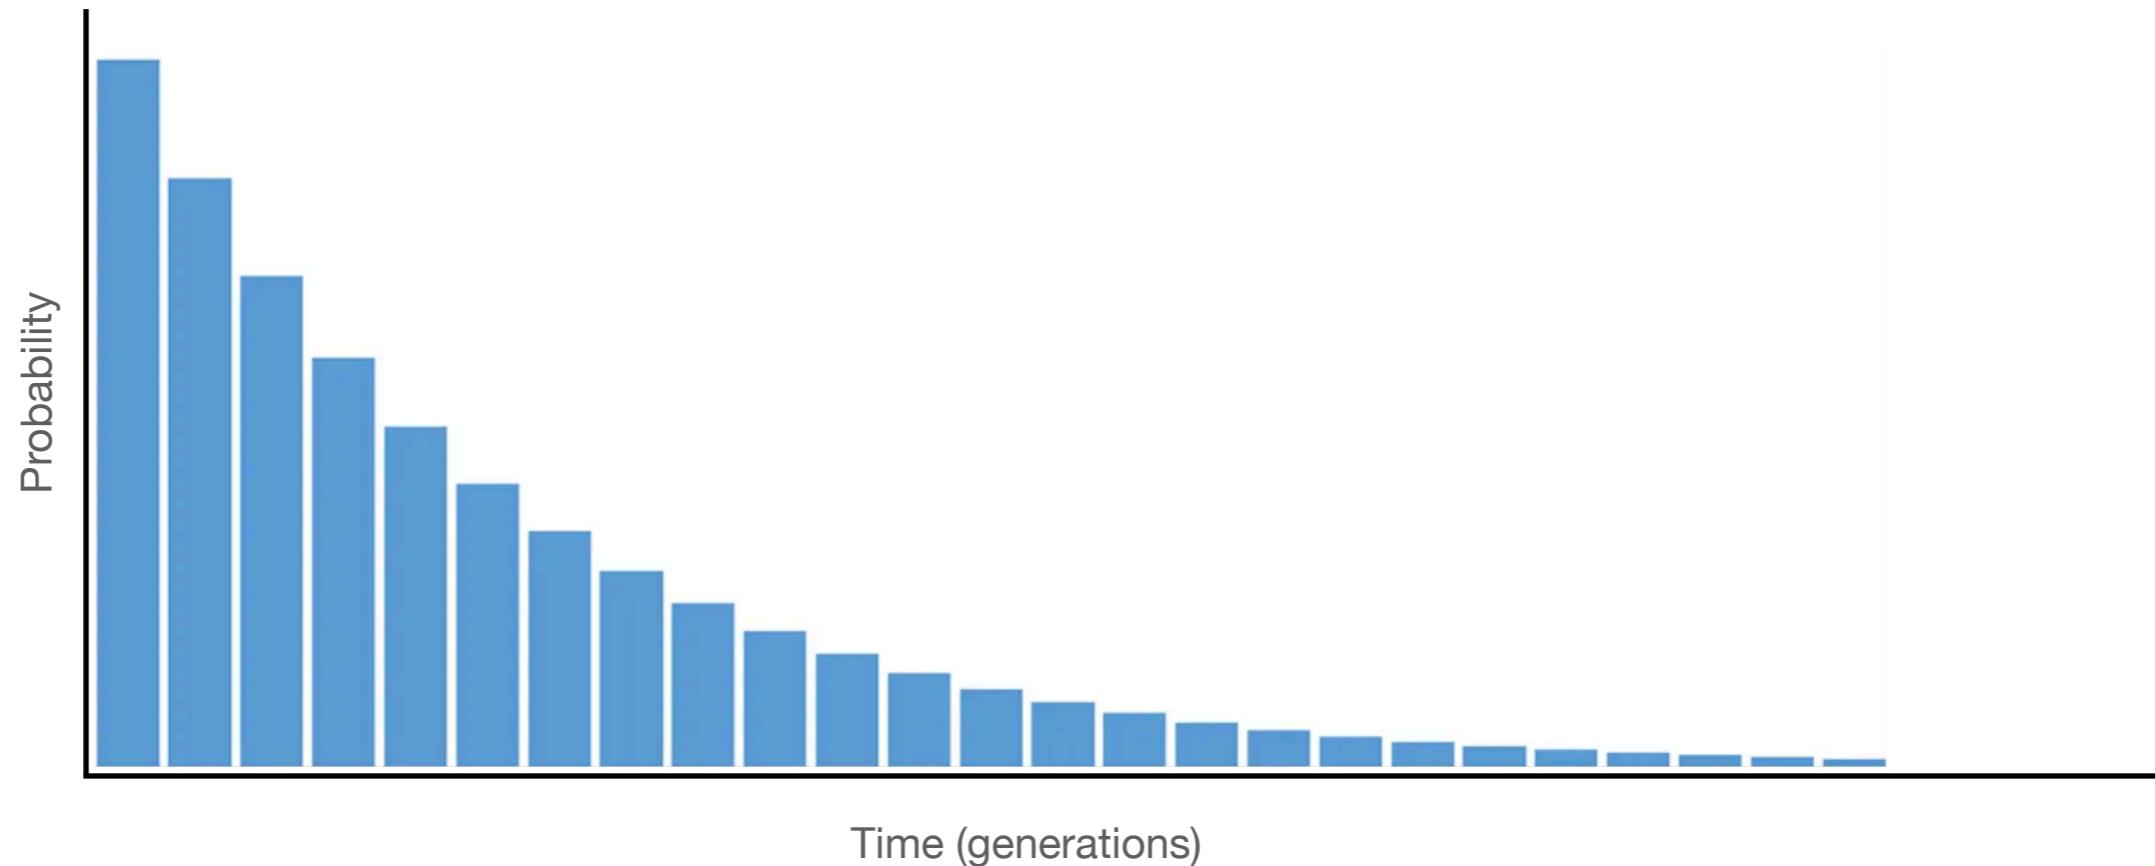
Distribution of TMRCA

For two sequences



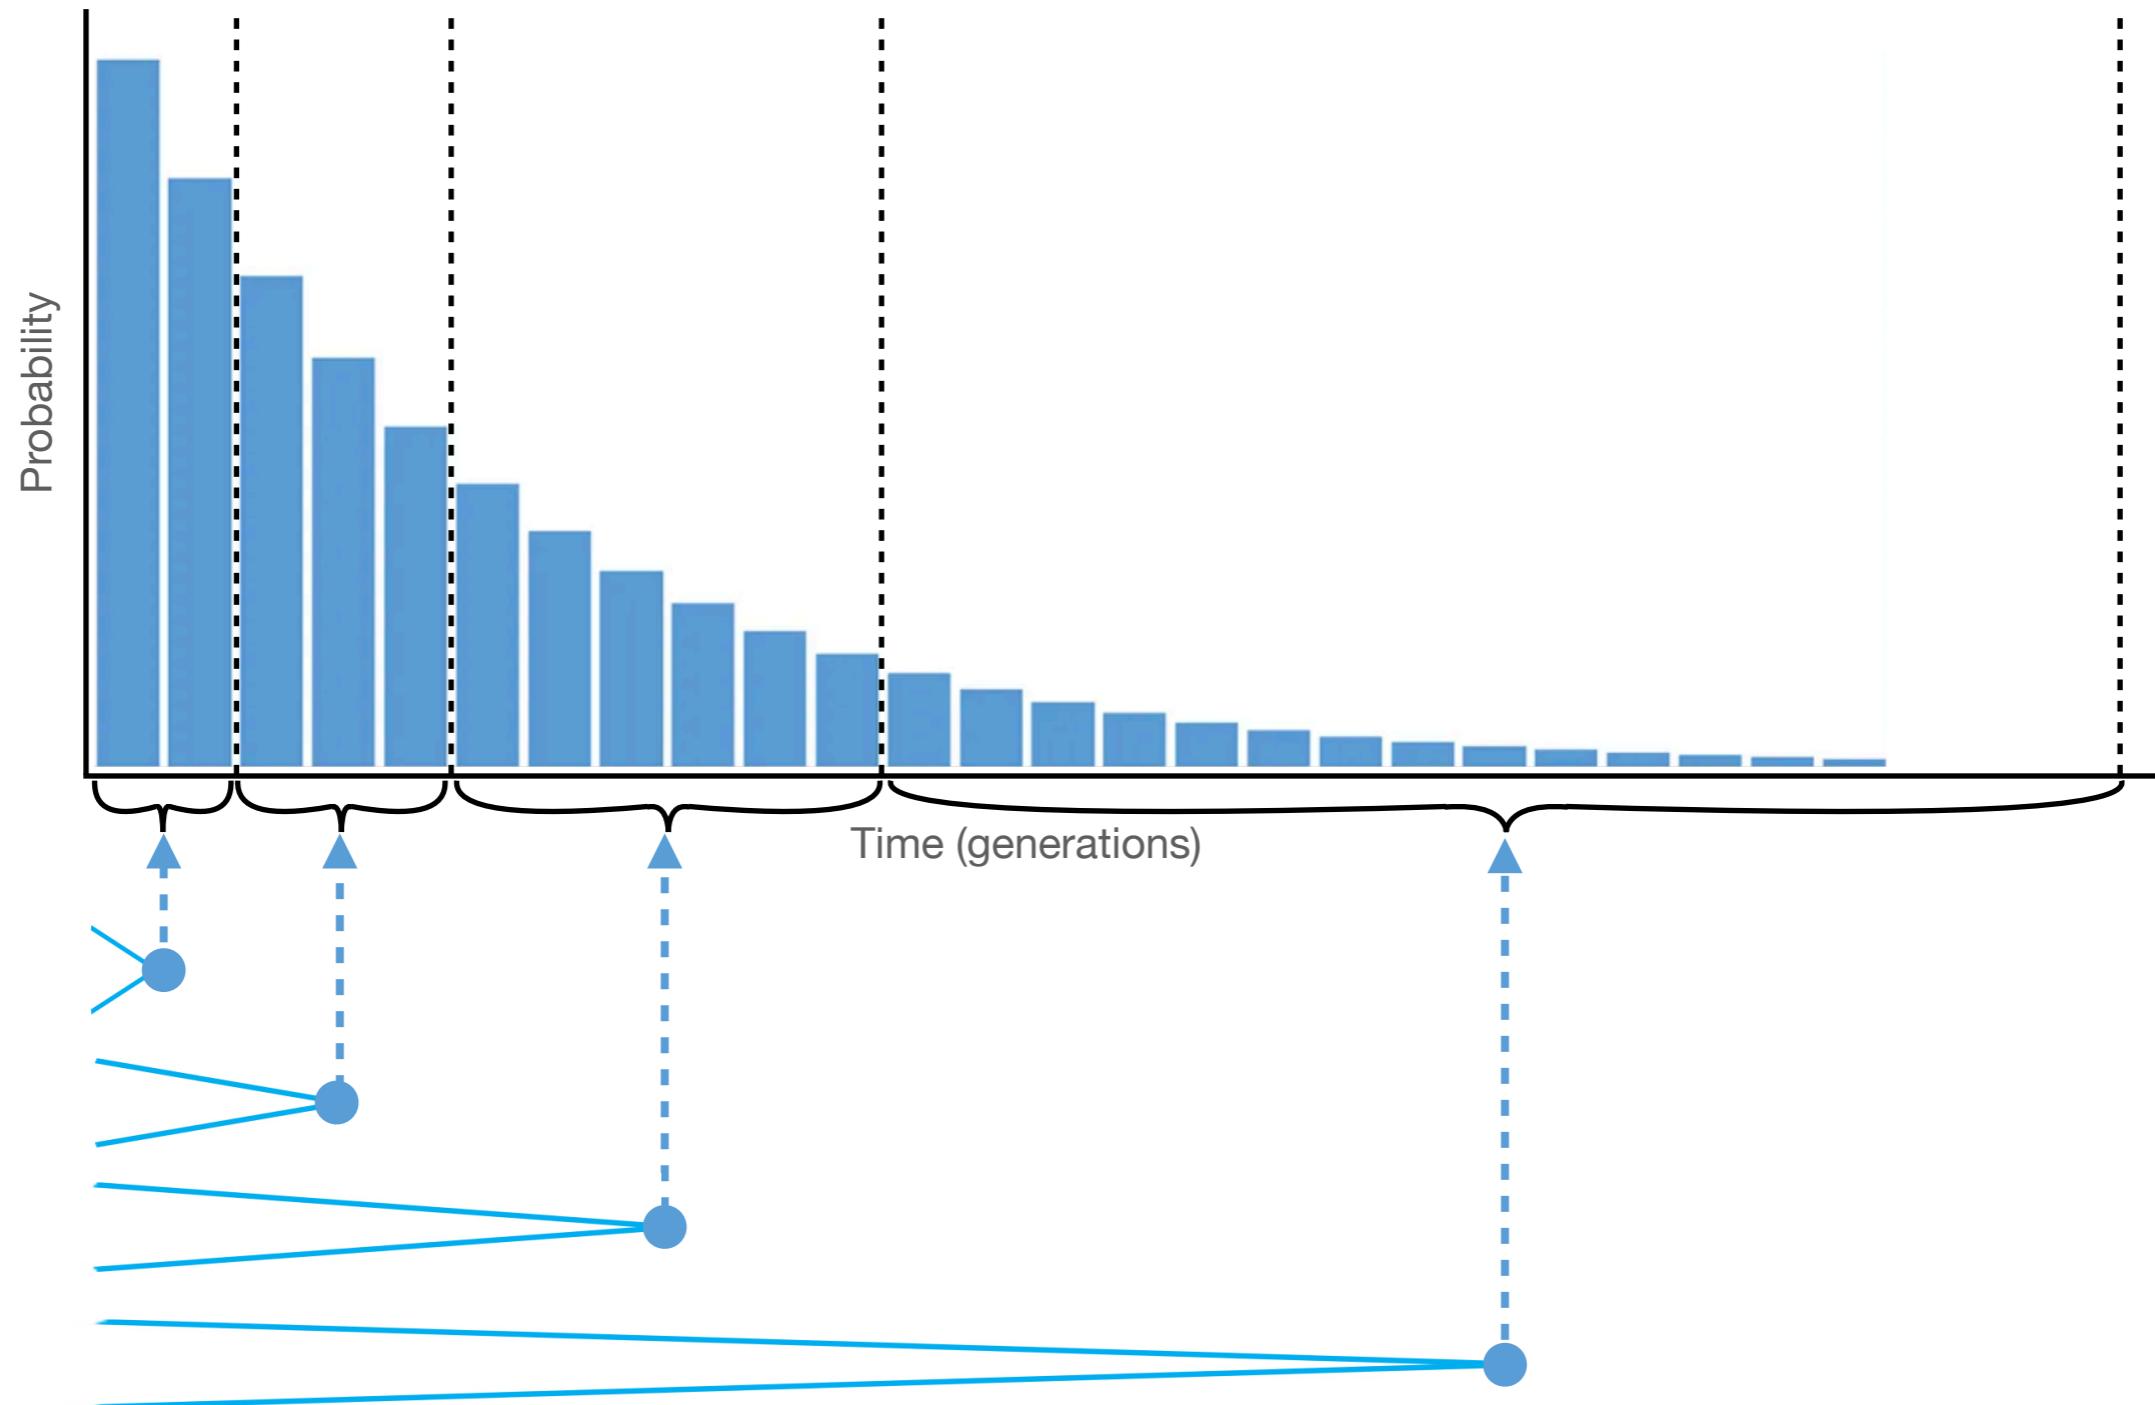
Distribution of TMRCA

For two sequences



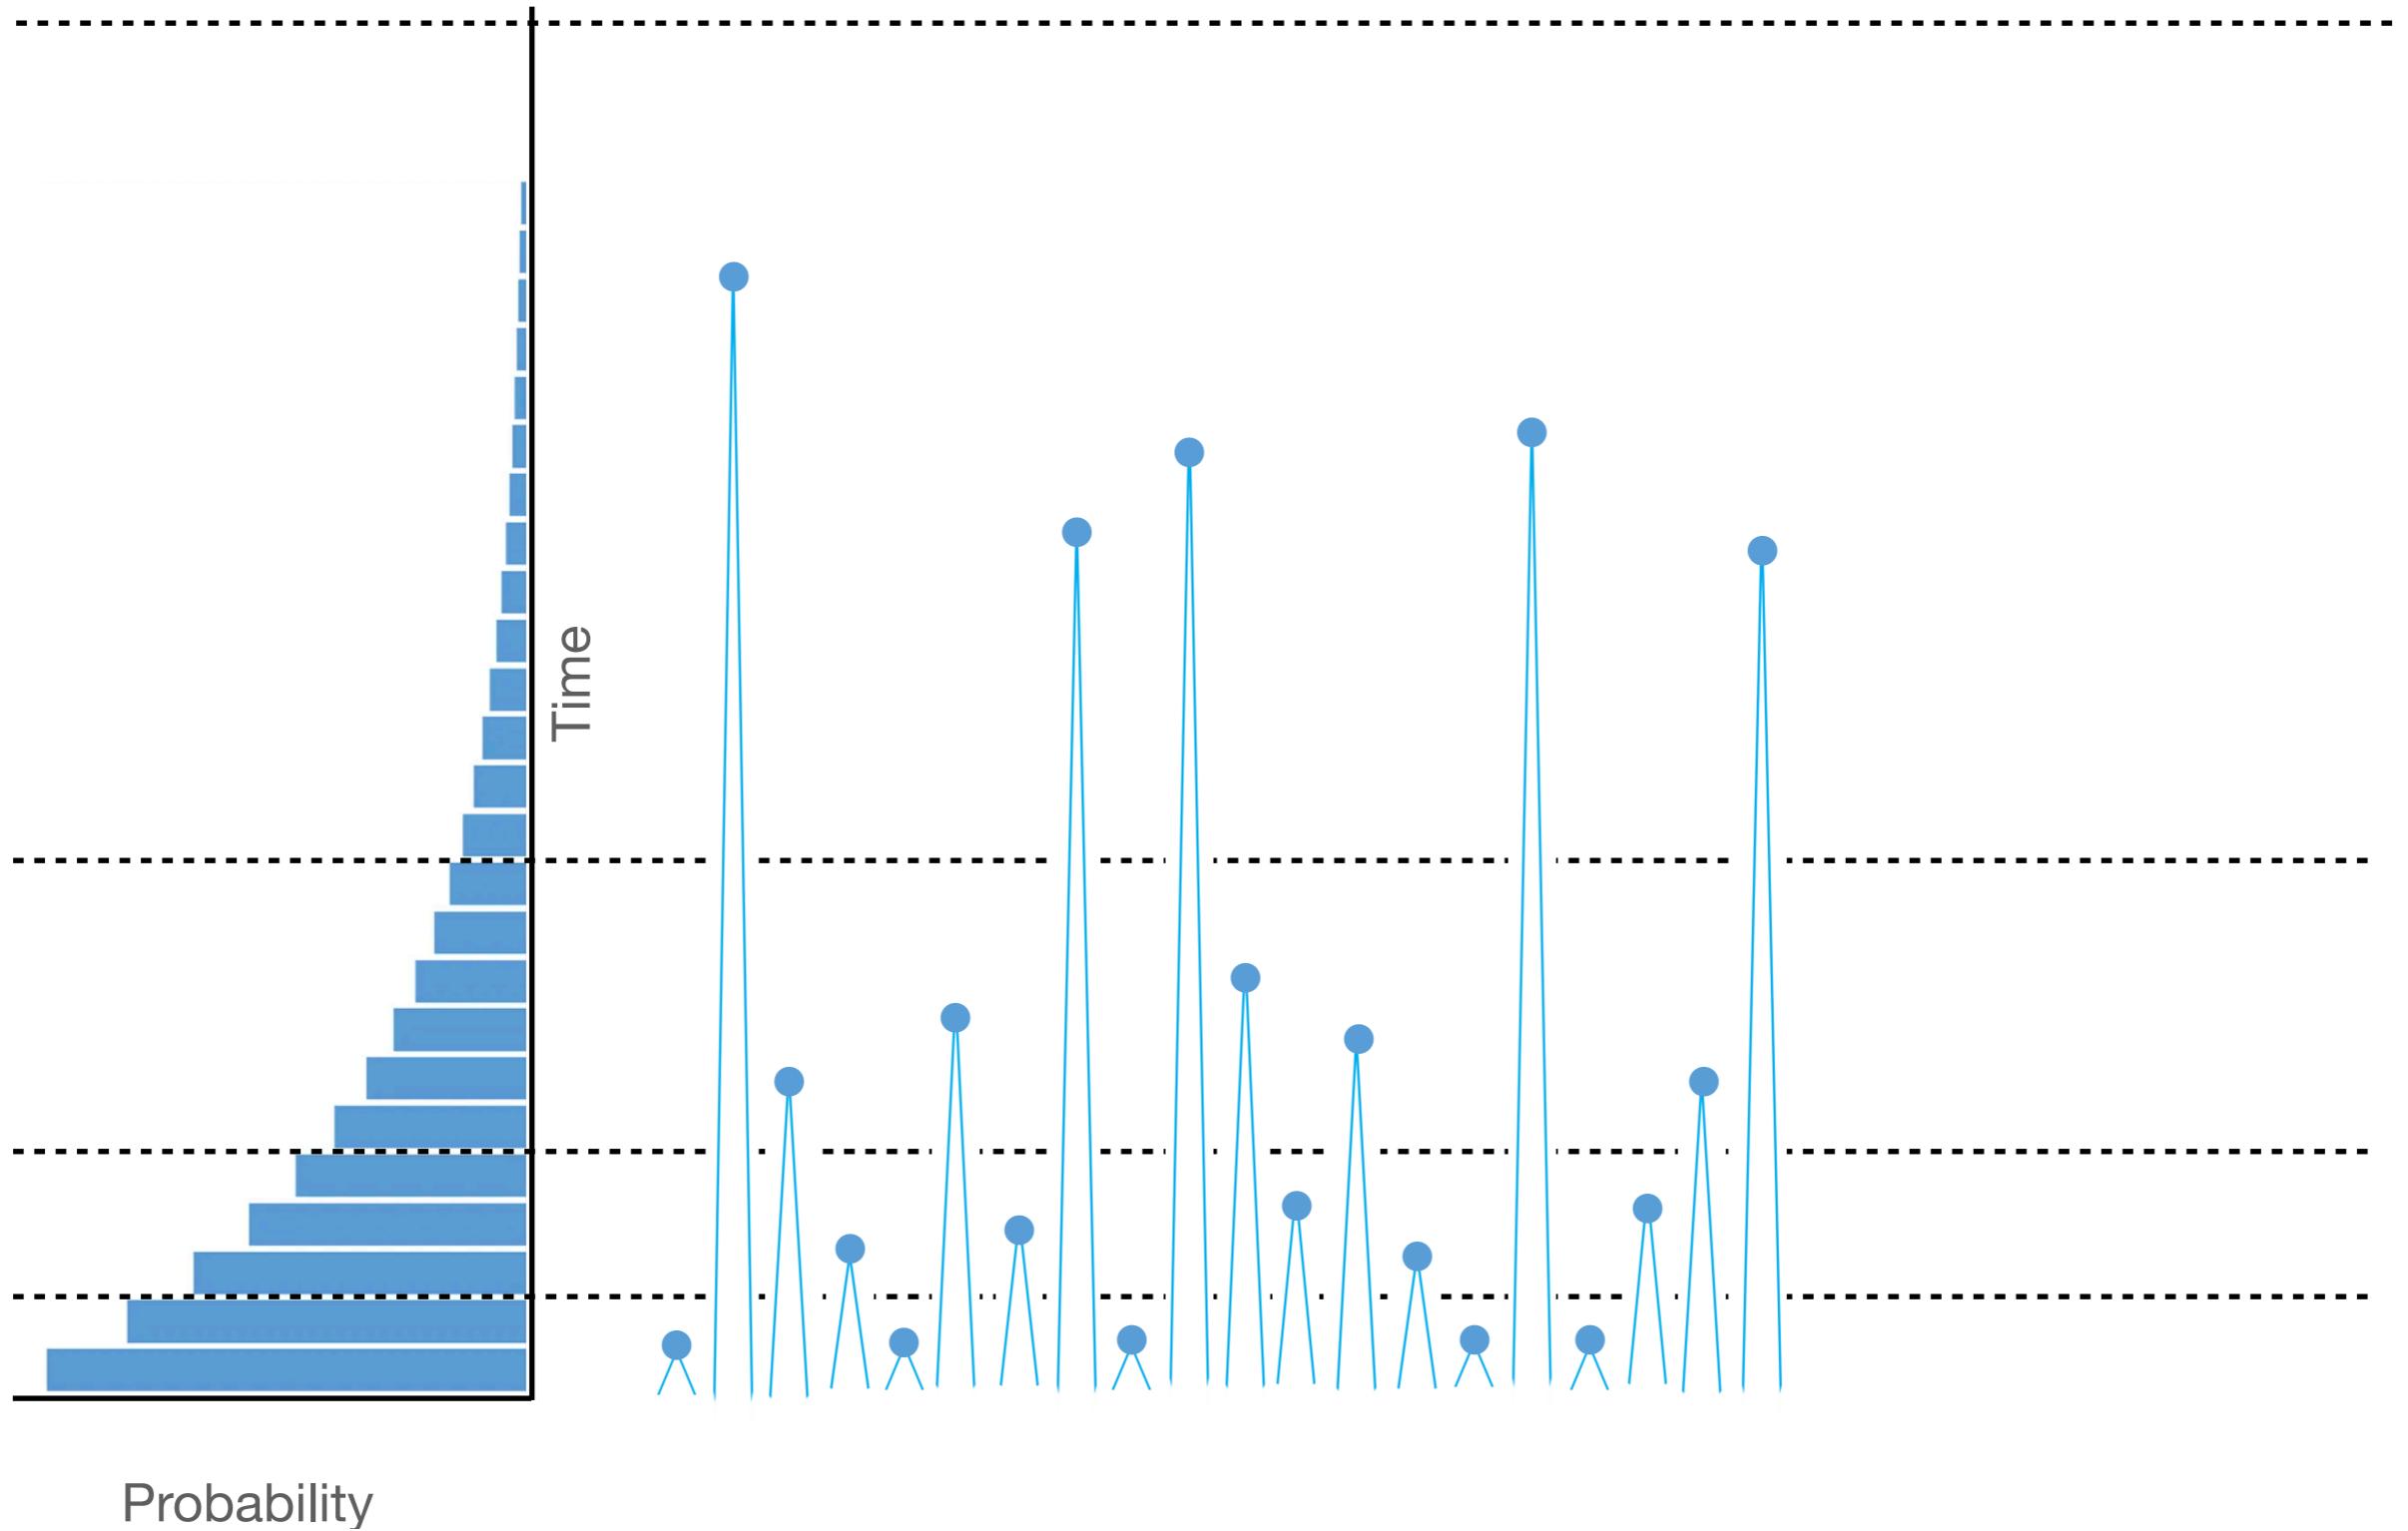
Logarithmic bins

Makes coalescence in each bin equally likely



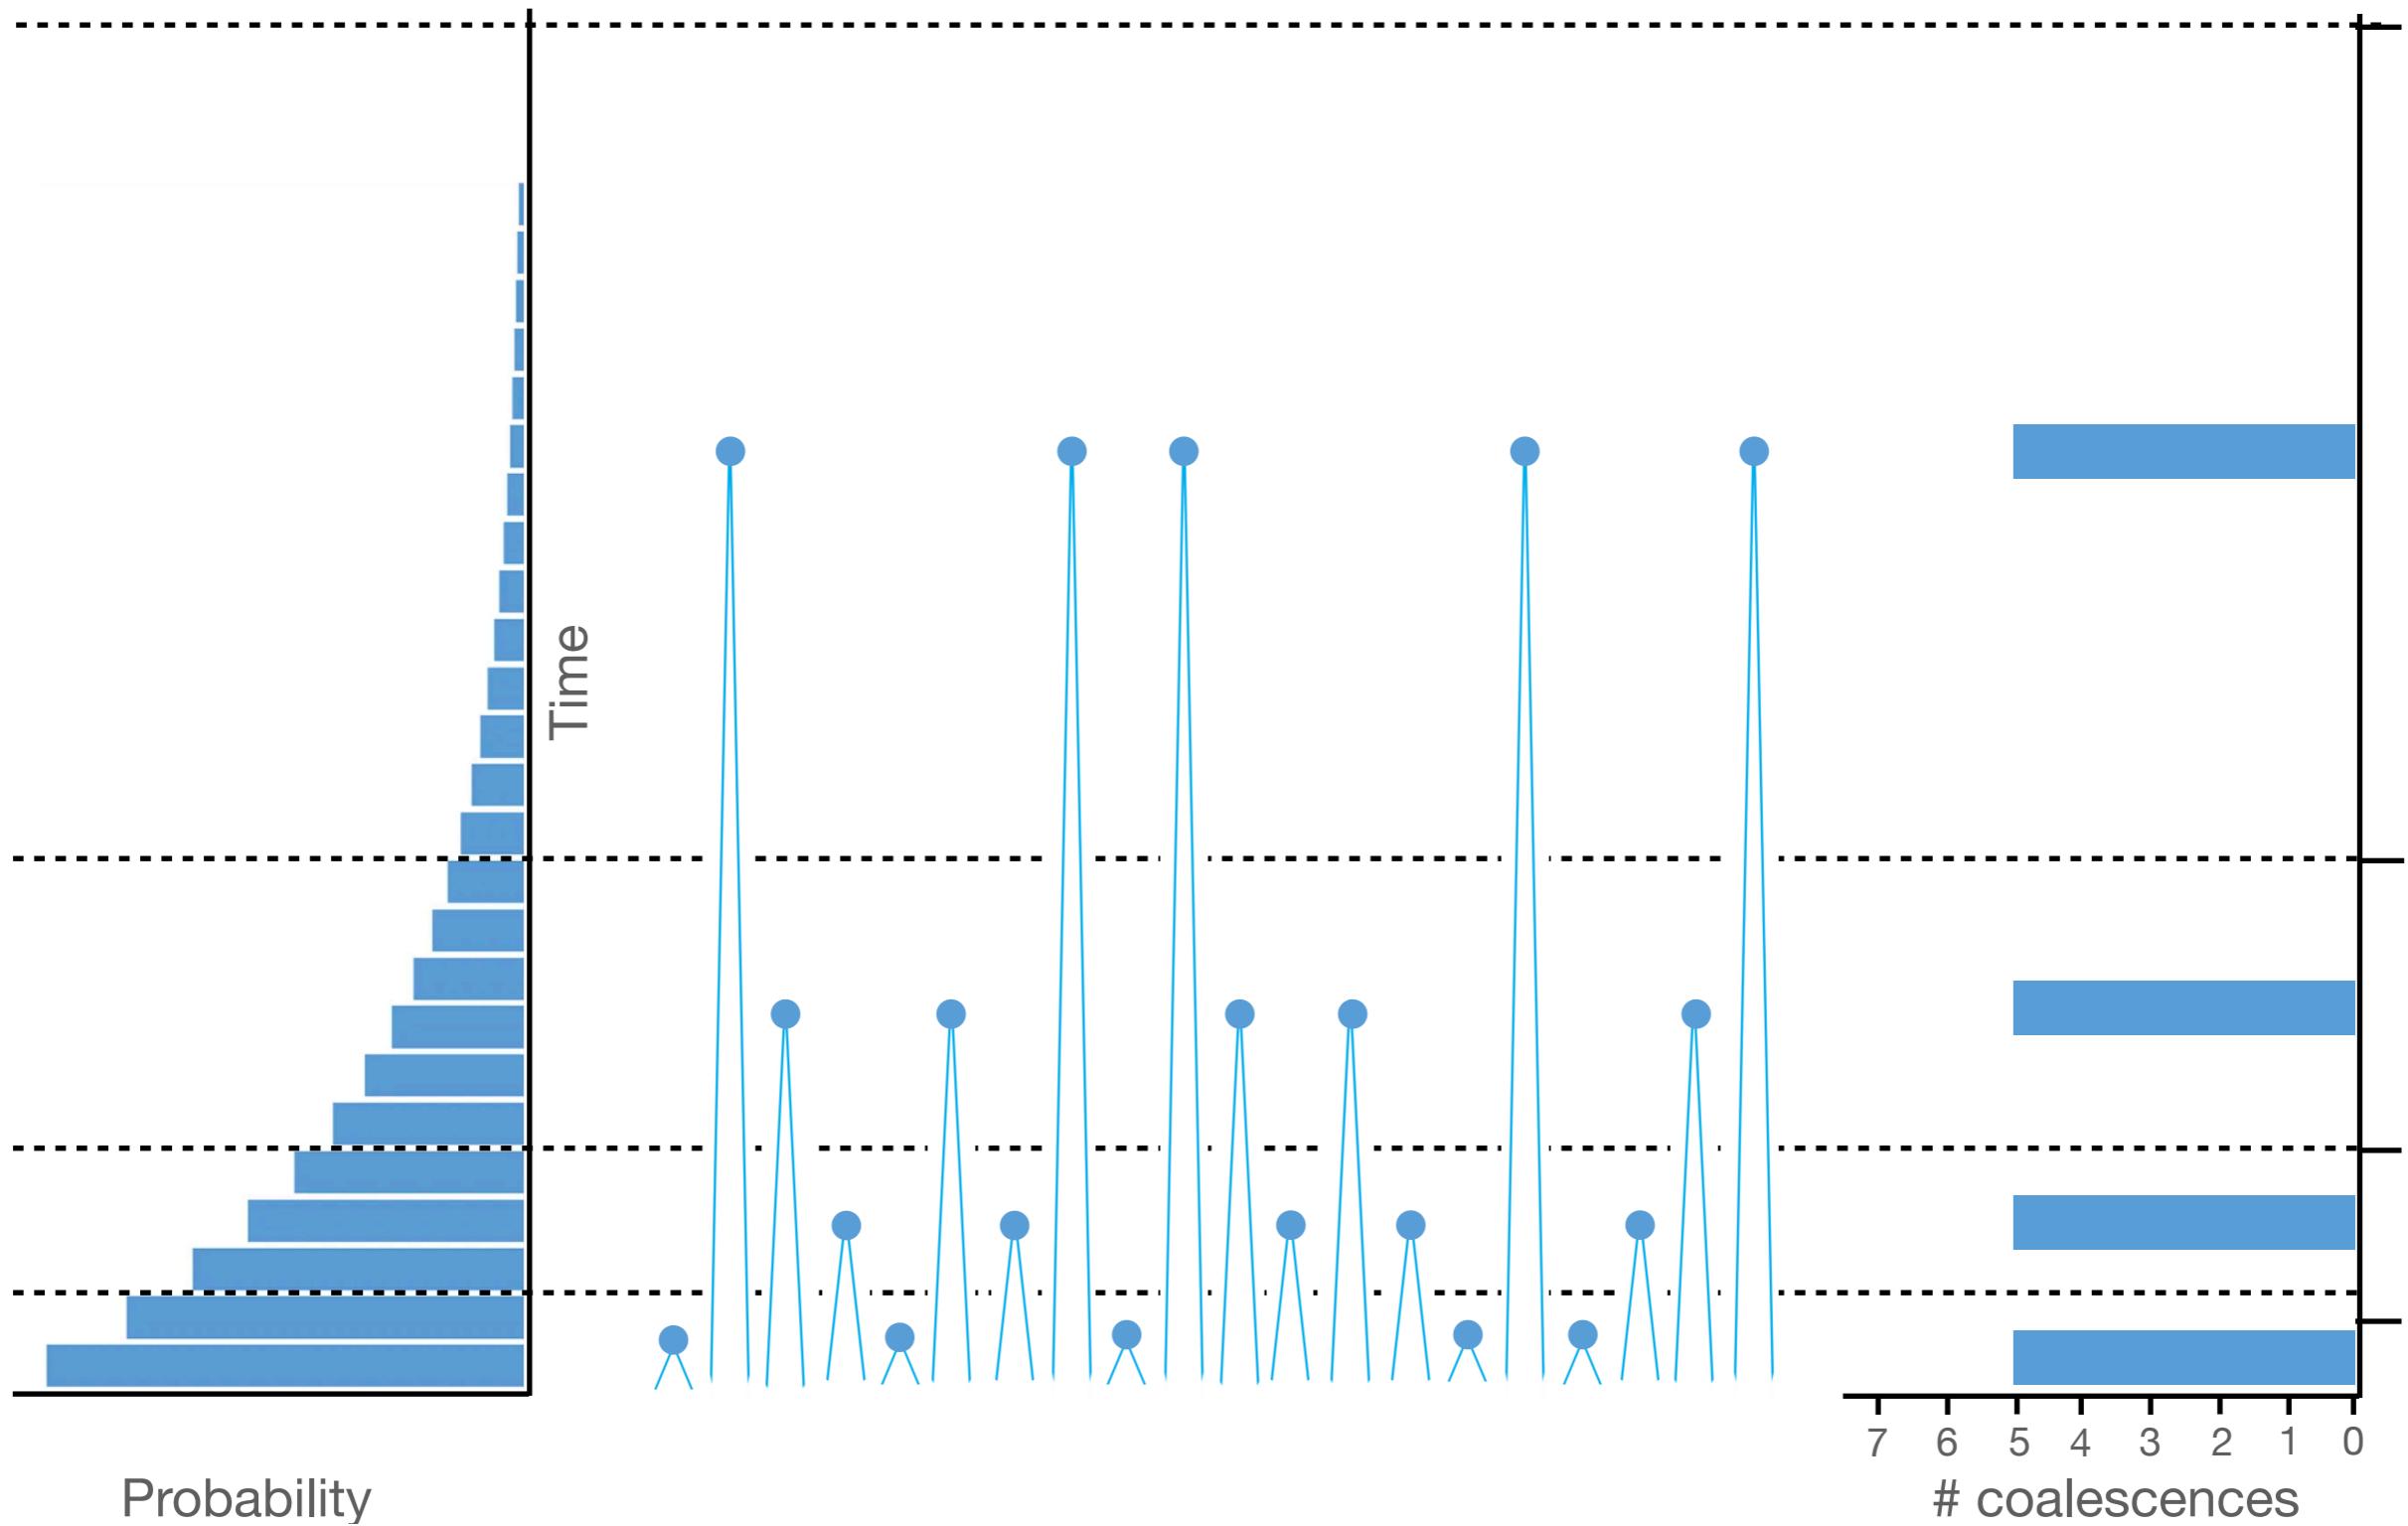
Thought experiment

Suppose we knew the TMRCA of each sequence segment



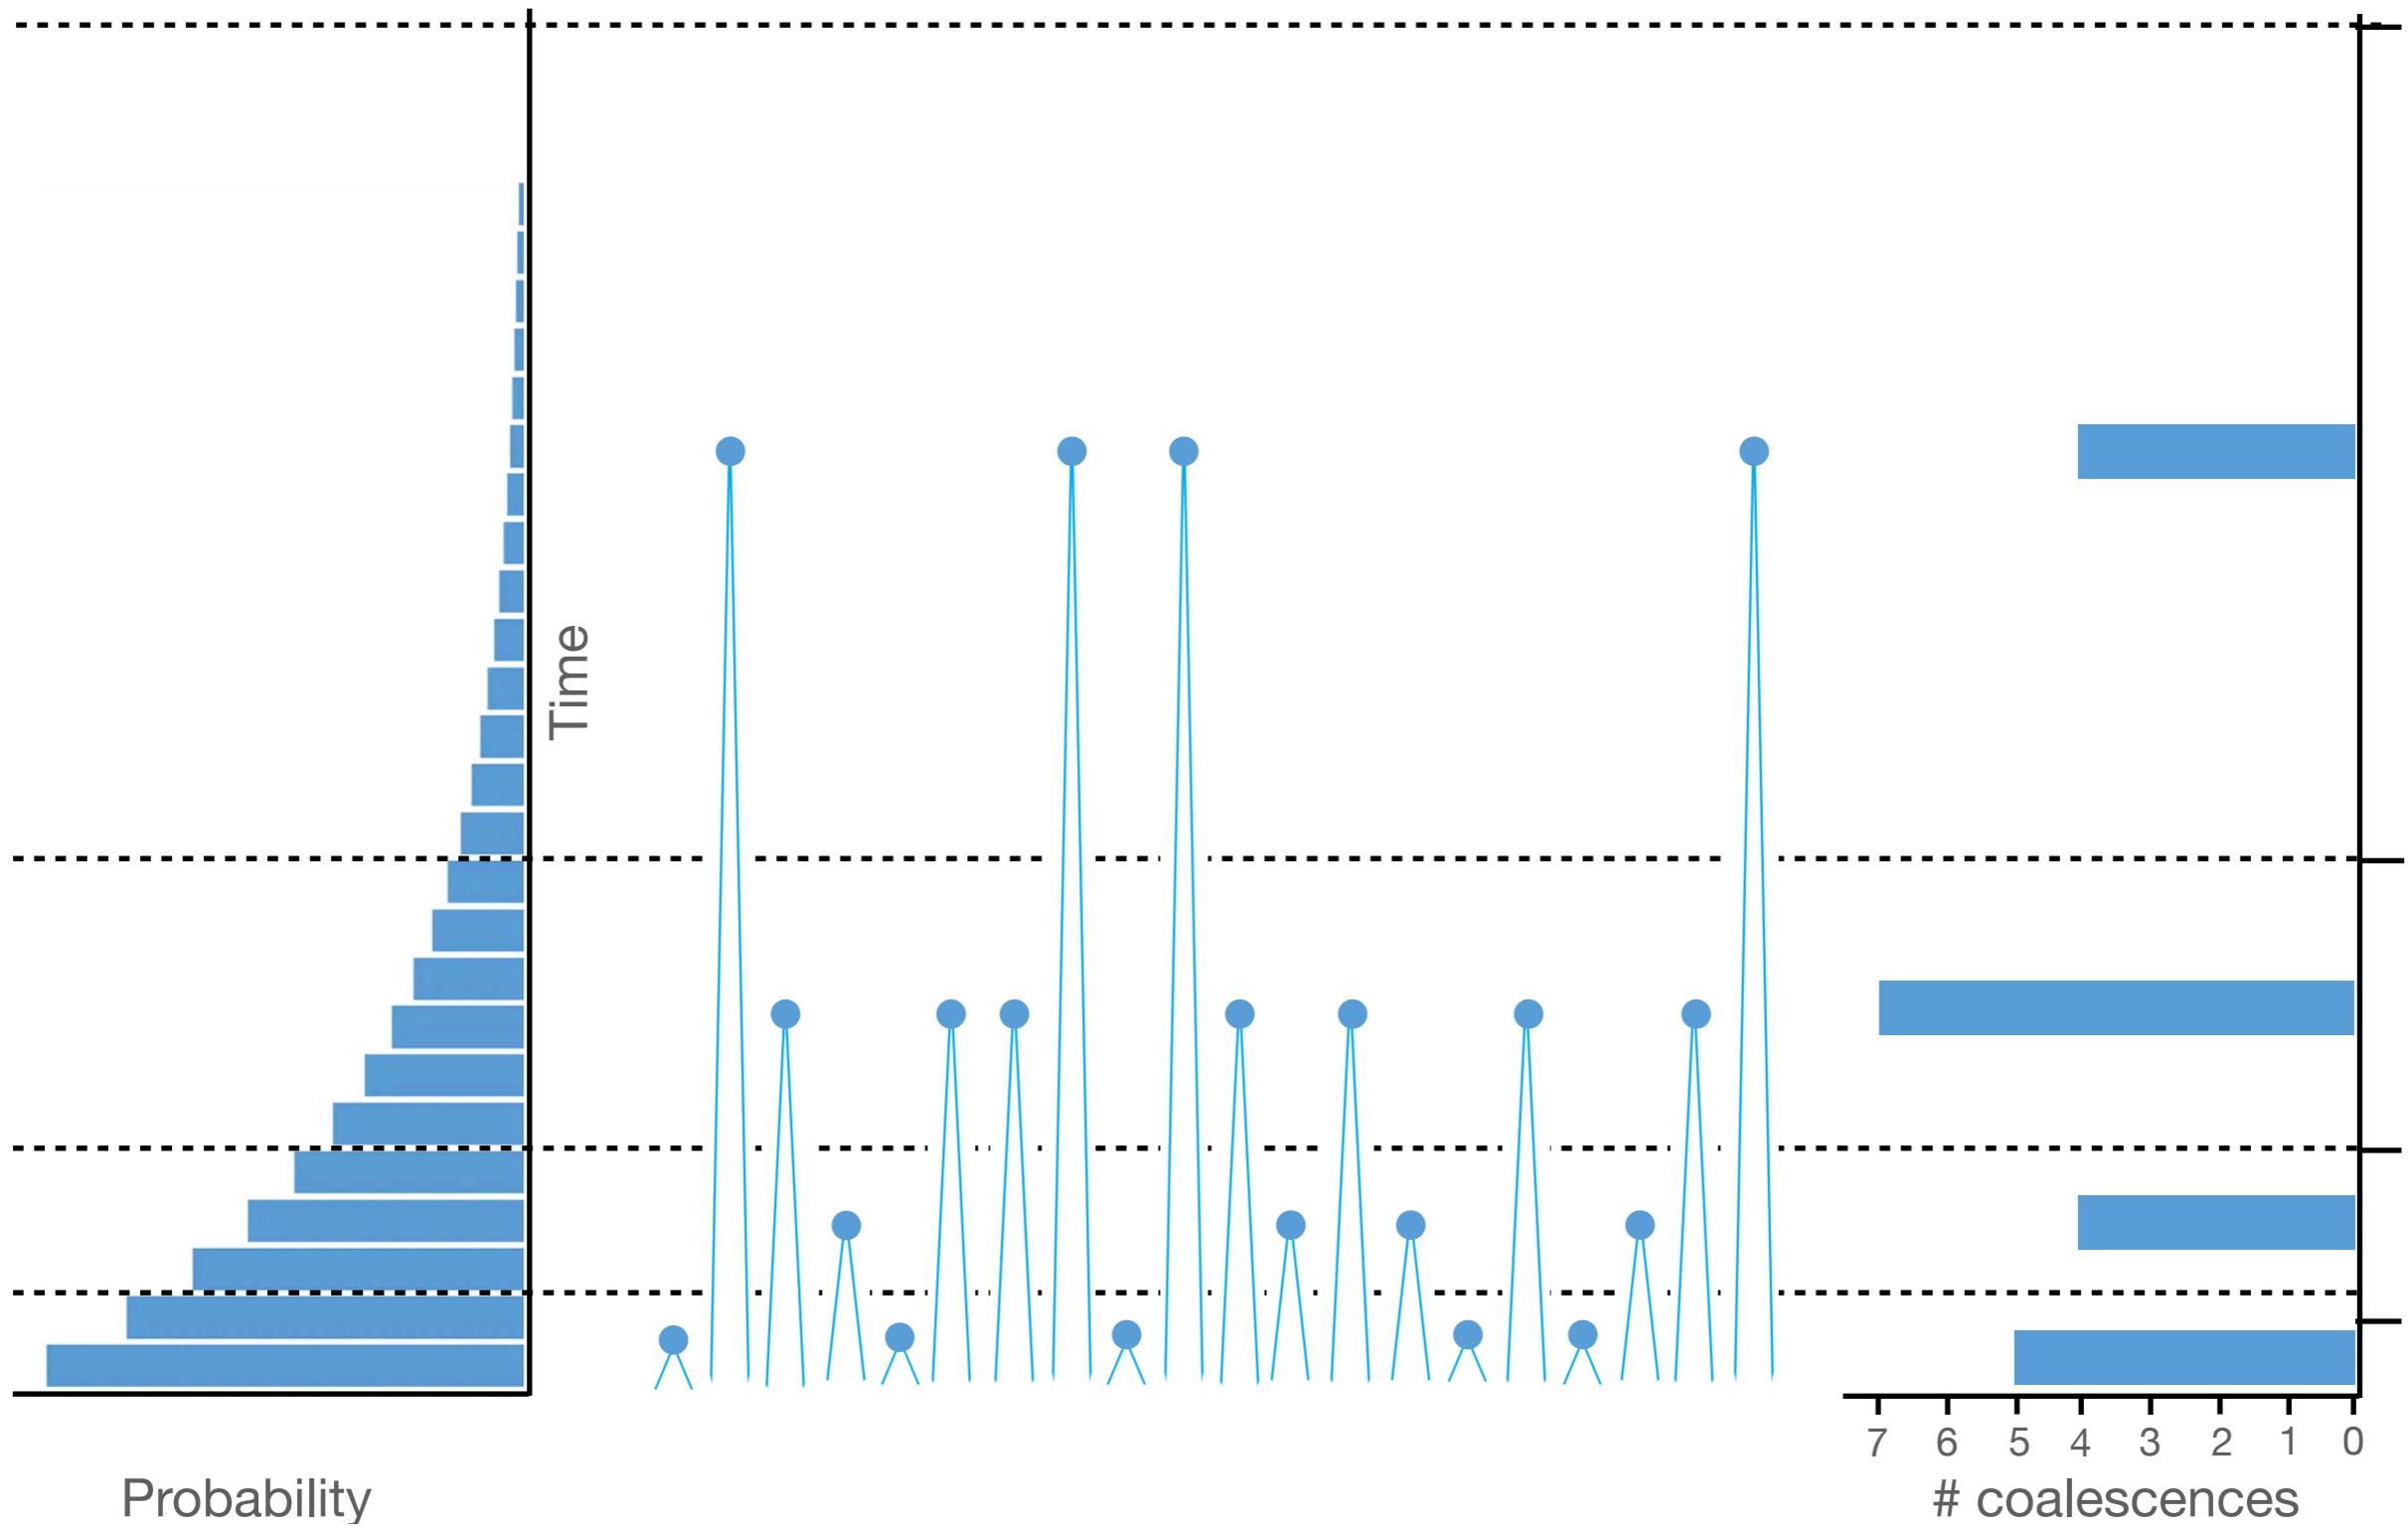
Thought experiment

Suppose we knew the TMRCA of each sequence segment



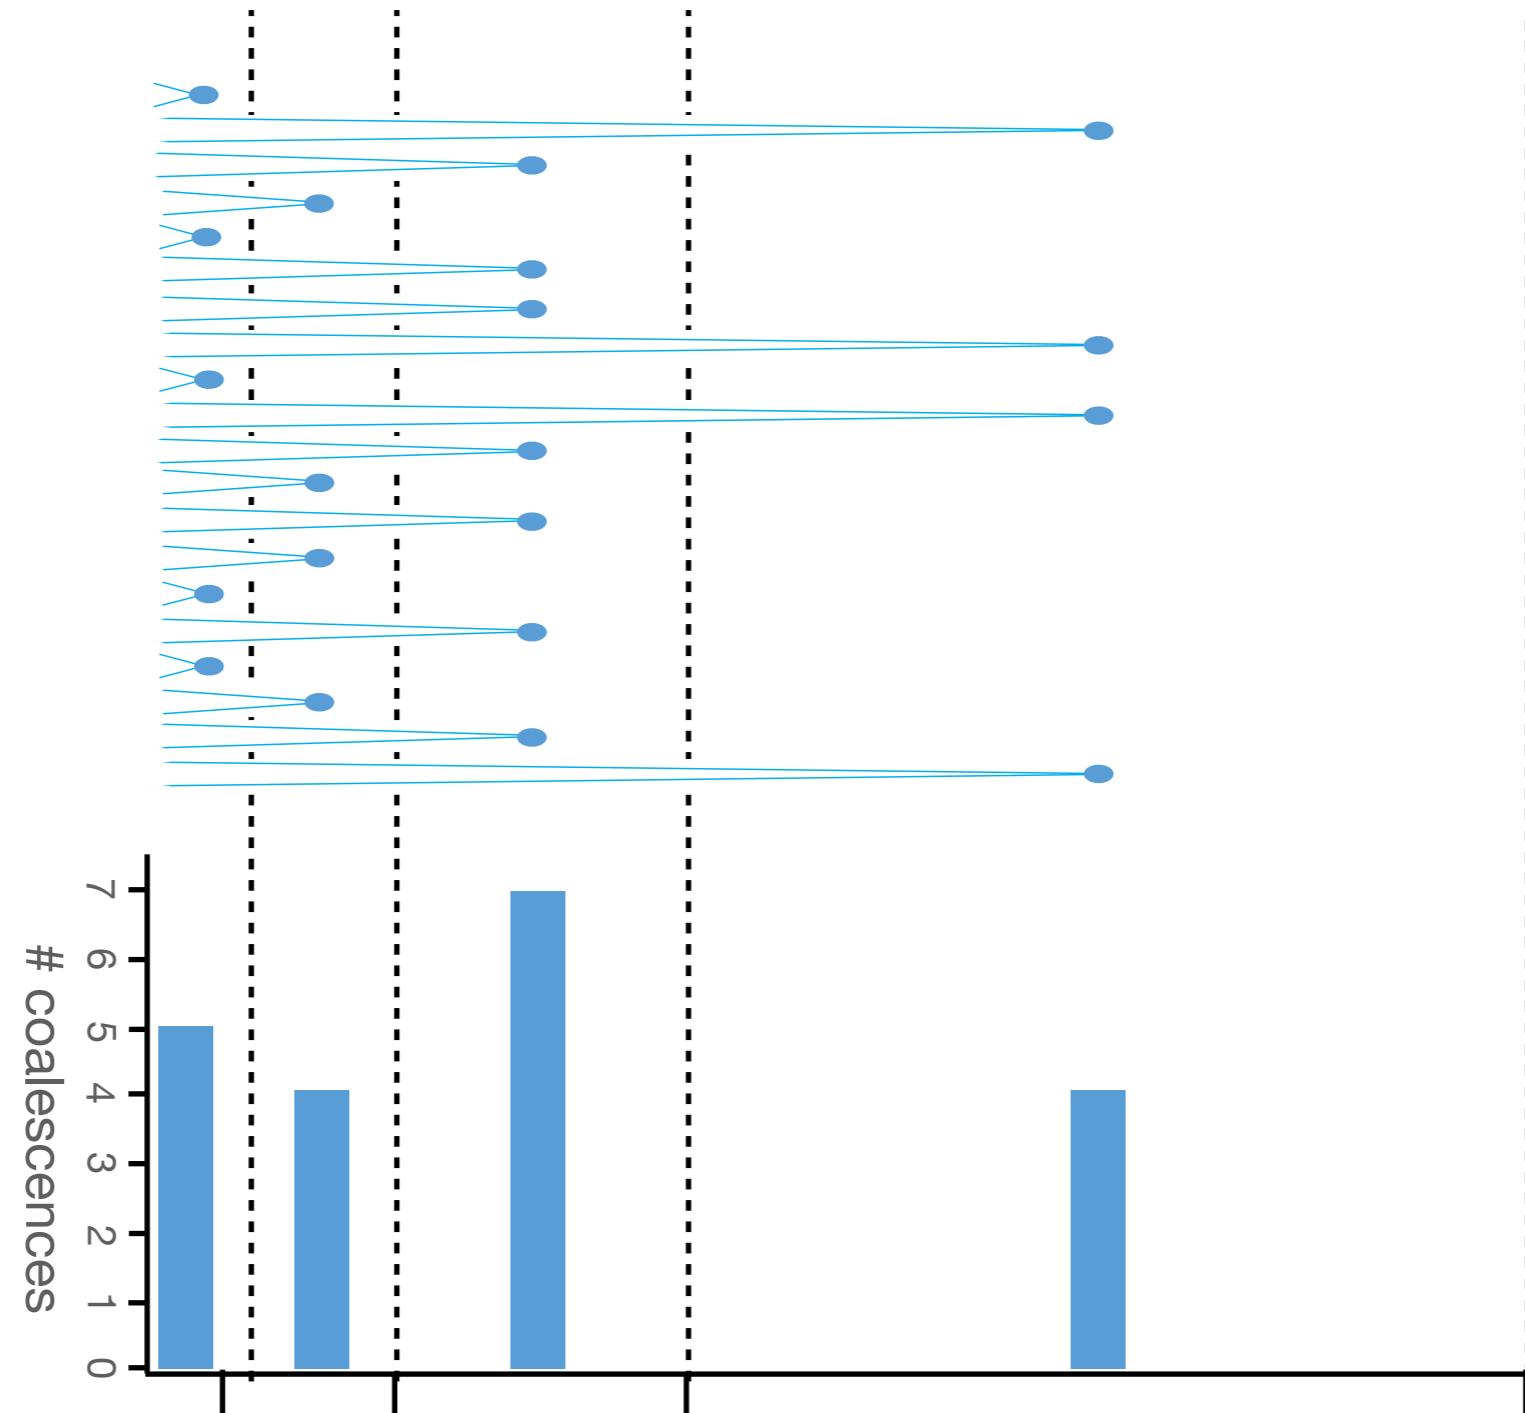
Thought experiment

Suppose we knew the TMRCA of each sequence segment



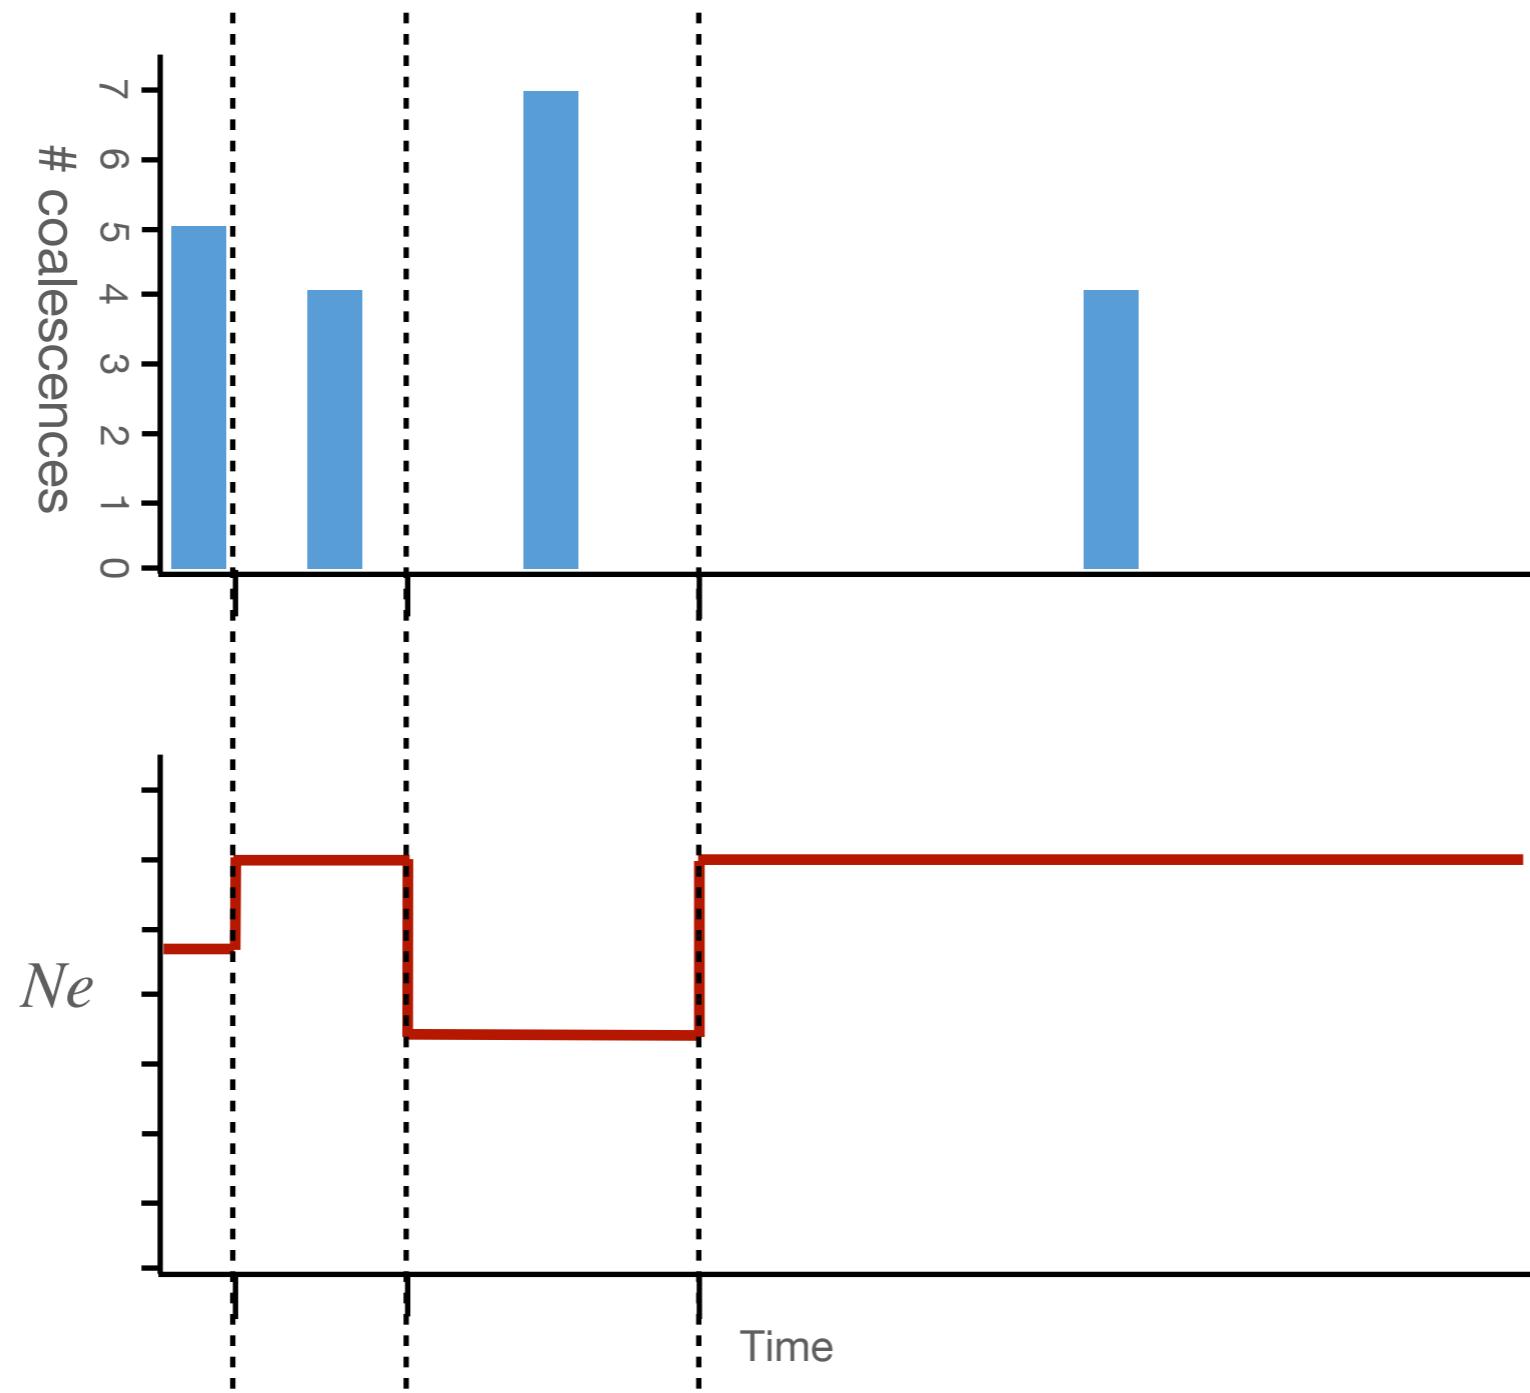
Thought experiment

Suppose we knew the TMRCA of each sequence segment



Thought experiment

Suppose we knew the TMRCA of each sequence segment

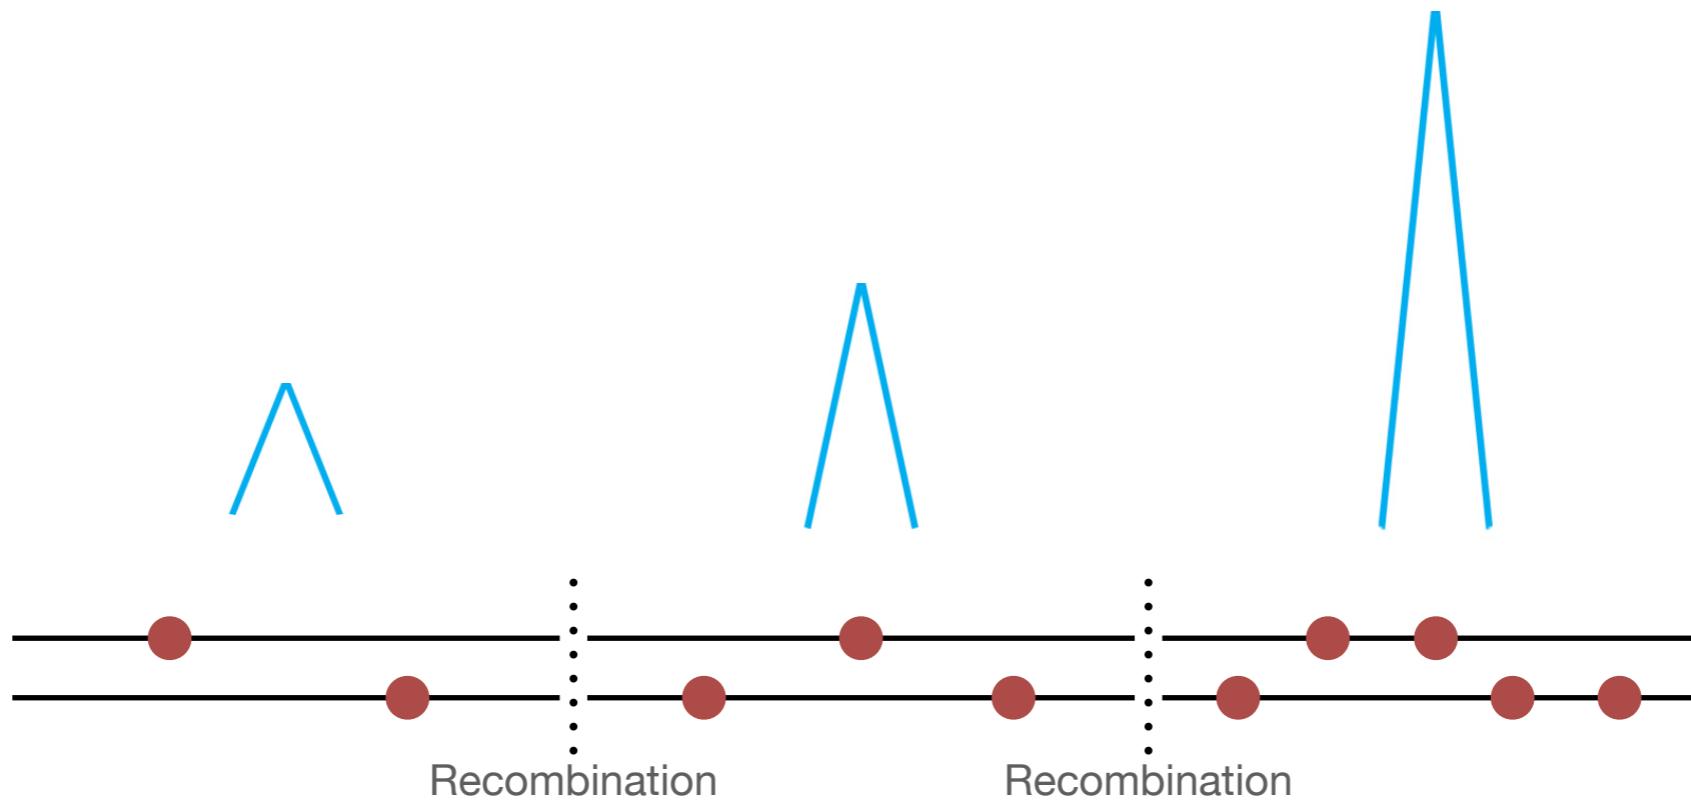


Back to reality

- **But don't despair**
- We **do not** have:
 - The segments between recombination events.
 - The TMRCA of each segment.
- We **do** have:
 - The sequence differences.
 - A model (the ARG) that relate sequence differences to segment length and TMRCA.

Mutations and TMRCA

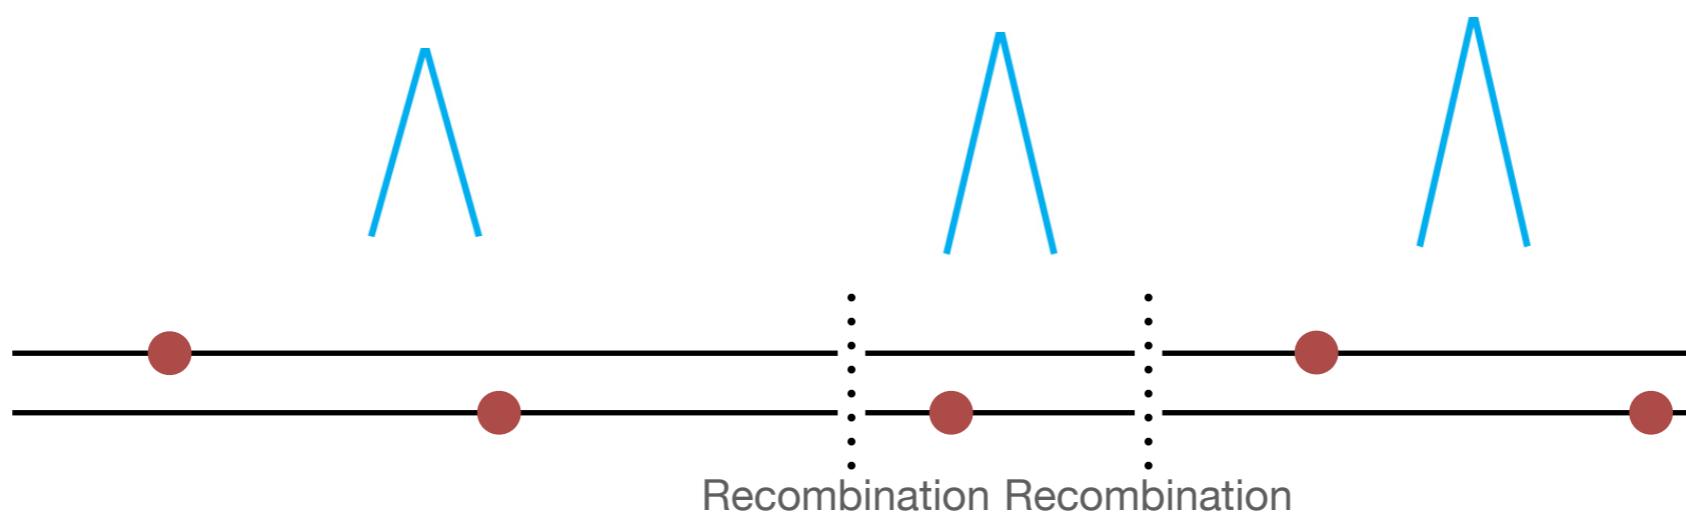
Which segment has the longest TMRCA?



- Why?

Mutations and TMRCA

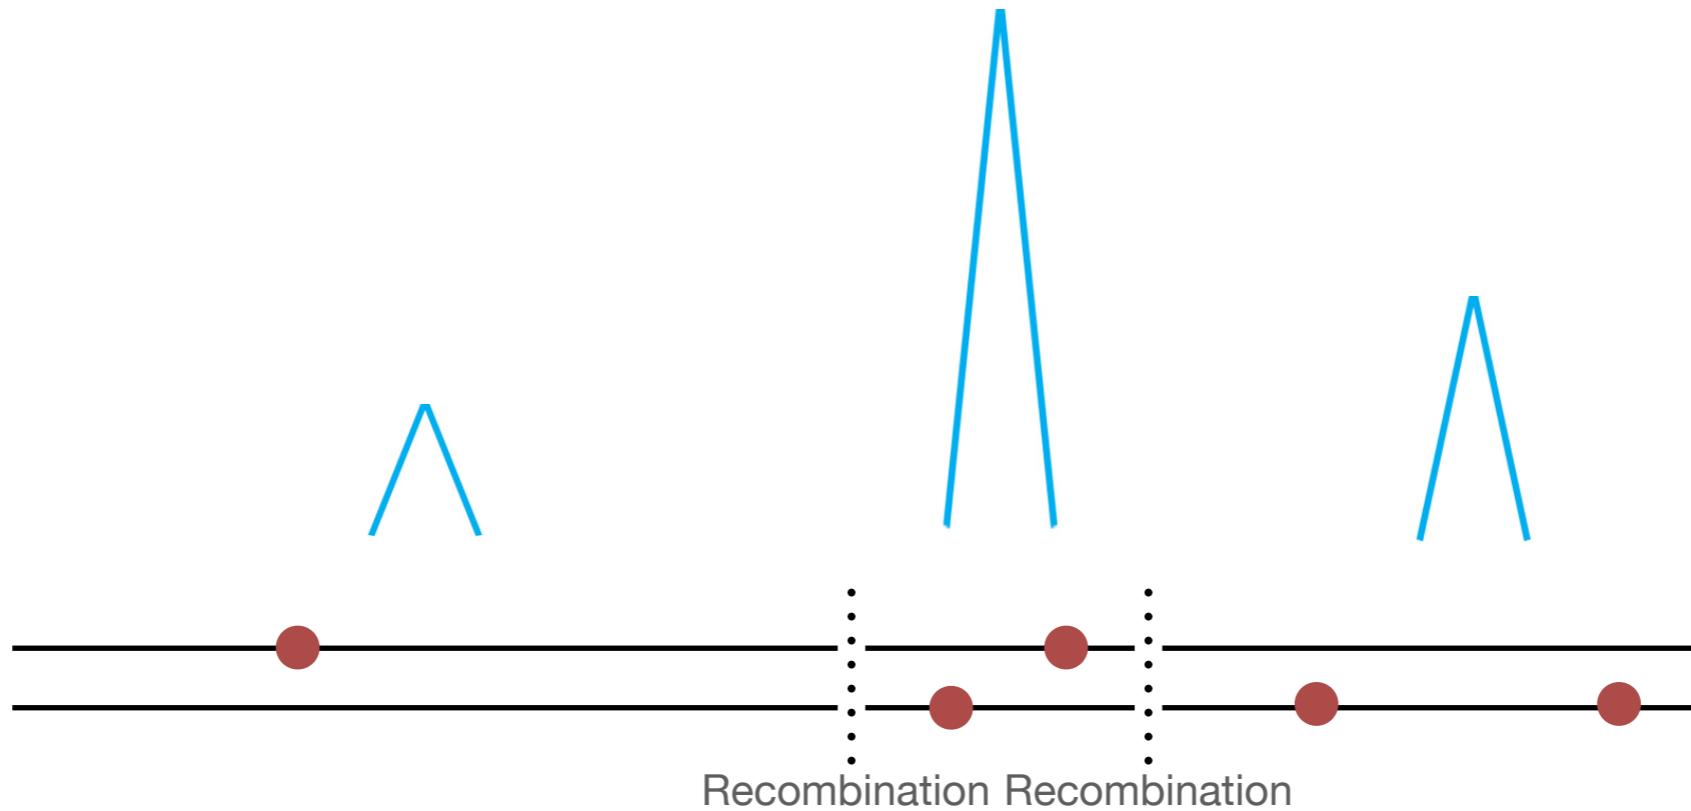
Which segment has the longest TMRCA?



- Why?

Mutations and TMRCA

Which segment has the longest TMRCA?



- Why?

Tree heights and sequence differences

Observed data:

Homozygote: 0

Heterozygote: 1

0 0 0 0 1 0 0 1 1 0

Medium tree: M



Tall tree: T



Short tree: S



Hidden Markov model

Tree heights and sequence differences

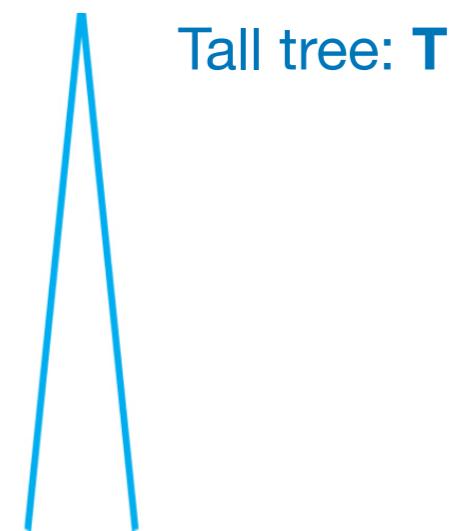
Observed data:

Homozygote: 0

Heterozygote: **1**

Hidden states:

Medium tree: **M**



0 0 0 0 **1** 0 0 1 1 0

Hidden Markov model

Tree heights and sequence differences

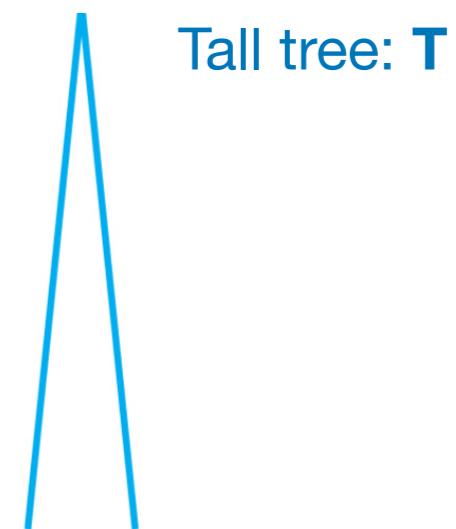
Observed data:

Homozygote: 0

Heterozygote: 1

Hidden states:

Medium tree: M

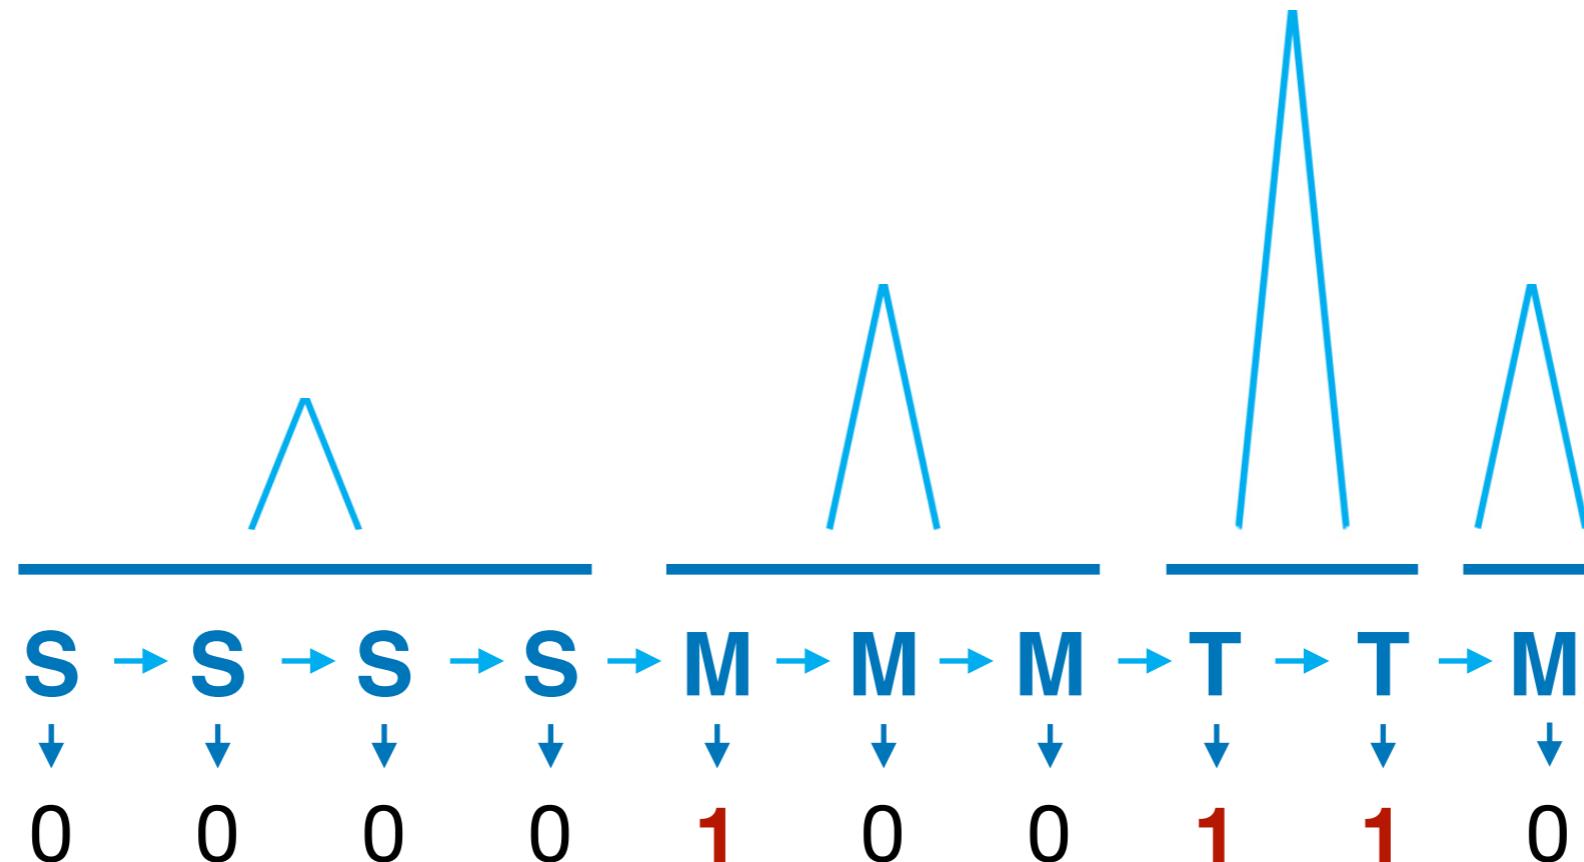


One possible hidden state path:

S → S → S → S → M → M → M → T → T → M
↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓
0 0 0 0 1 0 0 1 1 0

Hidden Markov model

Tree heights and sequence differences



Hidden Markov model

Tree heights and sequence differences

Emission probs.

$$e(1 | t) = e^{-\theta t}$$

$$e(0 | t) = 1 - e^{-\theta t}$$

Medium tree: **M**

Emission probs.

1:

0:

Transition probs.

$$p(t | s) = \dots$$

Tall tree: **T**

Emission probs.

1:

0:



Short tree: **S**

Emission probs.

1:

0:



Hidden Markov model

Tree heights and sequence differences

Emission probs.

$$e(1 | t) = 1 - e^{-\theta t}$$

$$e(0 | t) = e^{-\theta t}$$

Transition probs.

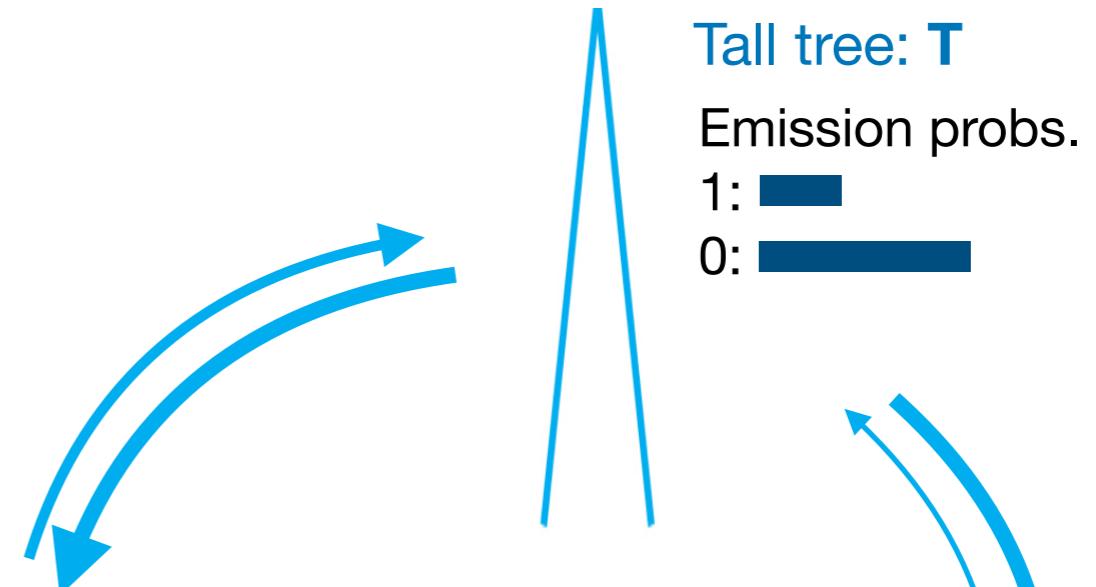
$$p(t | s) = \dots$$

Medium tree: **M**

Emission probs.

1:

0:



Tall tree: **T**

Emission probs.

1:

0:

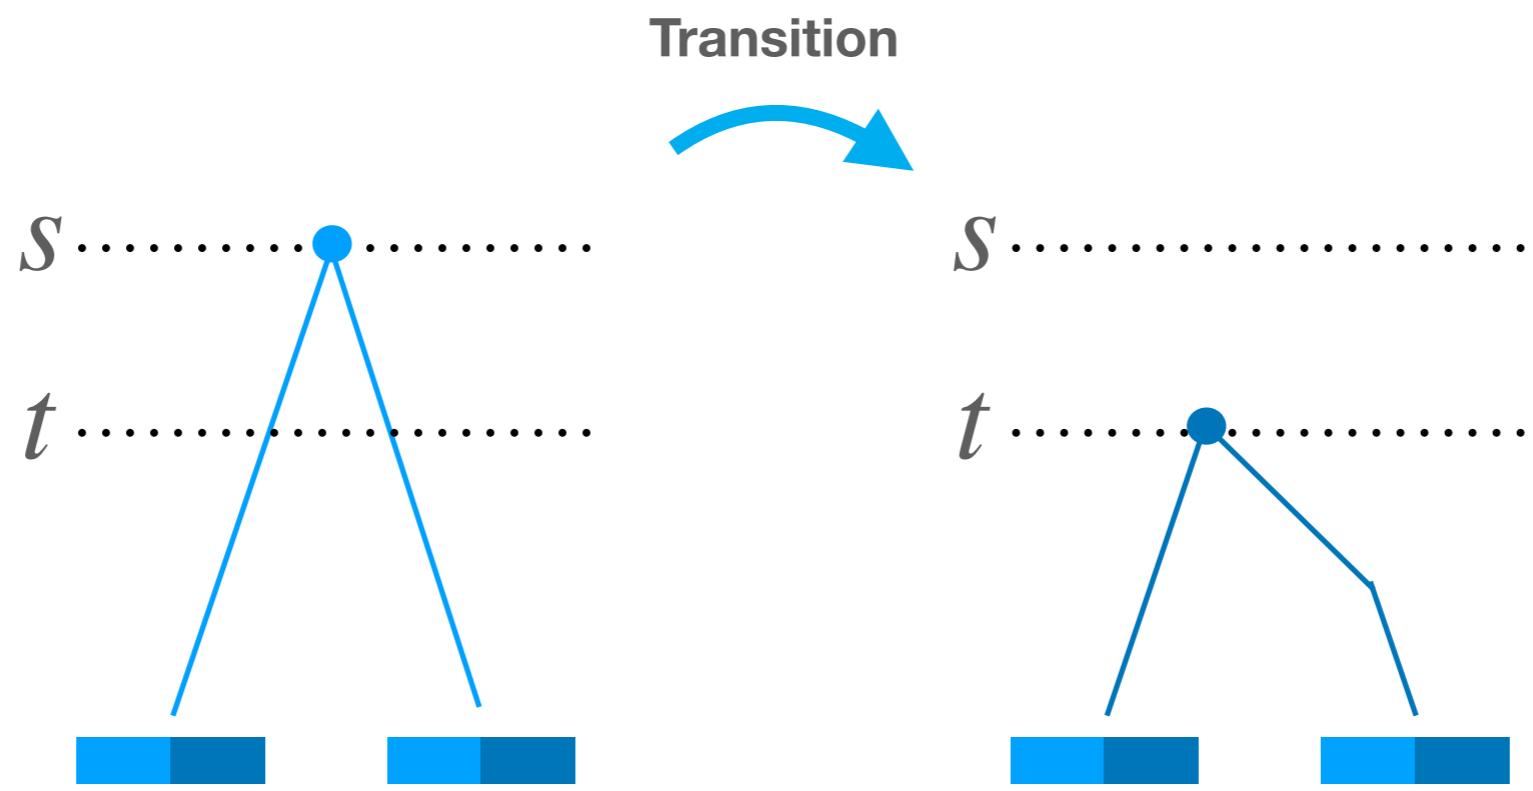
Short tree: **S**

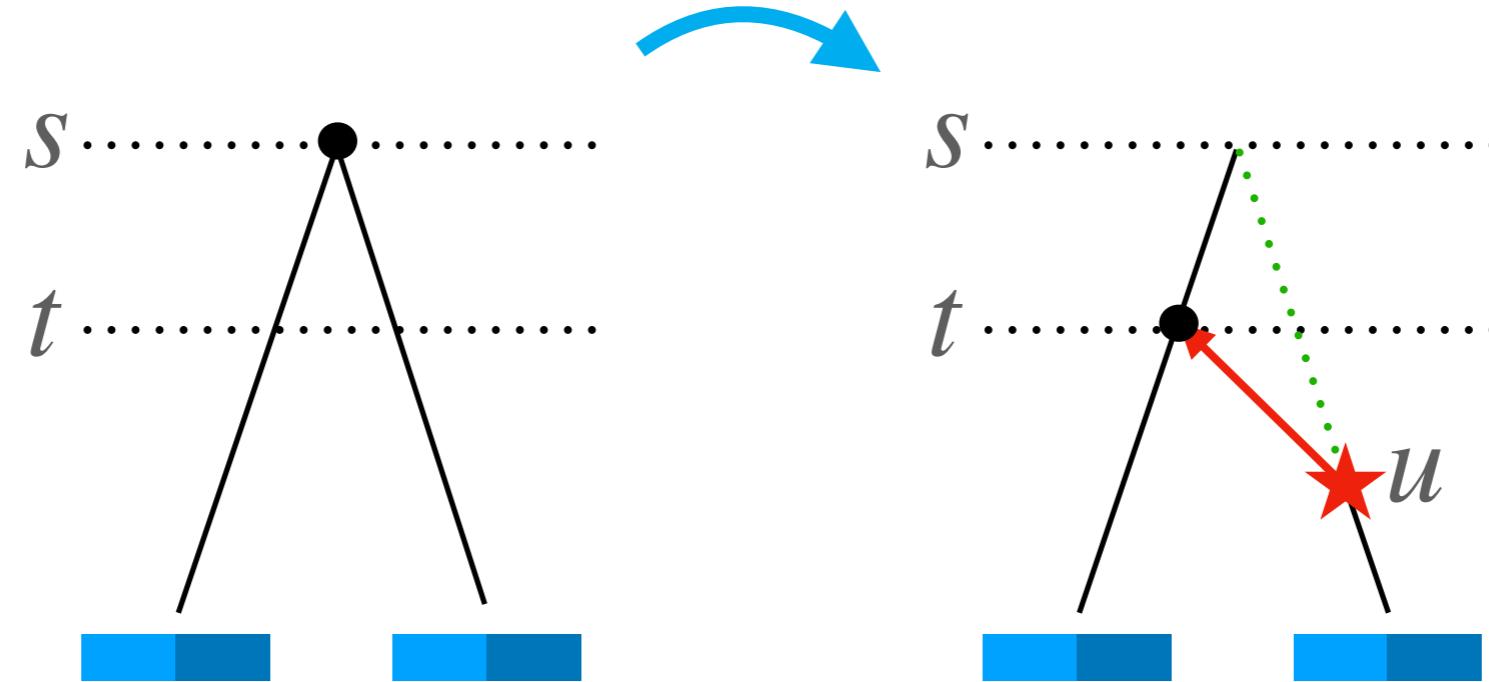
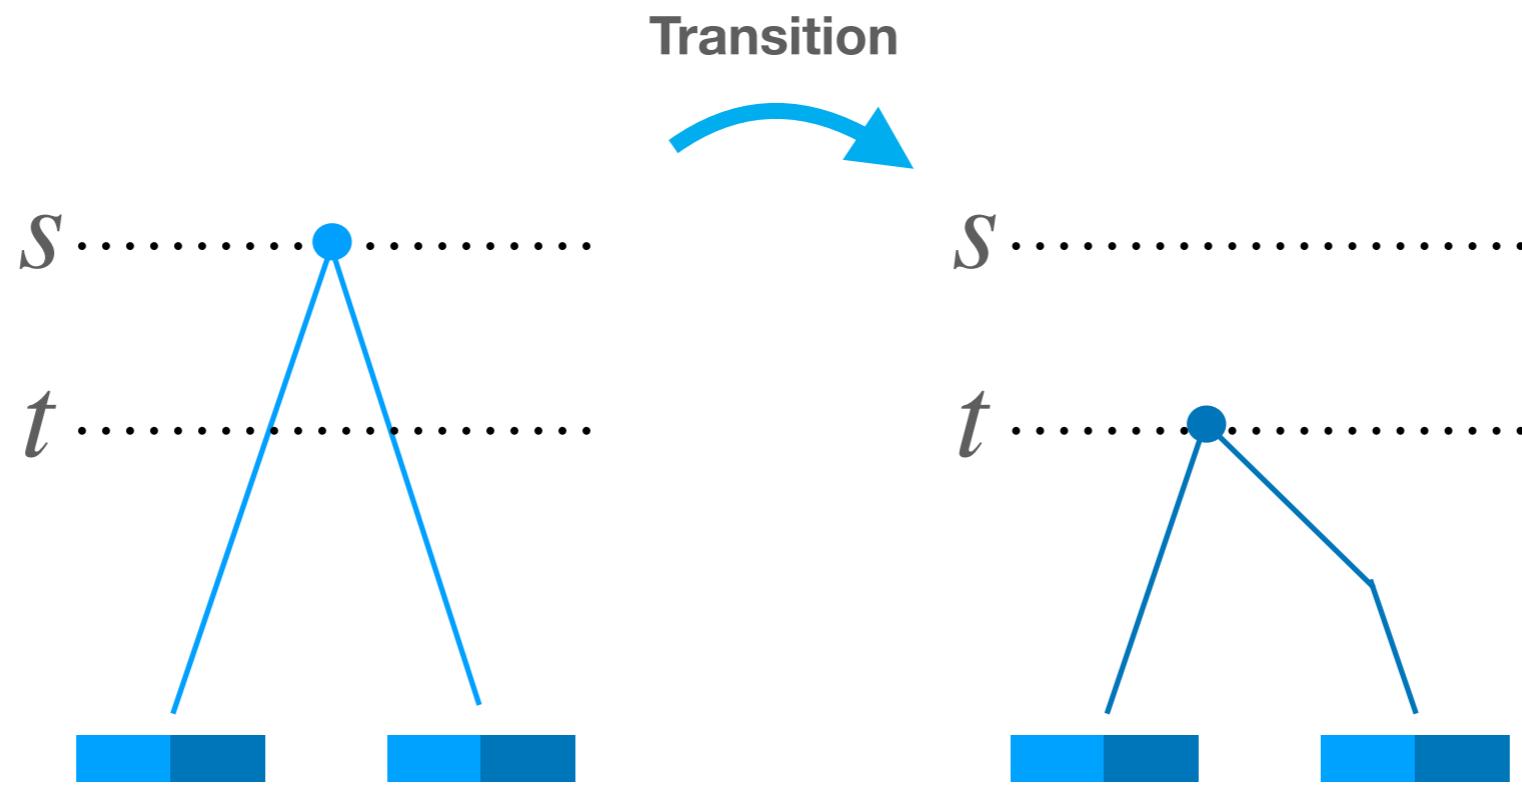
Emission probs.

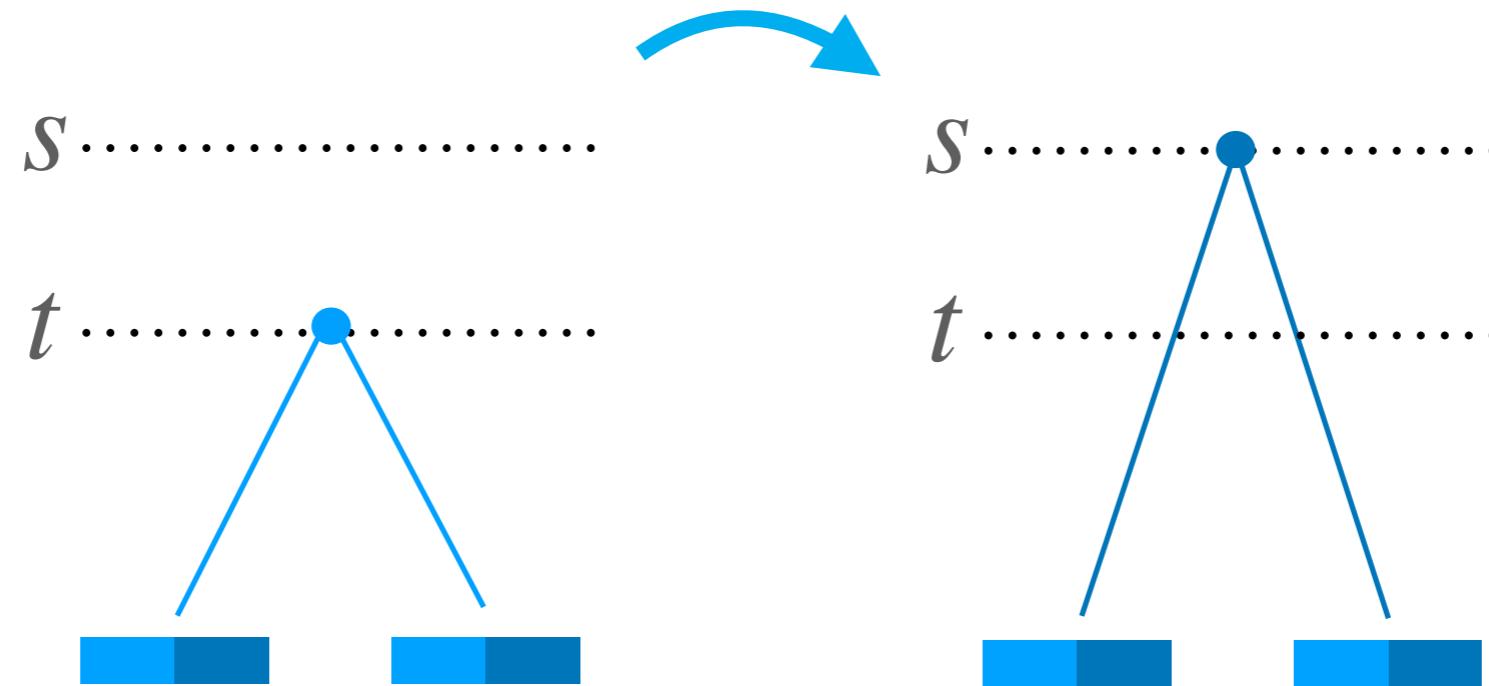
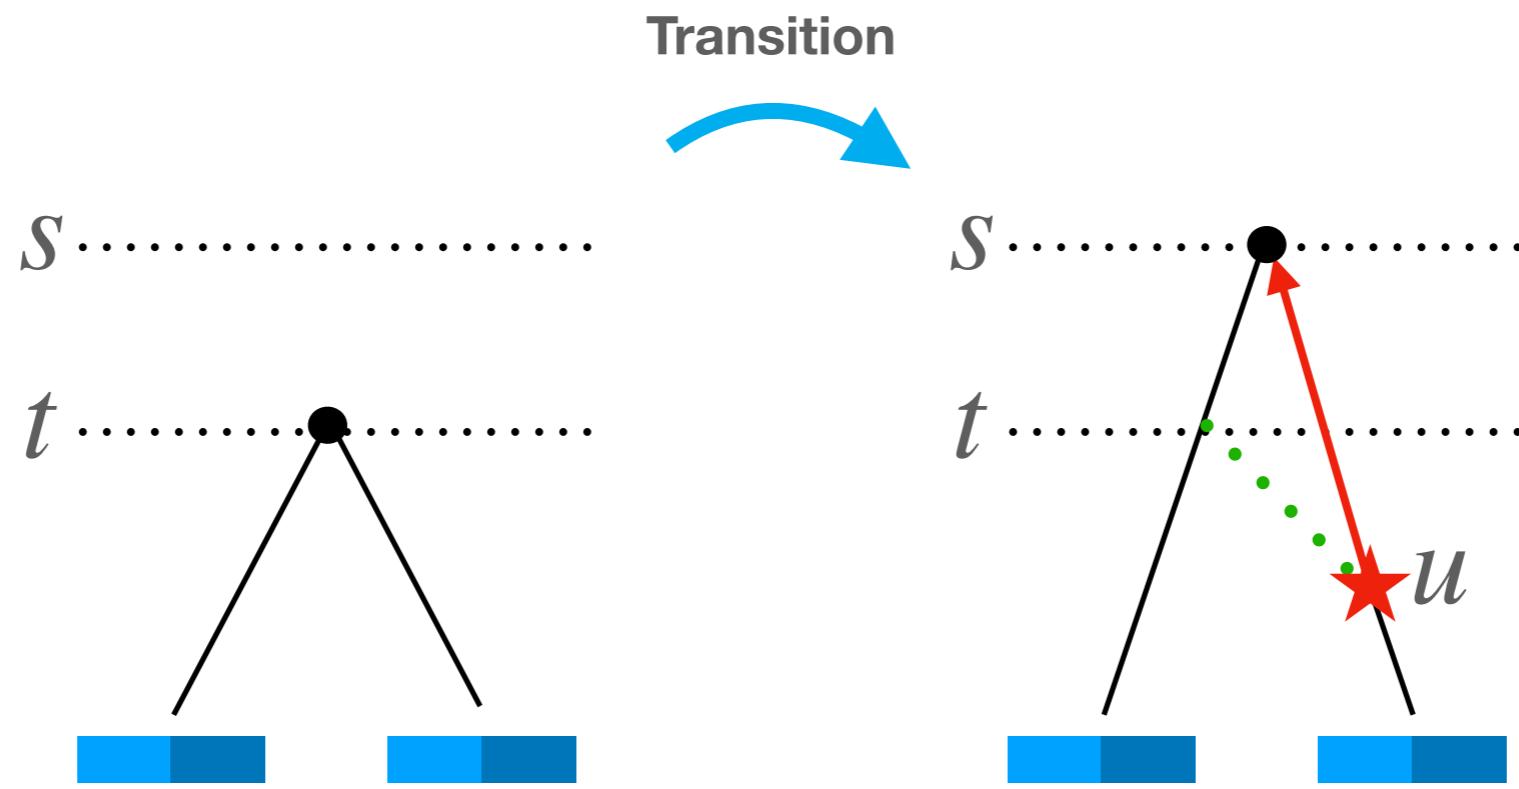
1:

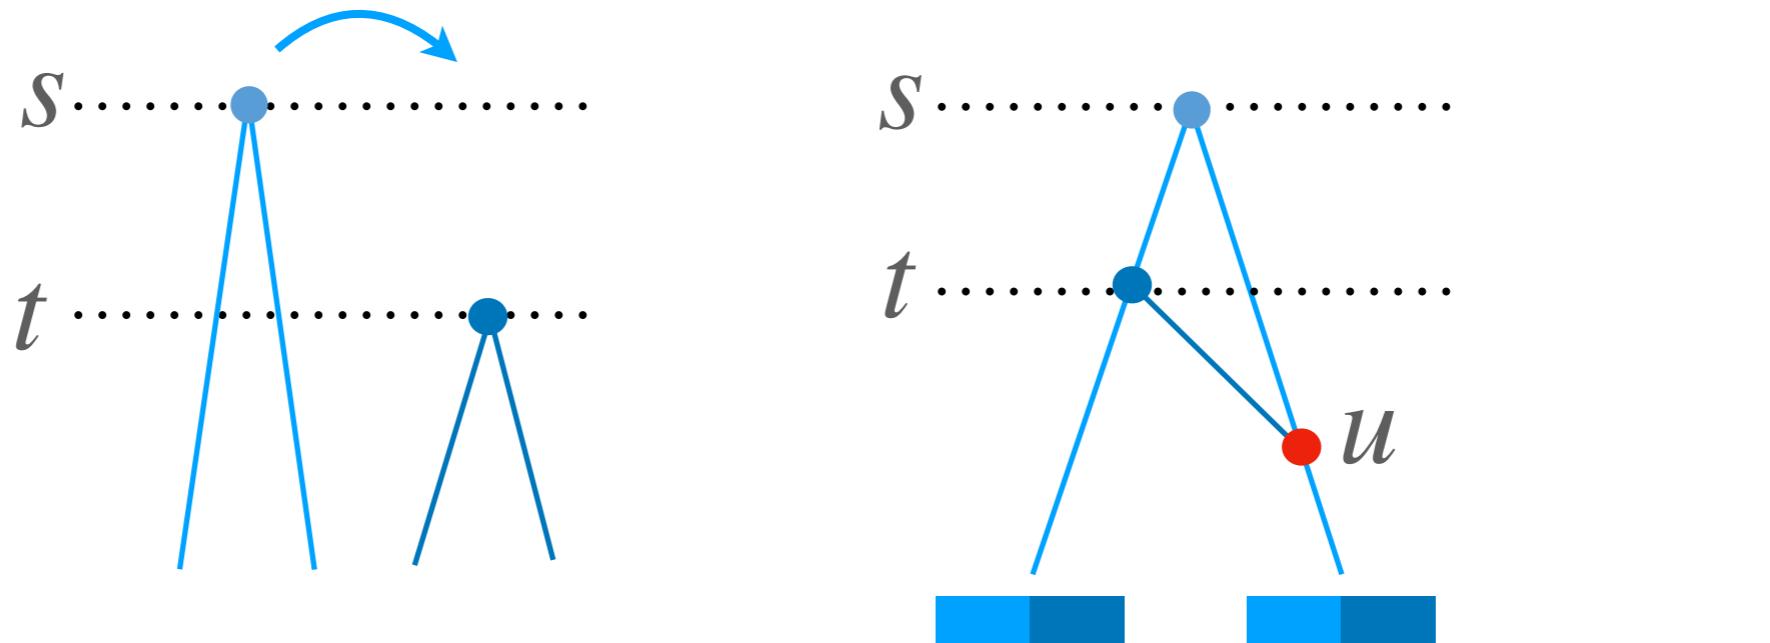
0:







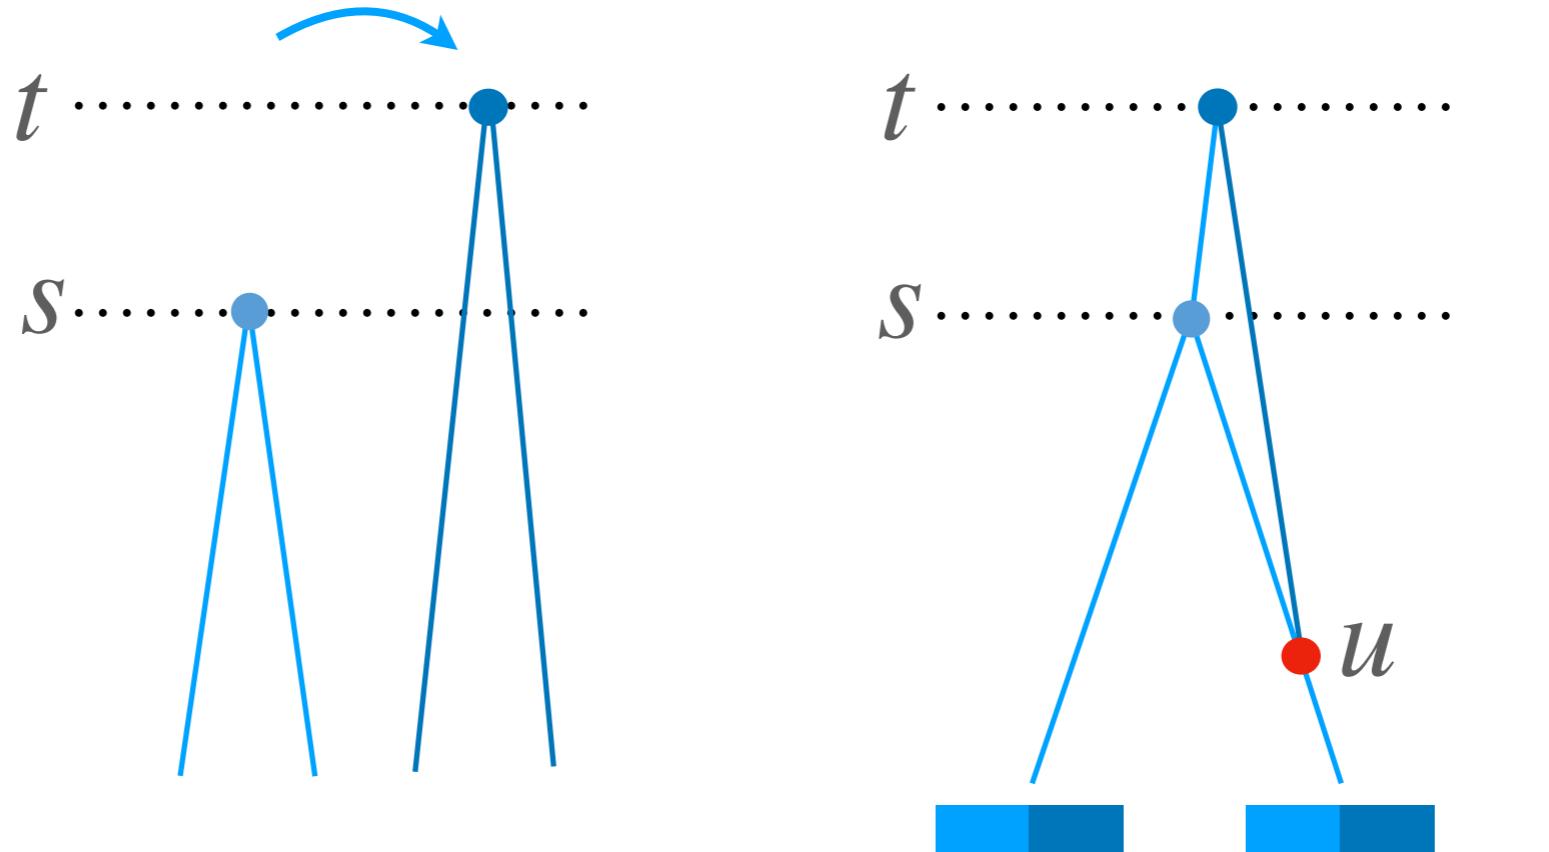




Transition probability from state s to state t → $p(t|s) = (1 - e^{-\rho t}) q(t|s) + e^{-\rho s} \delta(t - s)$

Prob. of change to state t upon recombination: → $q(t|s)$

Prob. of coalescing at time t Prob. of not coalescing between time u and t live (scattered) at time t



Transition probability from state s to state t → $p(t|s) = (1 - e^{-\rho t}) q(t|s) + e^{-\rho s} \delta(t - s)$

↓
Prob. of recombination ↓
Prob. of no recombination ↓
1 if $t = s$
else 0

Prob of change to state t upon recombination: → $q(t|s)$

Mean over recombination times u
(between 0 and t)

Prob. of coalescing at time t Prob. of not coalescing between time u and t

live (scattered) at time t

$$p(t|s) = (1 - e^{-\rho t}) q(t|s) + e^{-\rho s} \delta(t - s)$$

$$q(t|s) = \frac{1}{\lambda(t)} \int_0^{\min(s,t)} \frac{1}{s} \times e^{-\int_u^t \frac{1}{\lambda(v)} dv} du$$

↑
Prob. of coalescing
at time t

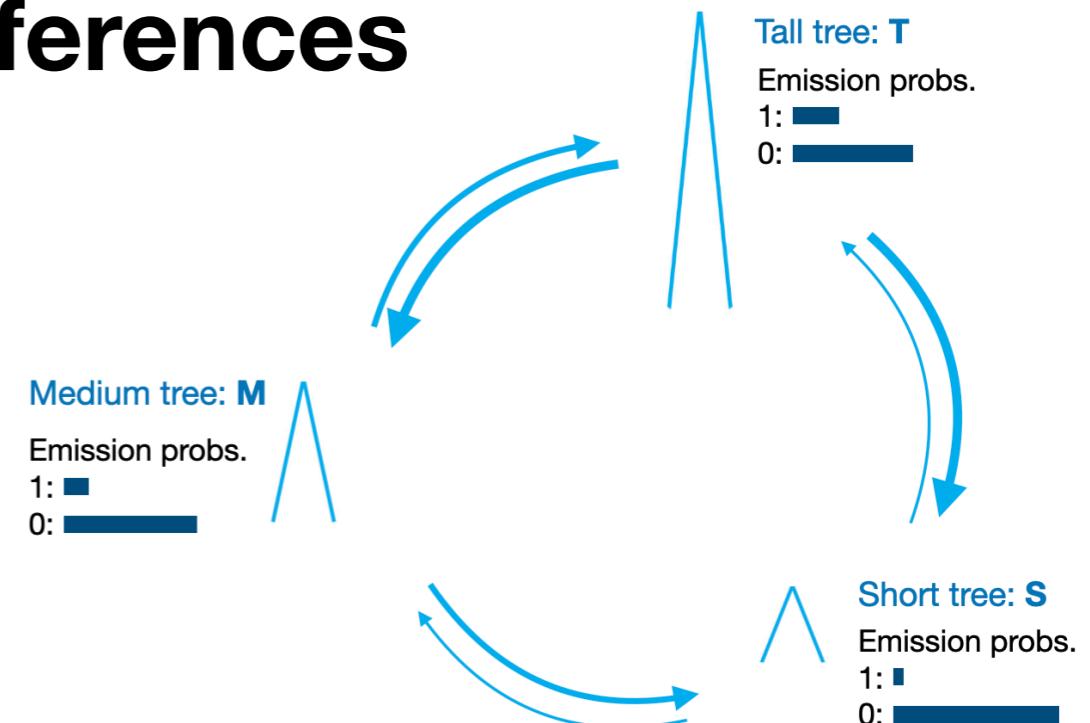
Prob. of not coalescing
between time u and t

$$\lambda(t) = \frac{N_e(t)}{N_0}$$

↑
N_e (scaled)
at time t

Hidden Markov model

Tree heights and sequence differences



One possible hidden state path:

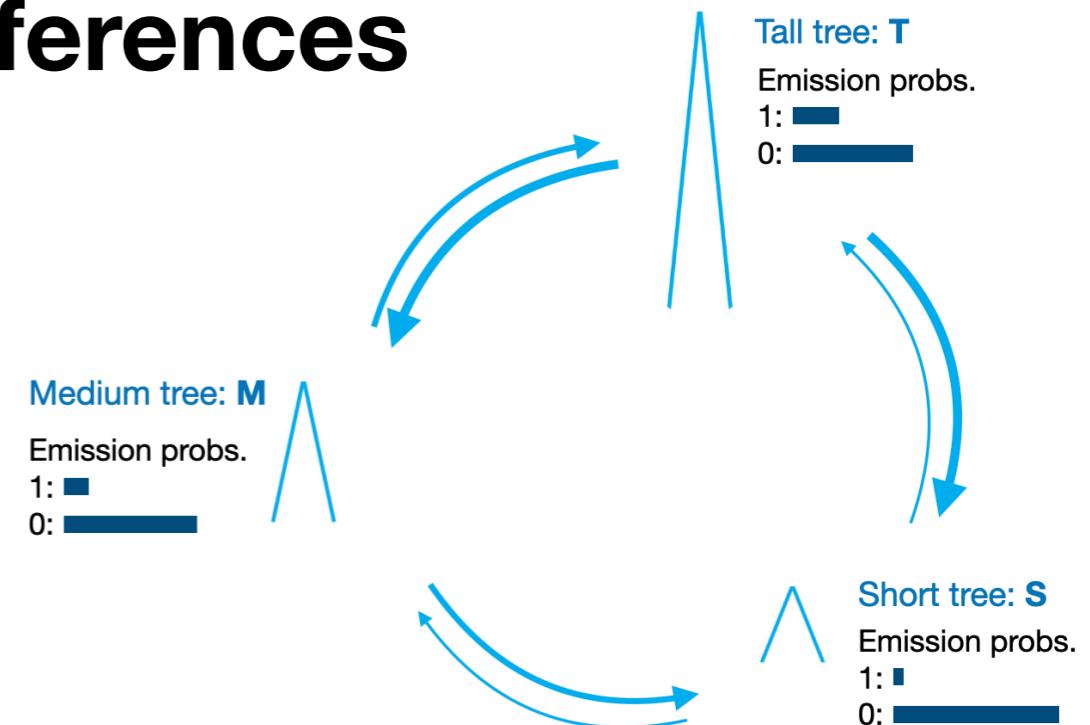
S → S → S → S → M → M → M → T → T → M
↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓
0 0 0 0 1 0 0 1 1 0

Another possible hidden state path:

M → M → S → S → S → S → S → T → T → T
↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓
0 0 0 0 1 0 0 1 1 0

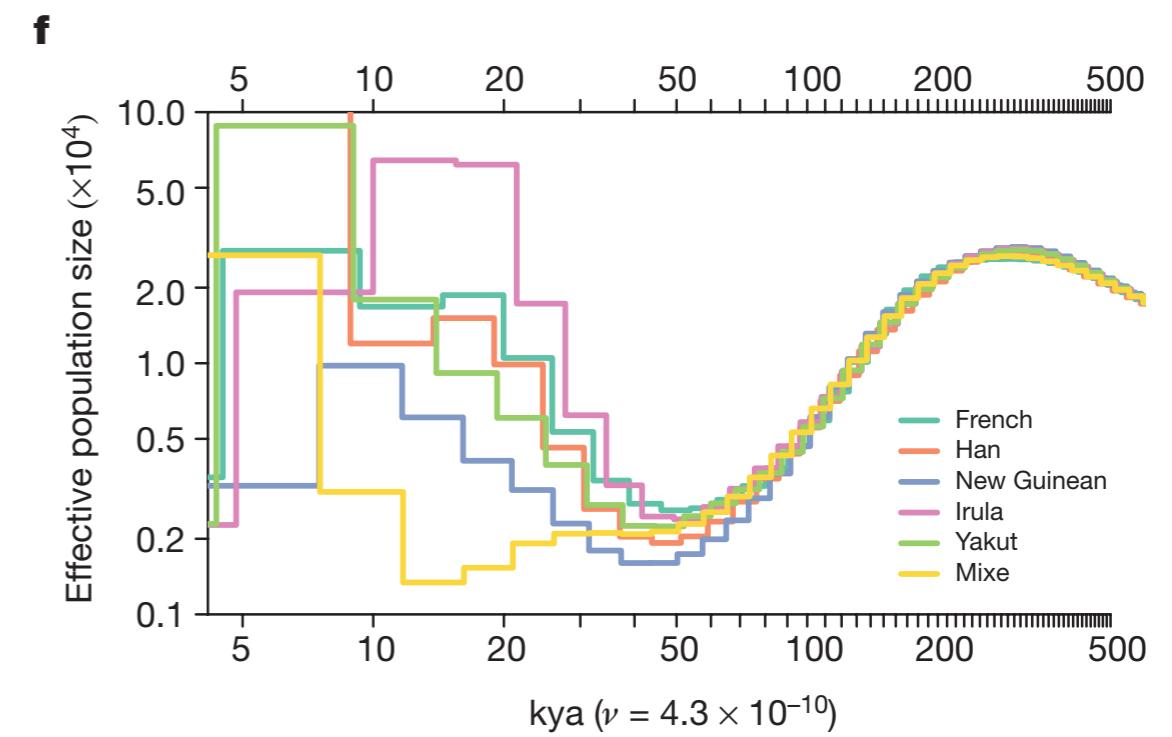
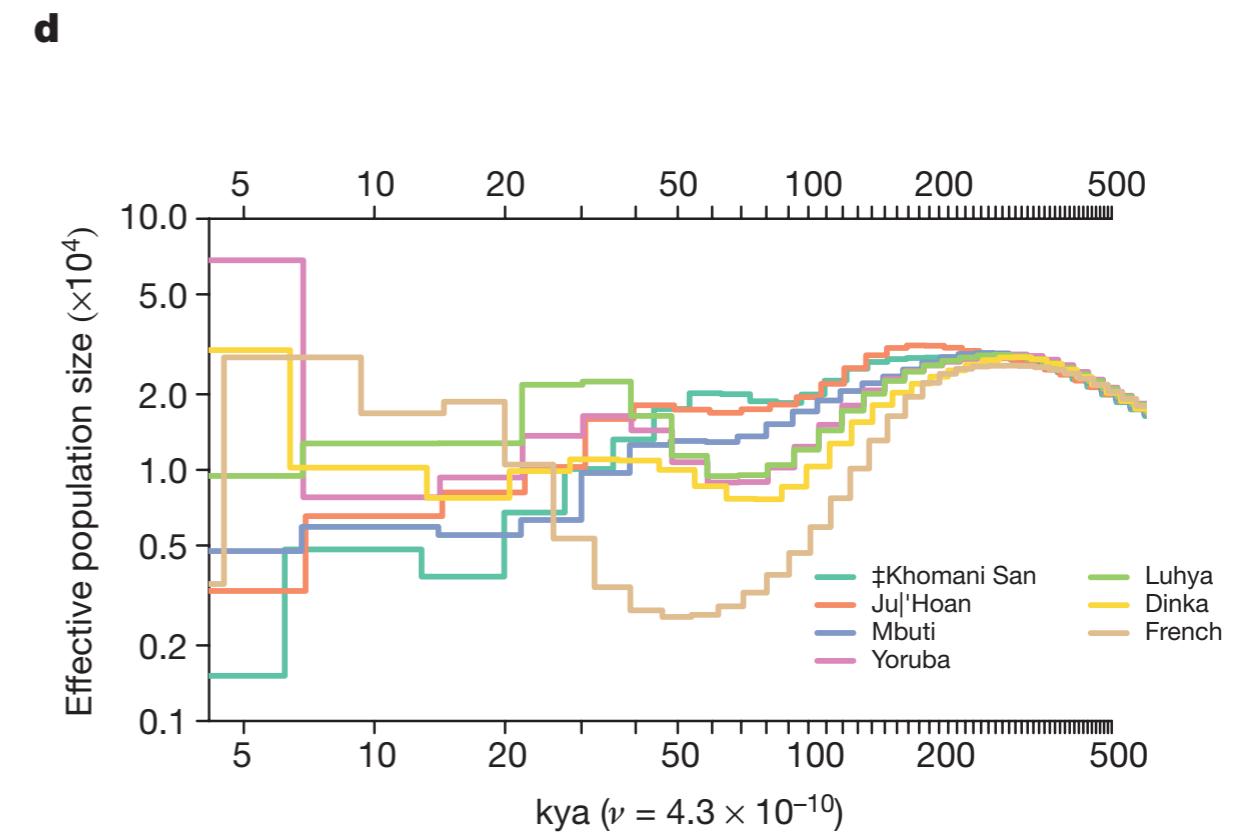
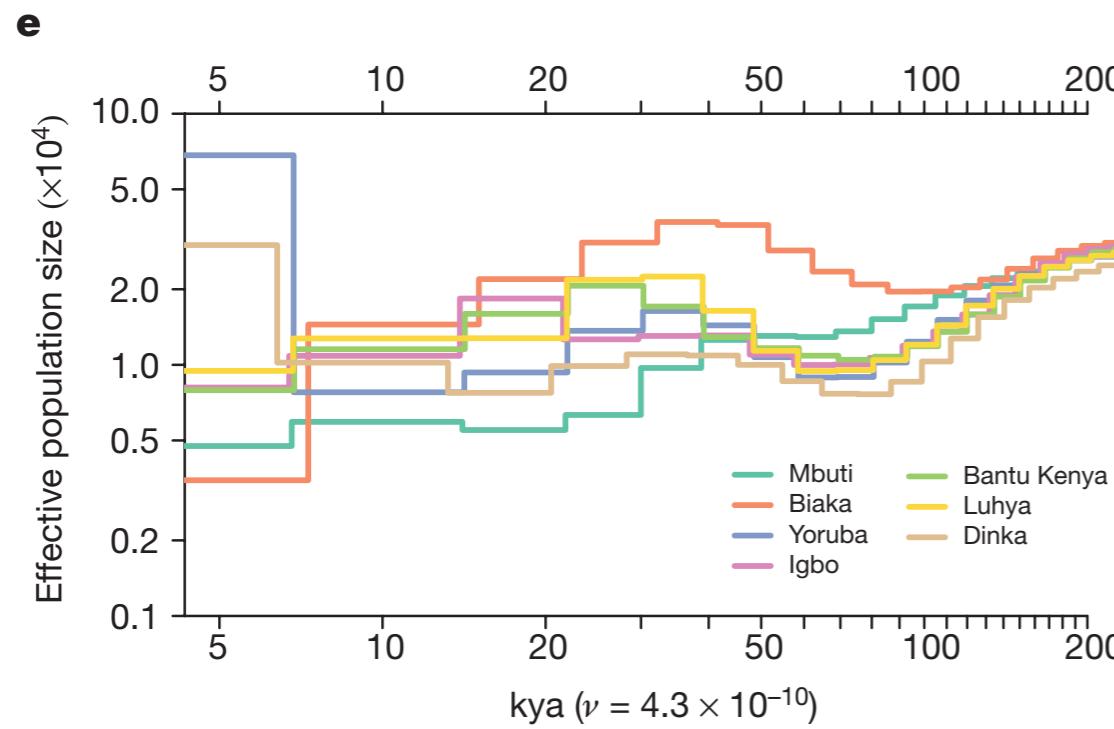
Hidden Markov model

Tree heights and sequence differences



- How do we find the past population sizes that makes our data most likely, across all possible hidden state paths?

SGDP paper



PSMC

